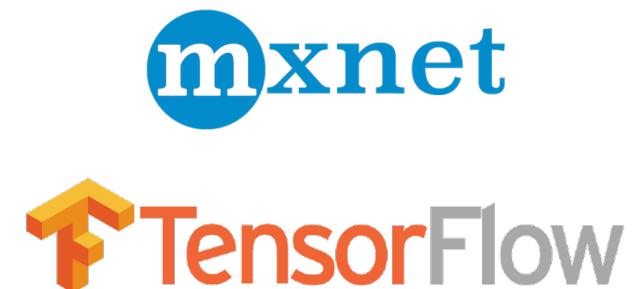


Benchmarking MXNet and TensorFlow

Presented by:

- **Matteo Paltenghi**
- **Kaoutar Chennaf**

Mentor: **Behrouz Derakhshan**



Agenda

1. Intro To Deep Learning (DL)
2. Mxnet And Tensorflow: History And Major Features
3. Api Architecture
4. Benchmark - Fundamental DL Operations
5. Benchmark - Complete Neural Network: Lenet
6. Summary And Future Work



1. Introduction to Deep Learning

Intro to Neural Network

Perceptron Model

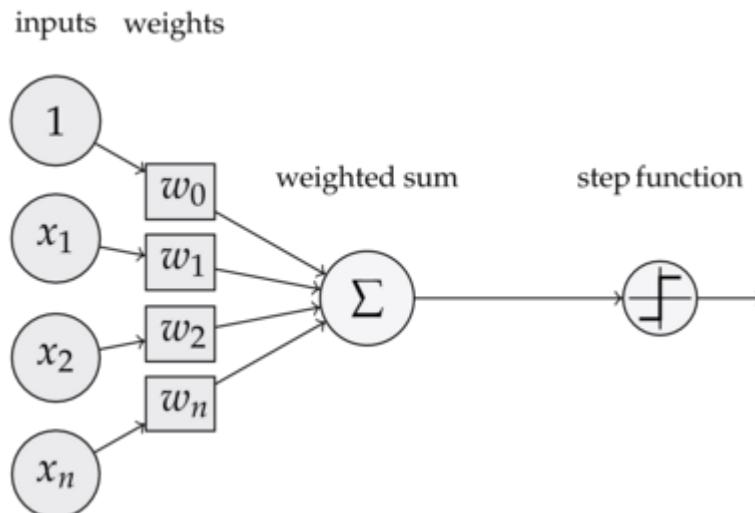


Image source: <https://blog.dbrgn.ch/2013/3/26/perceptrons-in-python/>

Multilayer Perceptron

(Feed Forward Neural Network = Signal flow only in one direction)
 (Deep Learning = when we have more than one hidden layer)

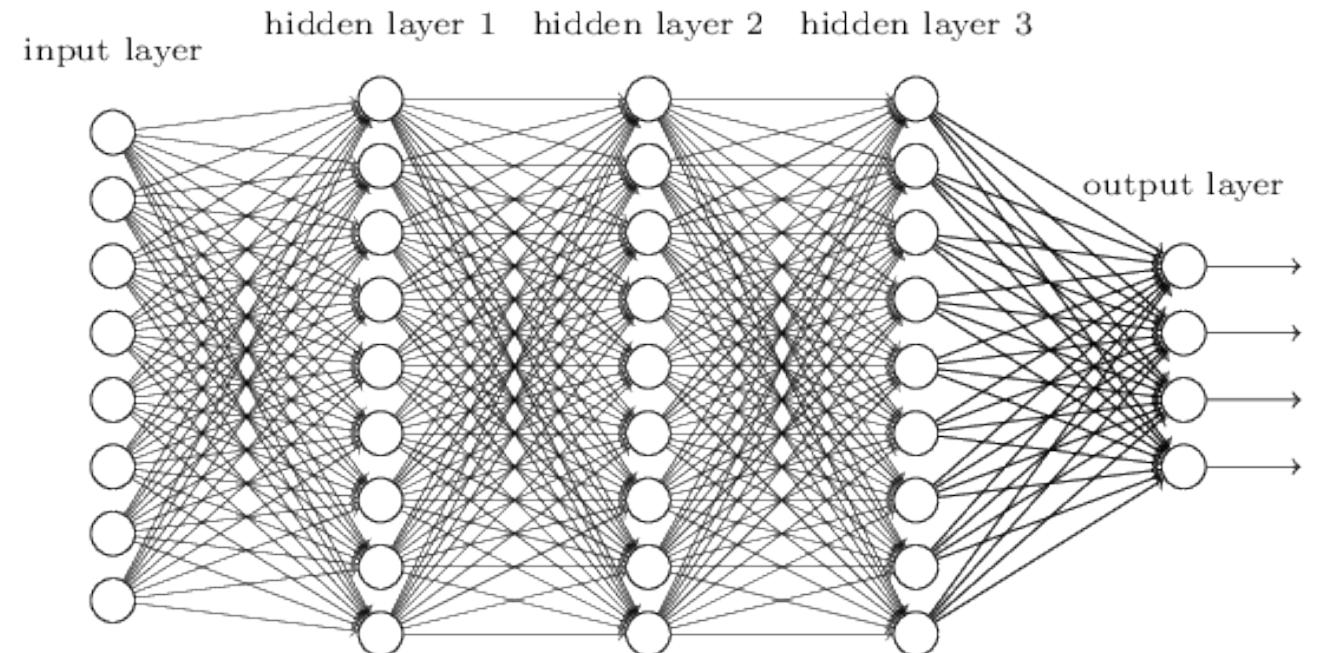


Image source: <http://neuralnetworksanddeeplearning.com/chap6.html>

Fundamental Operations

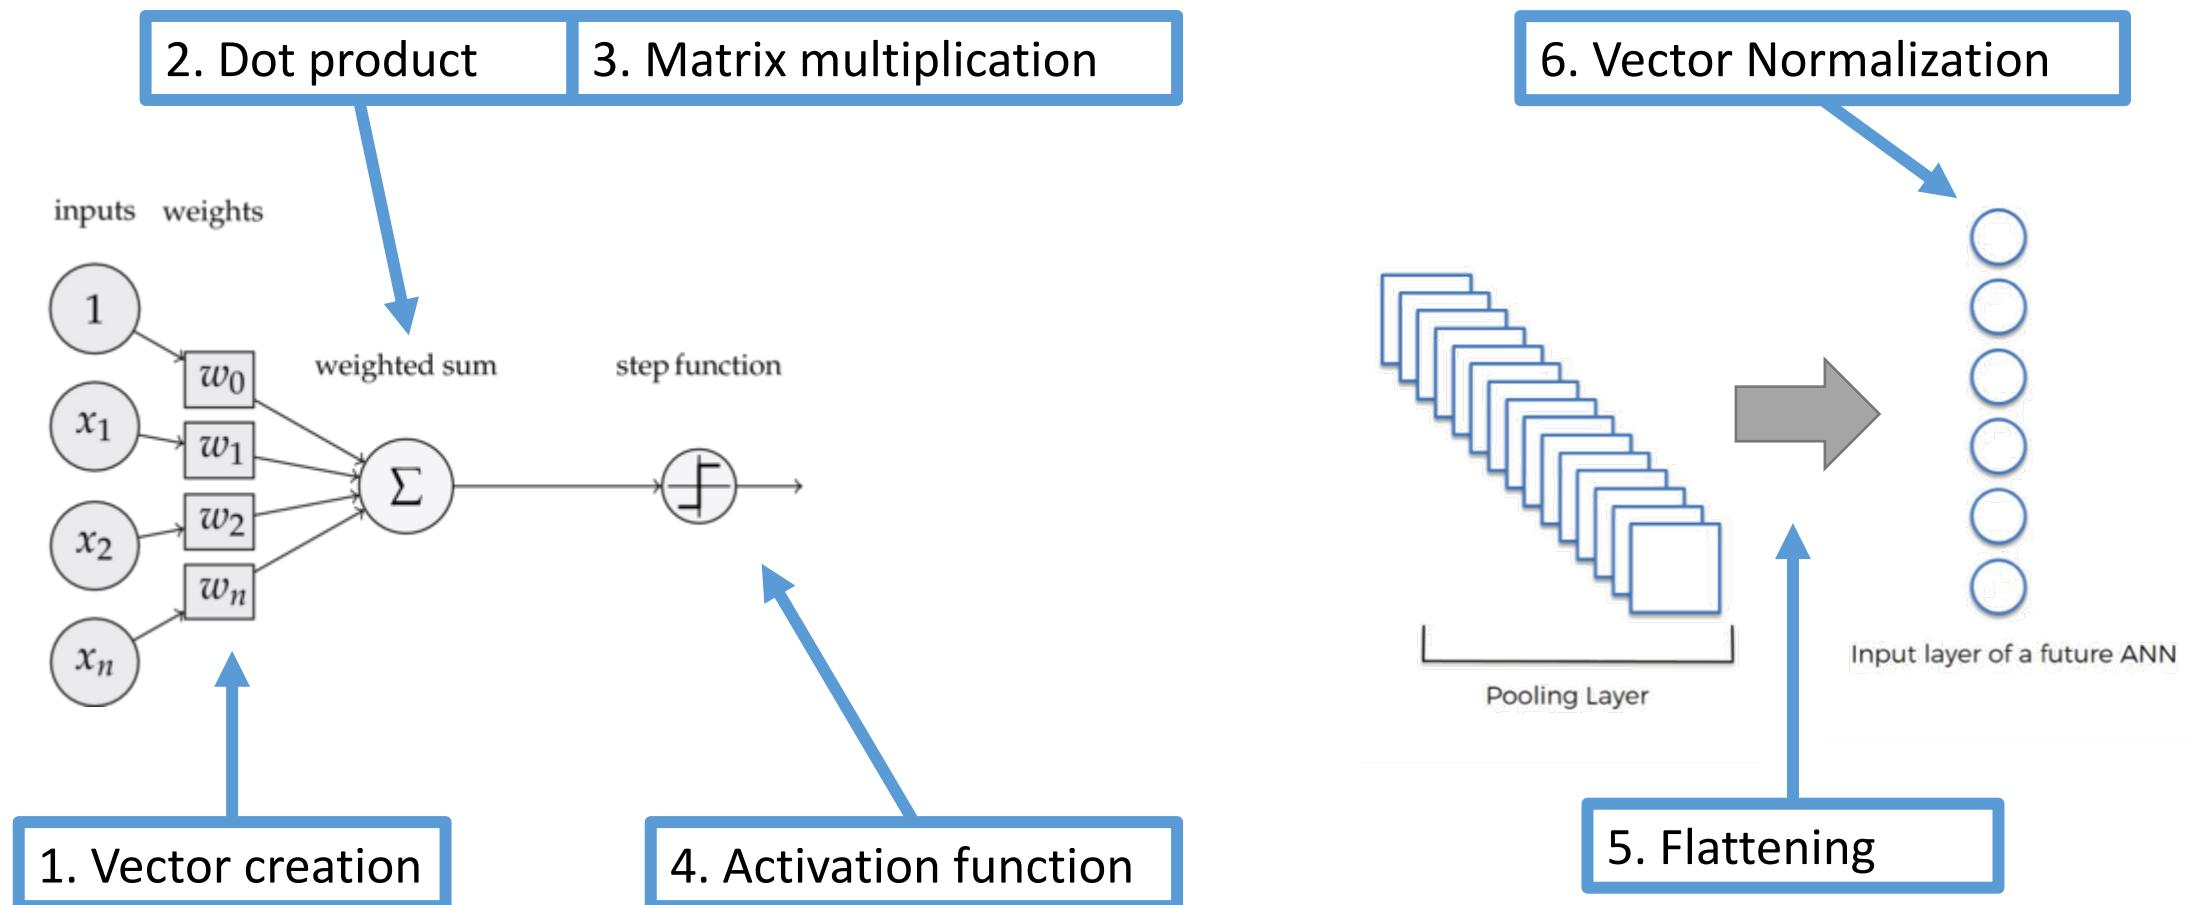


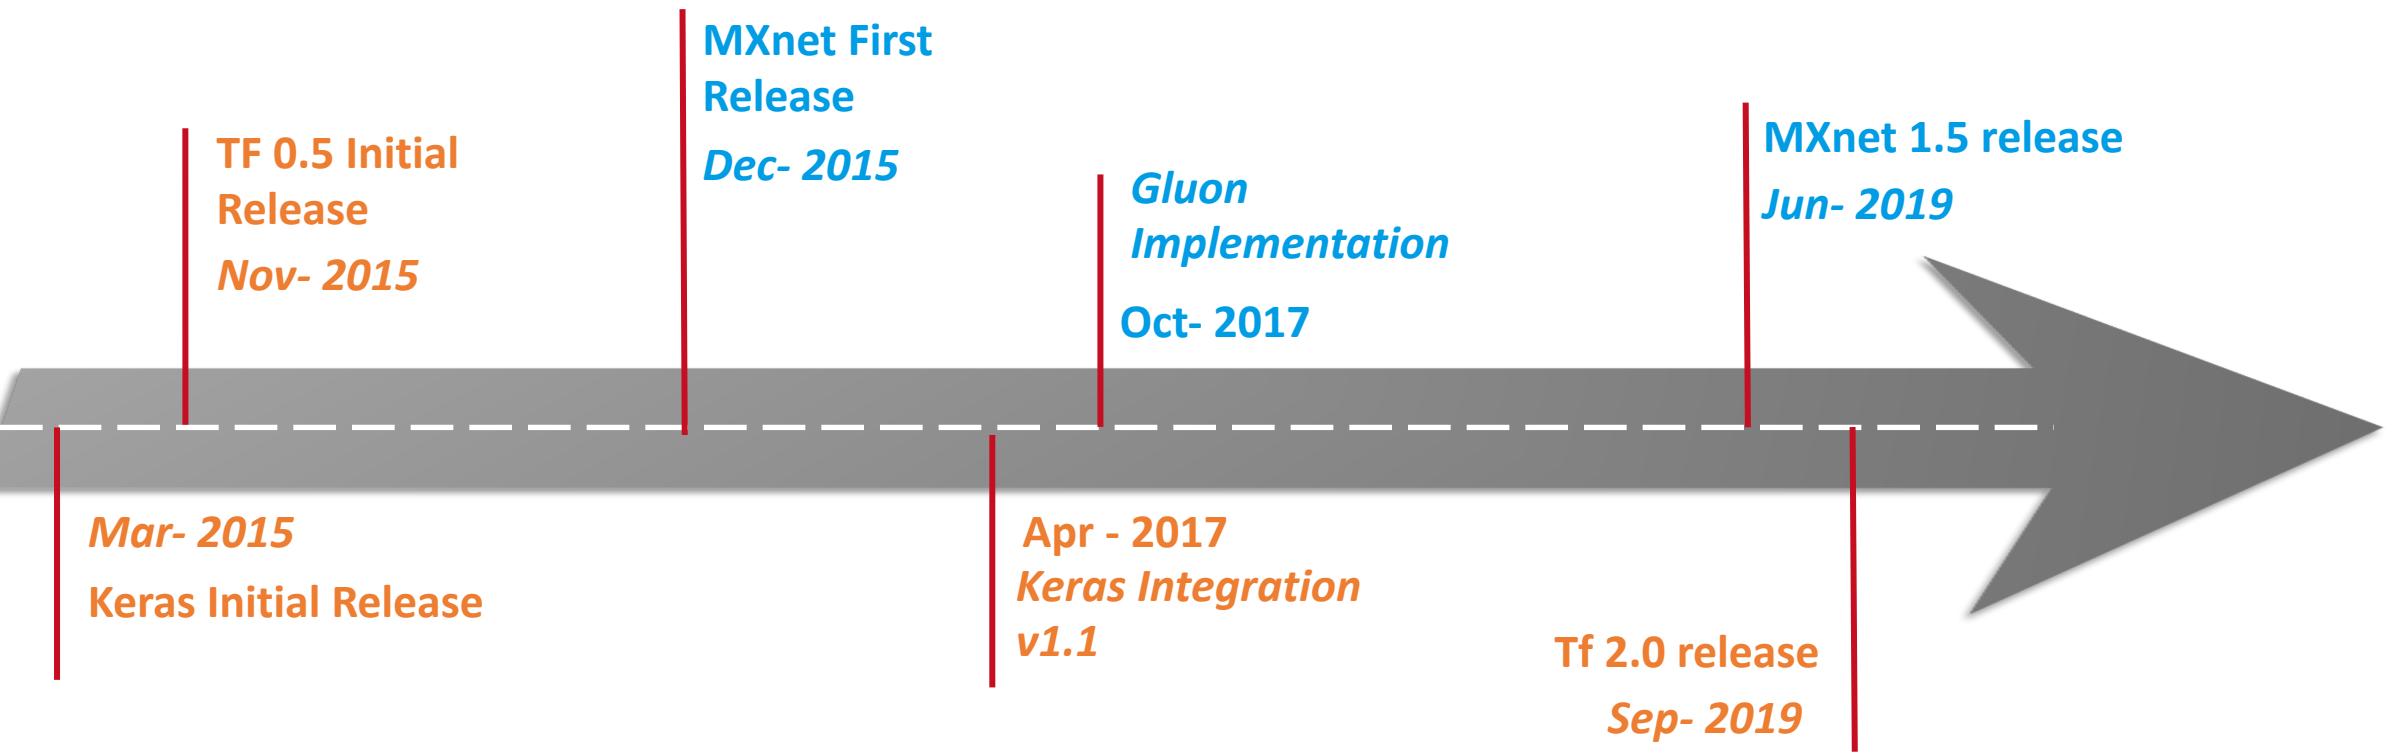
Image source: <https://blog.dbrgn.ch/2013/3/26/perceptrons-in-python/>

Image source: <https://www.superdatascience.com/blogs/convolutional-neural-networks-cnn-step-3-flattening>



2. History and Major Differences

System	Core Lang	Bindings Langs	Devices (beyond CPU)	Distributed	Imperative	Declarative
Tensorflow	C++	Python	GPU/TPU/Mobile	✓	✓ (with TF 2.0)	✓
MXNet	C++	Python/R/Julia/Go	GPU/Mobile	✓	✓	✓



Who is using



Image source: <https://www.hqew.net/news/news-36655>

Who is using mxnet



ABEJA, Inc.



Red Hat BIDS



smartbooks



Avito



Hewlett Packard
Enterprise



中国移动通信
CHINA MOBILE



Source: <https://enlyft.com/tech/products/mxnet>

Major Pros

Tensorflow

- It has a lot of documentation and guidelines;
- It offers monitoring for training processes of the models and visualization (**Tensorboard**);
- It's backed by a **large community** of devs and tech companies;
- It supports **TPU** when running on Google Cloud ServiceTensorflow Lite enables on-device inference with low latency for mobile devices;
- **Tensorflow JS** - enables deploying models in JavaScript environments, both frontend and Node.js backend.
 - TensorFlow.js also supports defining models in JavaScript and training them directly in the browser using a Keras-like API.

Mxnet

- It's quite **fast, flexible, and efficient** in terms of running DL algorithms;
- It features **advanced GPU support**, including multiple GPU mode;
- It has a **high-performance imperative API**;
- It's highly scalable;
- It provides rich support for **many programming languages**, such as Python, R, Scala, Javascript, and C++, among others;

Major Cons

TensorFlow

- It struggles with **poor results for speed** in benchmark tests compared with, for example, CNTK and MXNet,
- There is also one significant limitation: the **only fully supported language is Python**.

MXnet

- It has a much smaller community behind it compared with Tensorflow;
- It's **not so popular** among the research community.

Changes in TensorFlow



It brings us a bunch of exciting features, such as:

- It is possible to use **Keras inside TensorFlow**. It ensures that new Machine Learning models can be built with ease.
- Supports debugging your graphs and networks - TensorFlow 2.0 runs with eager execution by default for ease of use and smooth debugging. (**Imperative style supported**)
- **Robust model deployment** in production on any platform.
- Powerful experimentation for research.
- **Simplifying the API** by cleaning up deprecated APIs and reducing duplication.

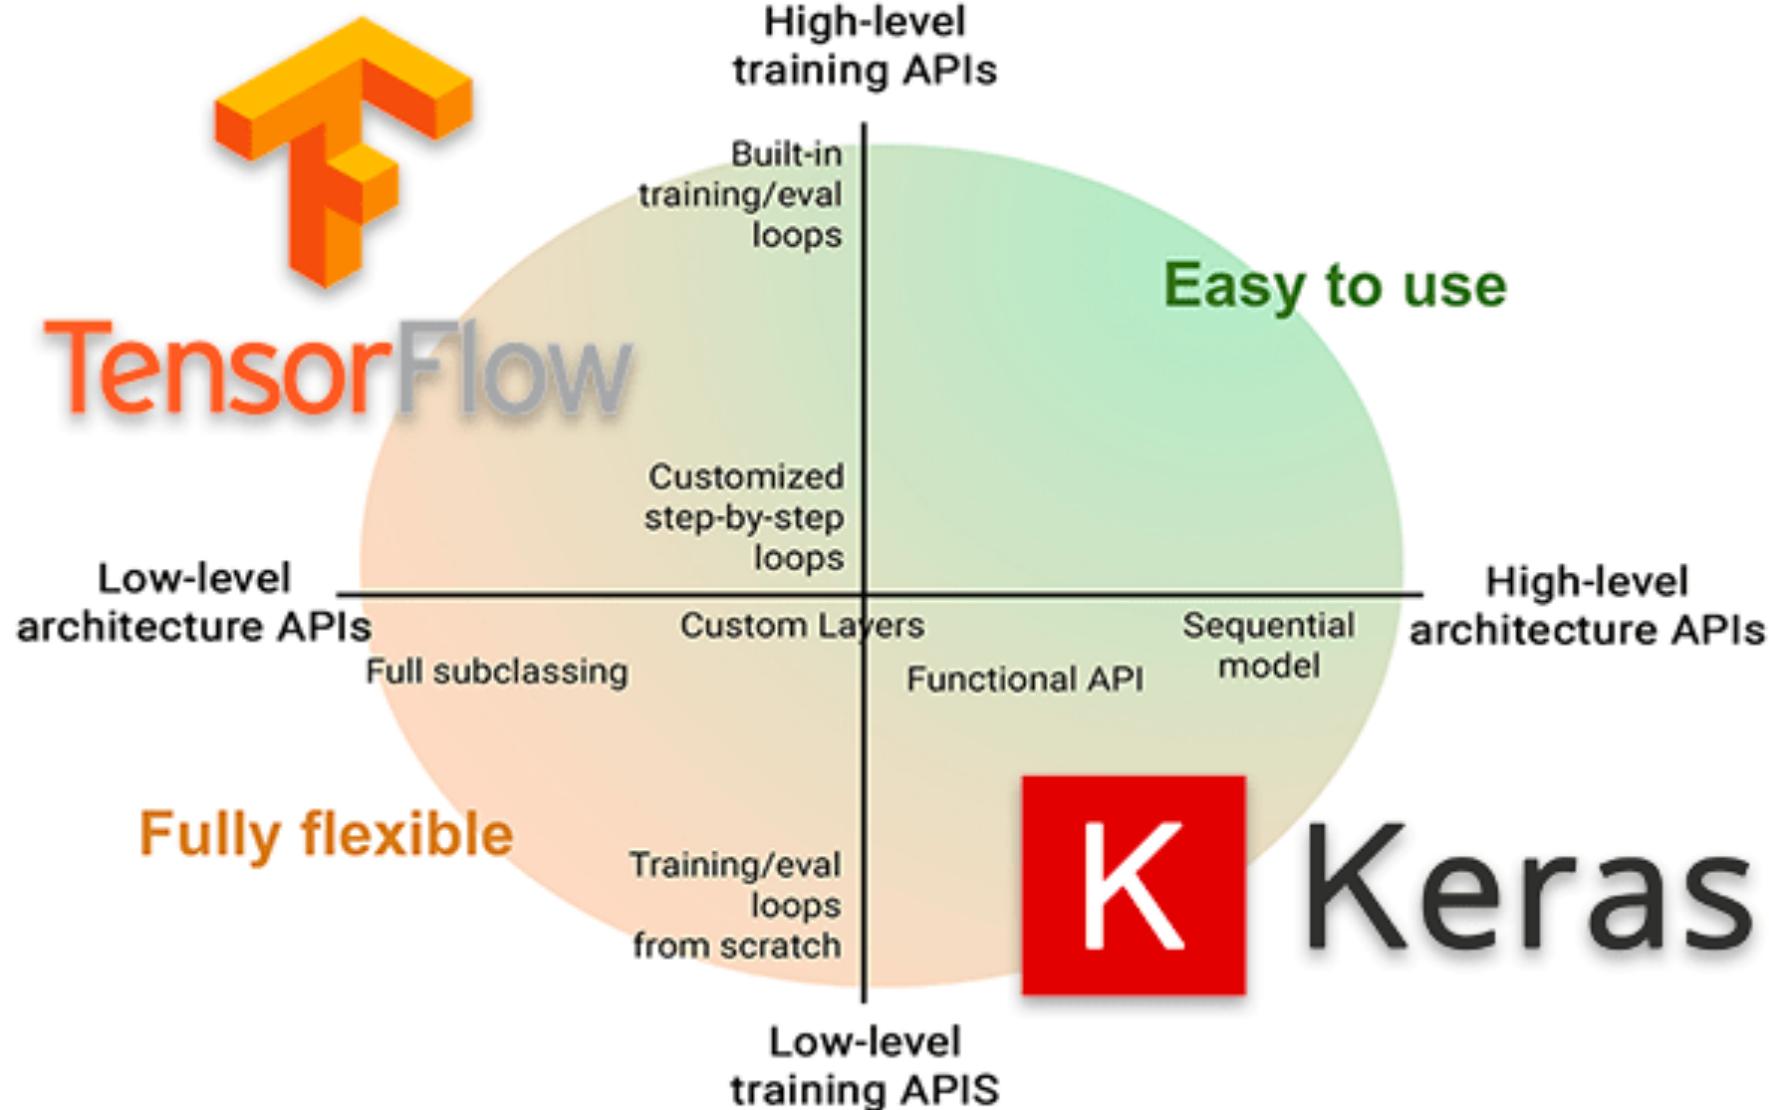


Image source: <https://www.pyimagesearch.com/2019/10/21/keras-vs-tf-keras-whats-the-difference-in-tensorflow-2-0/>



3. API Architectures

Imperative

```
# Imperative
import numpy as np

a = np.ones(10)
b = np.ones(10) * 2
c = b * a
d = c + 1
```

(+) Tend to be More Flexible

(+) Easy to try New Ideas (for Researches)

(+) Easy to specify Arbitrary Control Flow

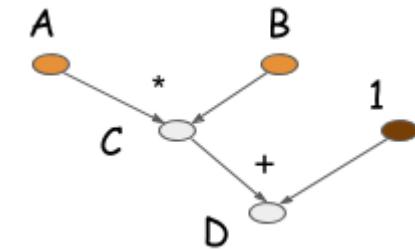
(-) Debugging happens during execution

(-) It can be More Difficult to Reuse

(-) Difficult to inspect, copy, or clone

Symbolic

```
# Symbolic
A = Variable('A')
B = Variable('B')
C = B * A
D = C + Constant(1)
# compiles the function
f = compile(D)
d = f(A=np.ones(10), B=np.ones(10)*2)
```



(+) Tend to be More Efficient

(+) Easy to inspect

(+) Run Extensive checks before Execution

(+) Easier to copy or clone

(+) Fit Most use cases

(-) Few Special use cases do not fit

There Is No Strict Boundary

[MXNet Documentation](#) : “it's possible to make an imperative program more like a traditional symbolic.” program or vice versa.

[TensorFlow Blog](#): “The two styles are fully interoperable as well, so you can mix and match (for example, you can nest one model type in another).”



Inventor of Keras (now part of TensorFlow 2.0)



François Chollet @fchollet · 16 ott 2018

The smarter thing to do is to blend imperative and symbolic paradigms like tf.Keras and Gluon do. You retain the core advantages of compiled graphs in most situations, and you still have access to the flexibility of dynamic graphs when you need it.

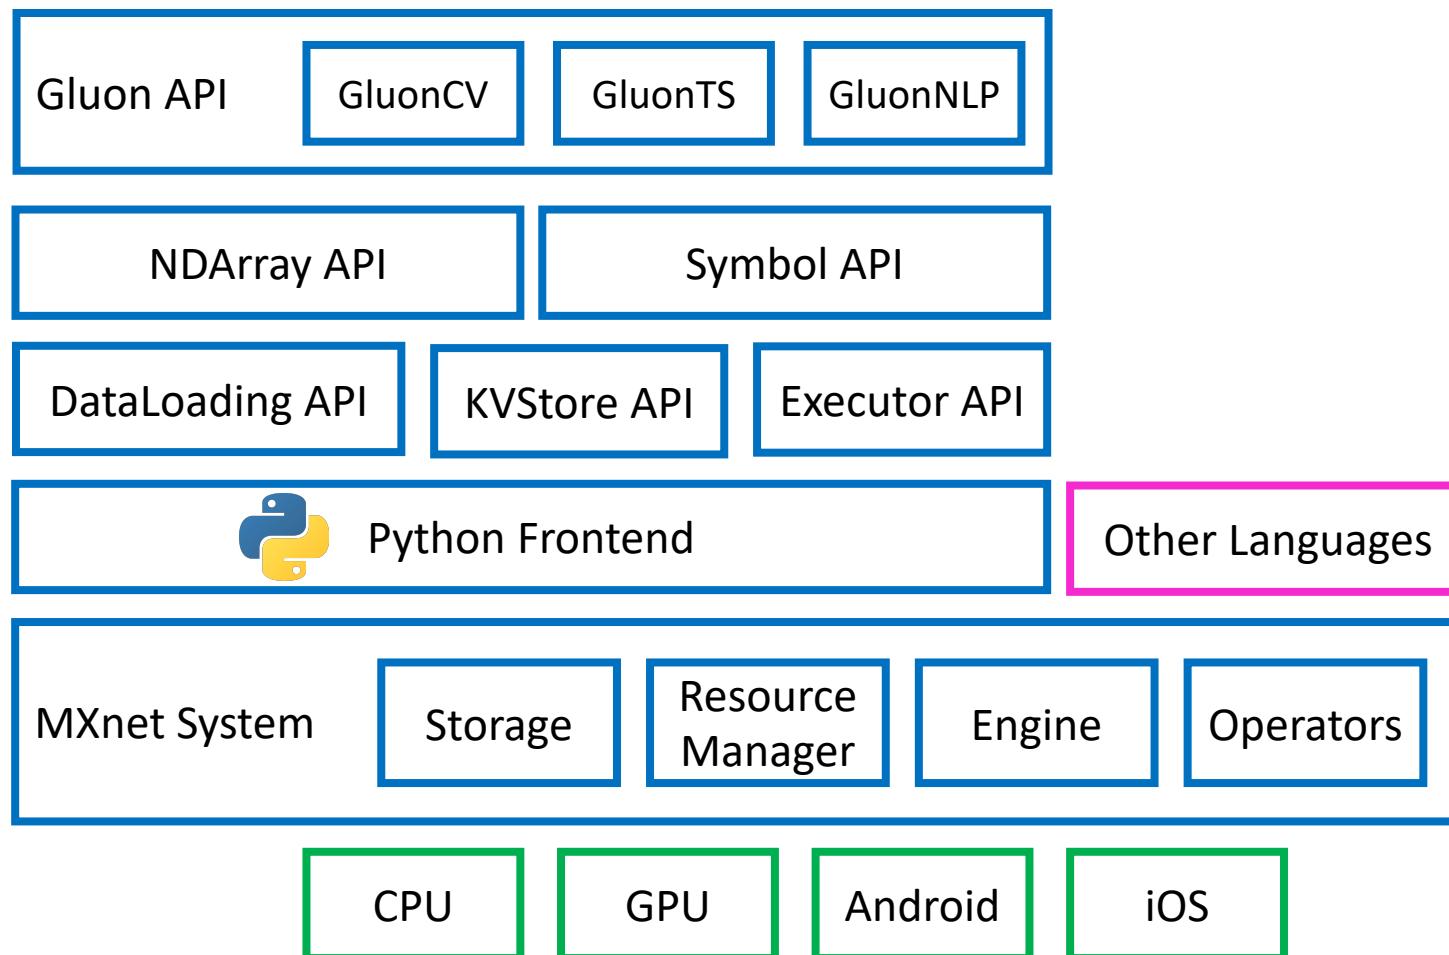
3

3

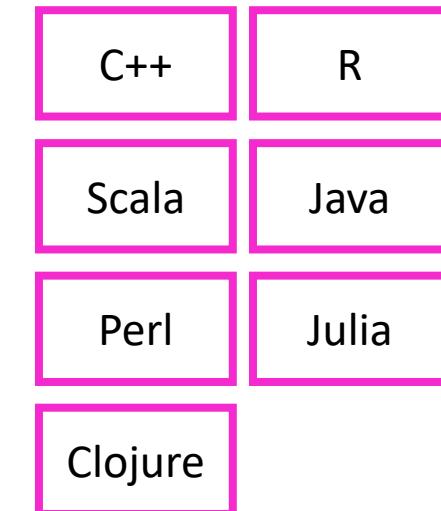
22

↑

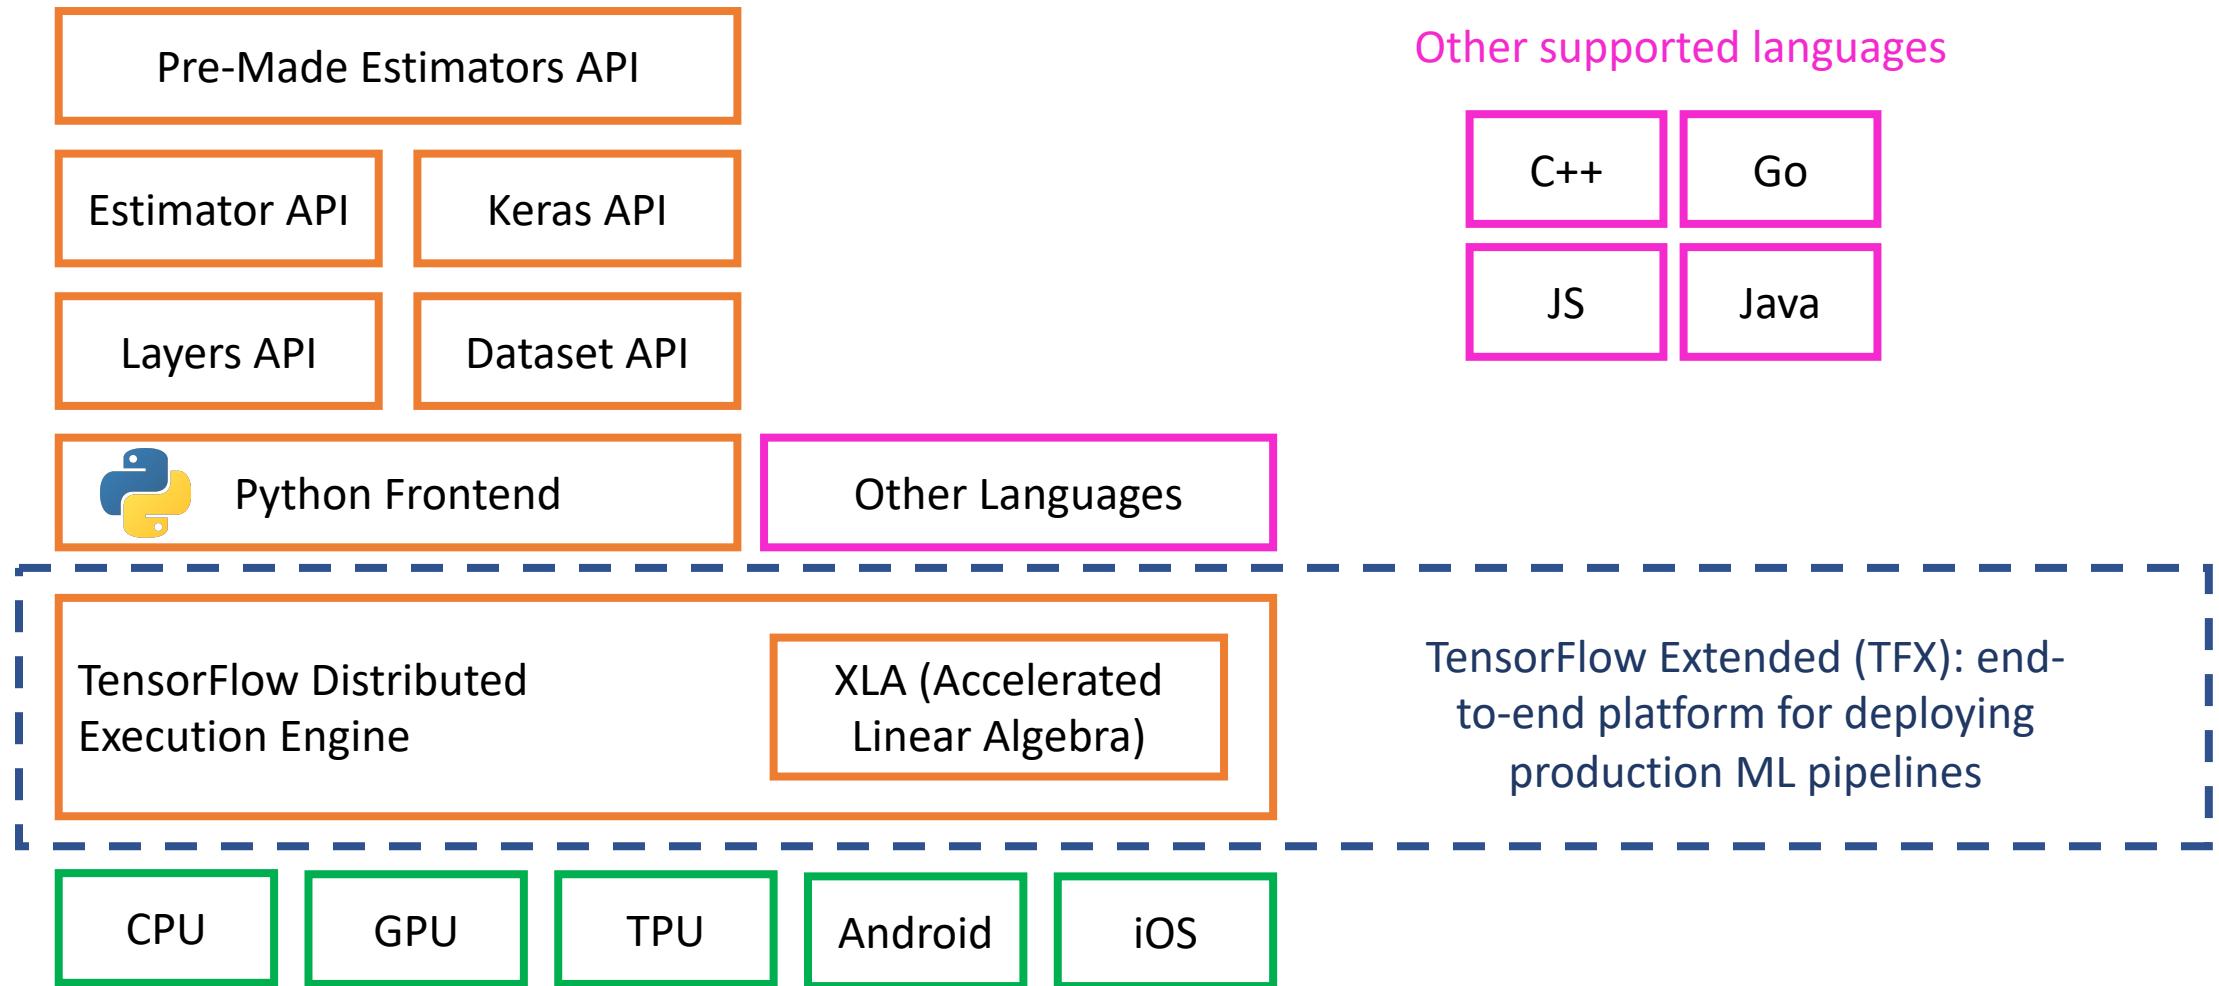
MXnet API Architecture



Other supported languages



TensorFlow API Architecture



Fundamental Operations

Operation	MXNet	TensorFlow
1. Vector creation	<code>mx.ndarray.from_numpy(np_vector)</code>	<code>tf.convert_to_tensor(np_vector)</code>
2. Dot Product	<code>mx.ndarray.dot(first, second)</code>	<code>tf.tensordot(first, tf.transpose(second), axes=1)</code>
3. Matrix Multiplication	<code>mx.nd.linalg.gemm2(first, second)</code>	<code>tf.tensordot(first, second, axes = [[1], [0]])</code>
4. Activation Function	<code>mx.ndarray.Activation(input_vector, act_type = "sigmoid")</code>	<code>tf.keras.activations.sigmoid(input_vector)</code>
5. Flattening	<code>matrix.reshape(shape=(-1,))</code>	<code>tf.reshape(matrix, [-1])</code>
6. Normalization	<code>matrix/mx.nd.norm(matrix, ord=2, axis=None)</code>	<code>tf.linalg.normalize(matrix, ord=2, axis=None)[0]</code>

MLP Construction

```
1 import numpy as np
2 from keras.models import Sequential
3 from keras.layers import Dense
4 from keras import optimizers
5
6 # Get your data
7 x_train = np.random.random((1000, 20))
8 y_train = np.random.randint(2, size=(1000, 1))
9
10 # Build the model
11 model = Sequential()
12 model.add(Dense(64, input_dim=20, activation='relu'))
13 model.add(Dense(64, activation='relu'))
14 model.add(Dense(1, activation='sigmoid'))
15
16 # Configures the model for training
17 model.compile(loss='binary_crossentropy',
18                 optimizer = optimizers.SGD(lr=0.01),
19                 metrics=['accuracy'])
20
21 # Start the training loop
22 model.fit(x_train, y_train,
23             epochs=20,
24             batch_size=128)
```

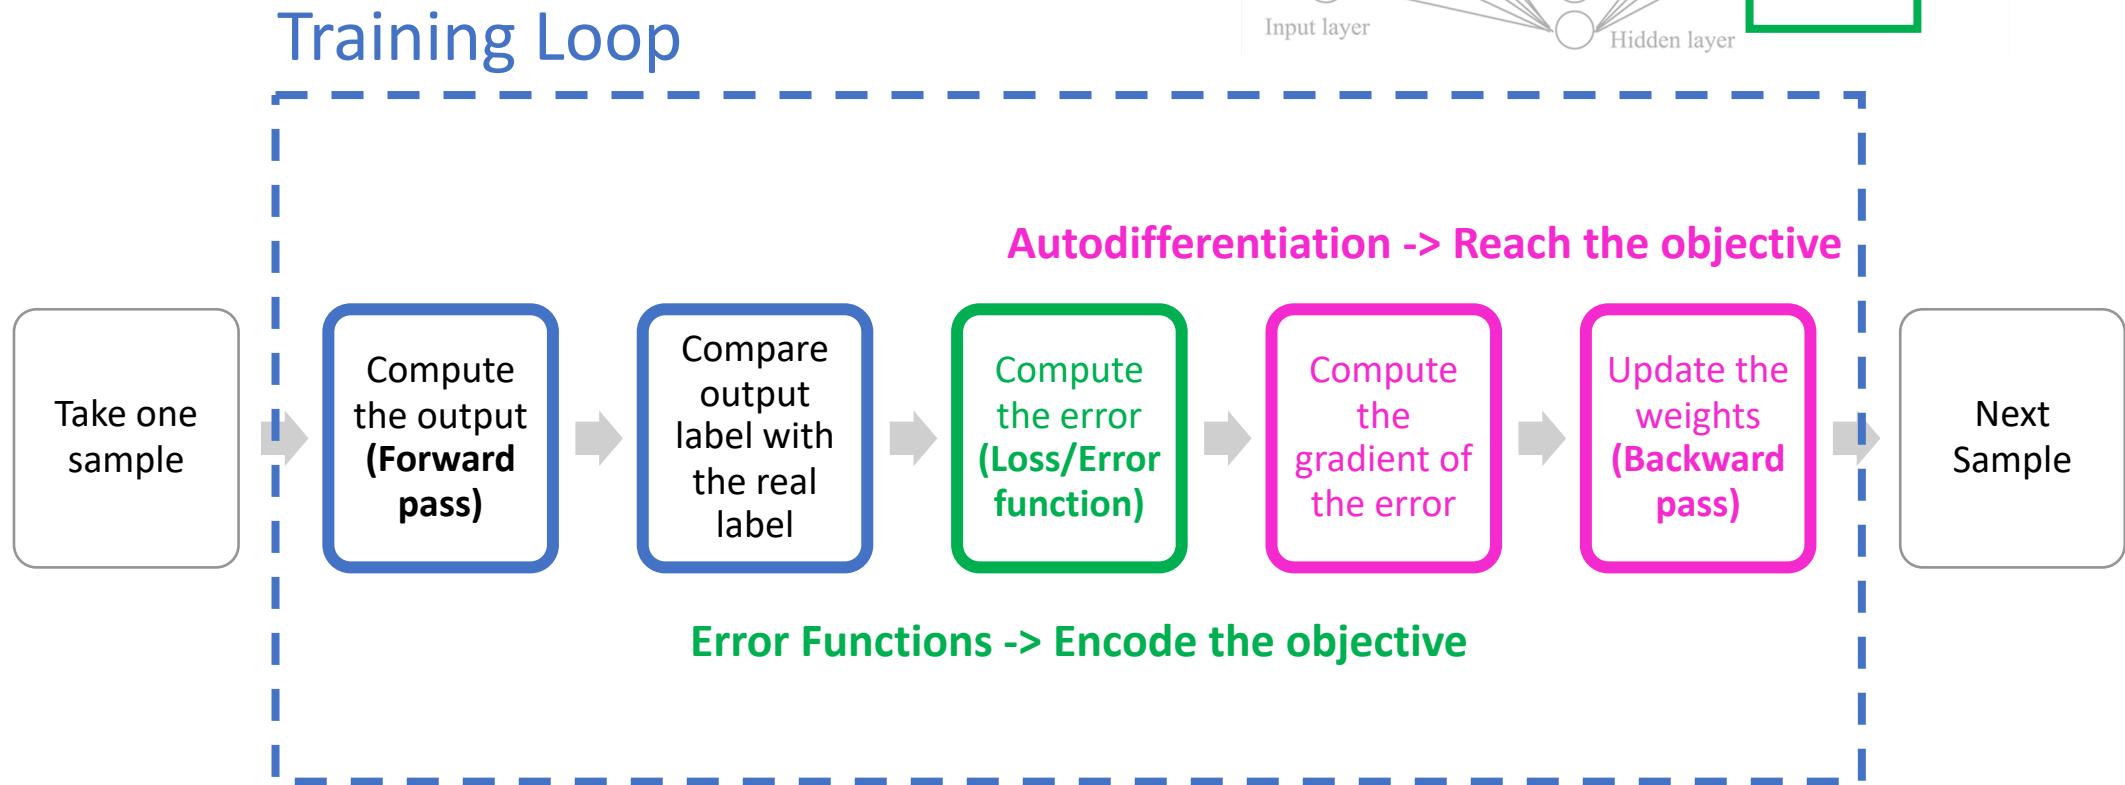


Keras

```
1 import mxnet as mx
2 from mxnet import gluon
3 from mxnet.gluon.contrib.estimator import estimator
4
5 # Get your data
6 x_train = np.random.random((1000, 20))
7 y_train = np.random.randint(2, size=(1000, 1))
8 dataset = mx.gluon.data.dataset.ArrayDataset(x_train, y_train)
9 train_data_loader = gluon.data.DataLoader(dataset, batch_size=128)
10
11 # Build the model
12 model = gluon.nn.Sequential()
13 with model.name_scope():
14     model.add(gluon.nn.Dense(64, activation="relu"))
15     model.add(gluon.nn.Dense(64, activation="relu"))
16     model.add(gluon.nn.Dense(1, activation="sigmoid"))
17
18 # Configures the model for training
19 loss_fn = gluon.loss.SigmoidBinaryCrossEntropyLoss
20 trainer = gluon.Trainer(model.collect_params(), 'sgd', {'learning_rate': 0.01})
21 train_acc = mx.metric.Accuracy() # Metric to monitor
22
23 # Define the estimator, by passing to it
24 # the model, loss function, metrics, trainer object and context
25 est = estimator.Estimator(net = model,
26                           loss = loss_fn,
27                           metrics = train_acc,
28                           trainer = trainer)
29
30 # Start the training loop
31 est.fit(train_data = train_data_loader, epochs=20)
```



Inside the fit() method



Error Functions

y(weights) = functions to be derivated

Mean Square Error (regression):

$$= \frac{\sum_{i=1}^n (y_i - \hat{y}_i)^2}{n}$$

Binary Cross Entropy (classification):

$$= -(y_i \log(\hat{y}_i) + (1 - y_i) \log(1 - \hat{y}_i))$$



- binary_crossentropy(...)
- categorical_crossentropy(...)
- categorical_hinge(...)
- cosine_similarity(...)
- hinge(...)
- kullback_leibler_divergence(...)
- logcosh(...) mean_absolute_error(...)
- mean_absolute_percentage_error(...)
- mean_squared_error(...)
- mean_squared_logarithmic_error(...)
- poisson(...)
- sparse_categorical_crossentropy(...)
- squared_hinge(...)



- L2Loss()
- L1Loss()
- SigmoidBinaryCrossEntropyLoss()
- SigmoidBCELoss
- SoftmaxCrossEntropyLoss()
- SoftmaxCELoss
- KLDivLoss()
- CTCLoss()
- HuberLoss()
- HingeLoss()
- SquaredHingeLoss()
- LogisticLoss()
- TripletLoss()
- PoissonNLLLoss()
- CosineEmbeddingLoss()

Automatic Differentiation implementation

mxnet.autograd API

What does it do? “It expedites the gradient calculation by automatically calculating derivatives.”

```
from mxnet import nd
from mxnet import autograd

x = nd.array([[1, 2], [3, 4]])
# tell what to monitor
x.attach_grad()
# record operation
with autograd.record():
    # do the operation you want
    y = 2 * x * x
# compute backward pass
y.backward()
# derivative is stored in:
x.grad
```

tf.GradientTape API

What does it do? “It computes the gradient of a computation with respect to its input variables.”

```
import tensorflow as tf

x = tf.constant([[1, 2], [3, 4]])
# record the tape
with tf.GradientTape() as tape:
    # tell what to monitor
    tape.watch(x)
    # do the operation you want
    y = 2 * x * x
# read the gradient from the tape
dy_dx = tape.gradient(y, x)
```

Variables are automatically watched by GradientTape



4. Microbenchmark results

Fundamental Operations

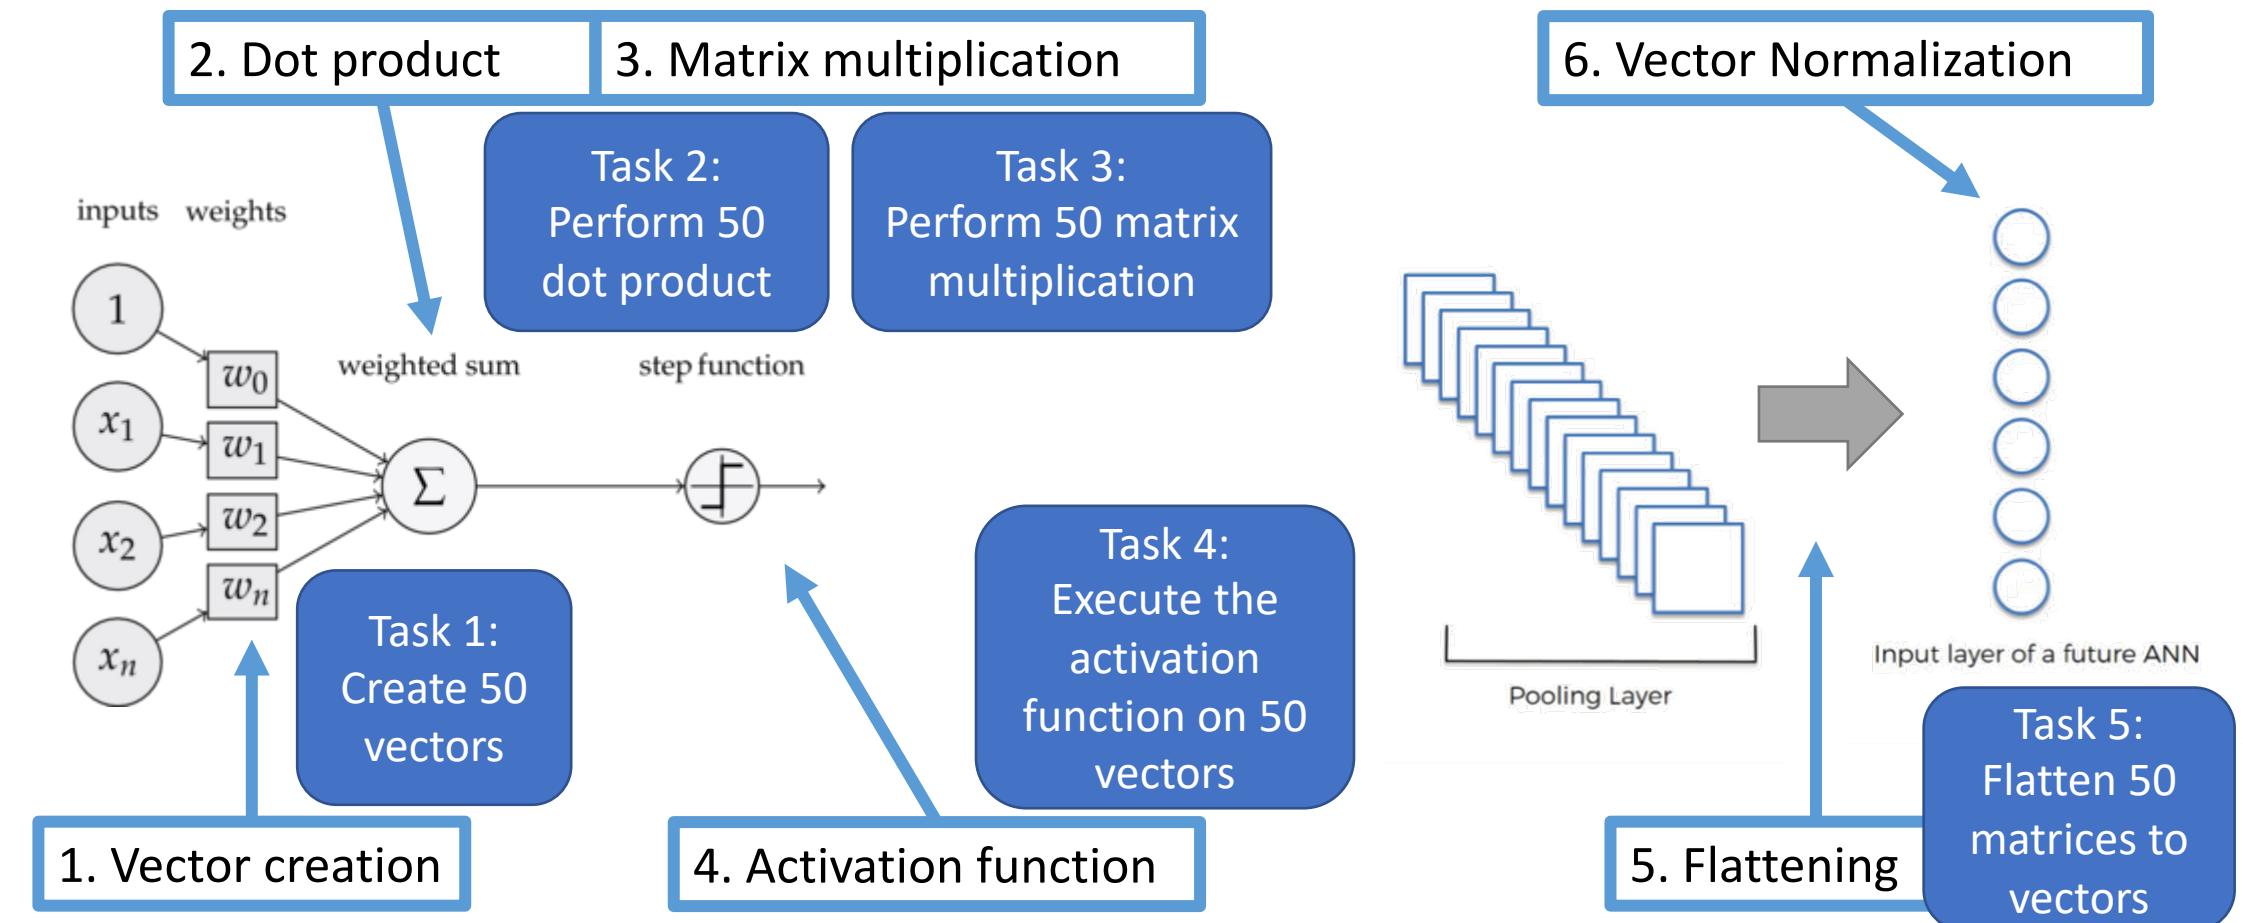


Image source: <https://blog.dbrgn.ch/2013/3/26/perceptrons-in-python/>

Image source: <https://www.superdatascience.com/blogs/convolutional-neural-networks-cnn-step-3-flattening>

Tasks 1 – Vector Creation

Task 1:
Create 50
vectors

Task 1.1:
from Numpy



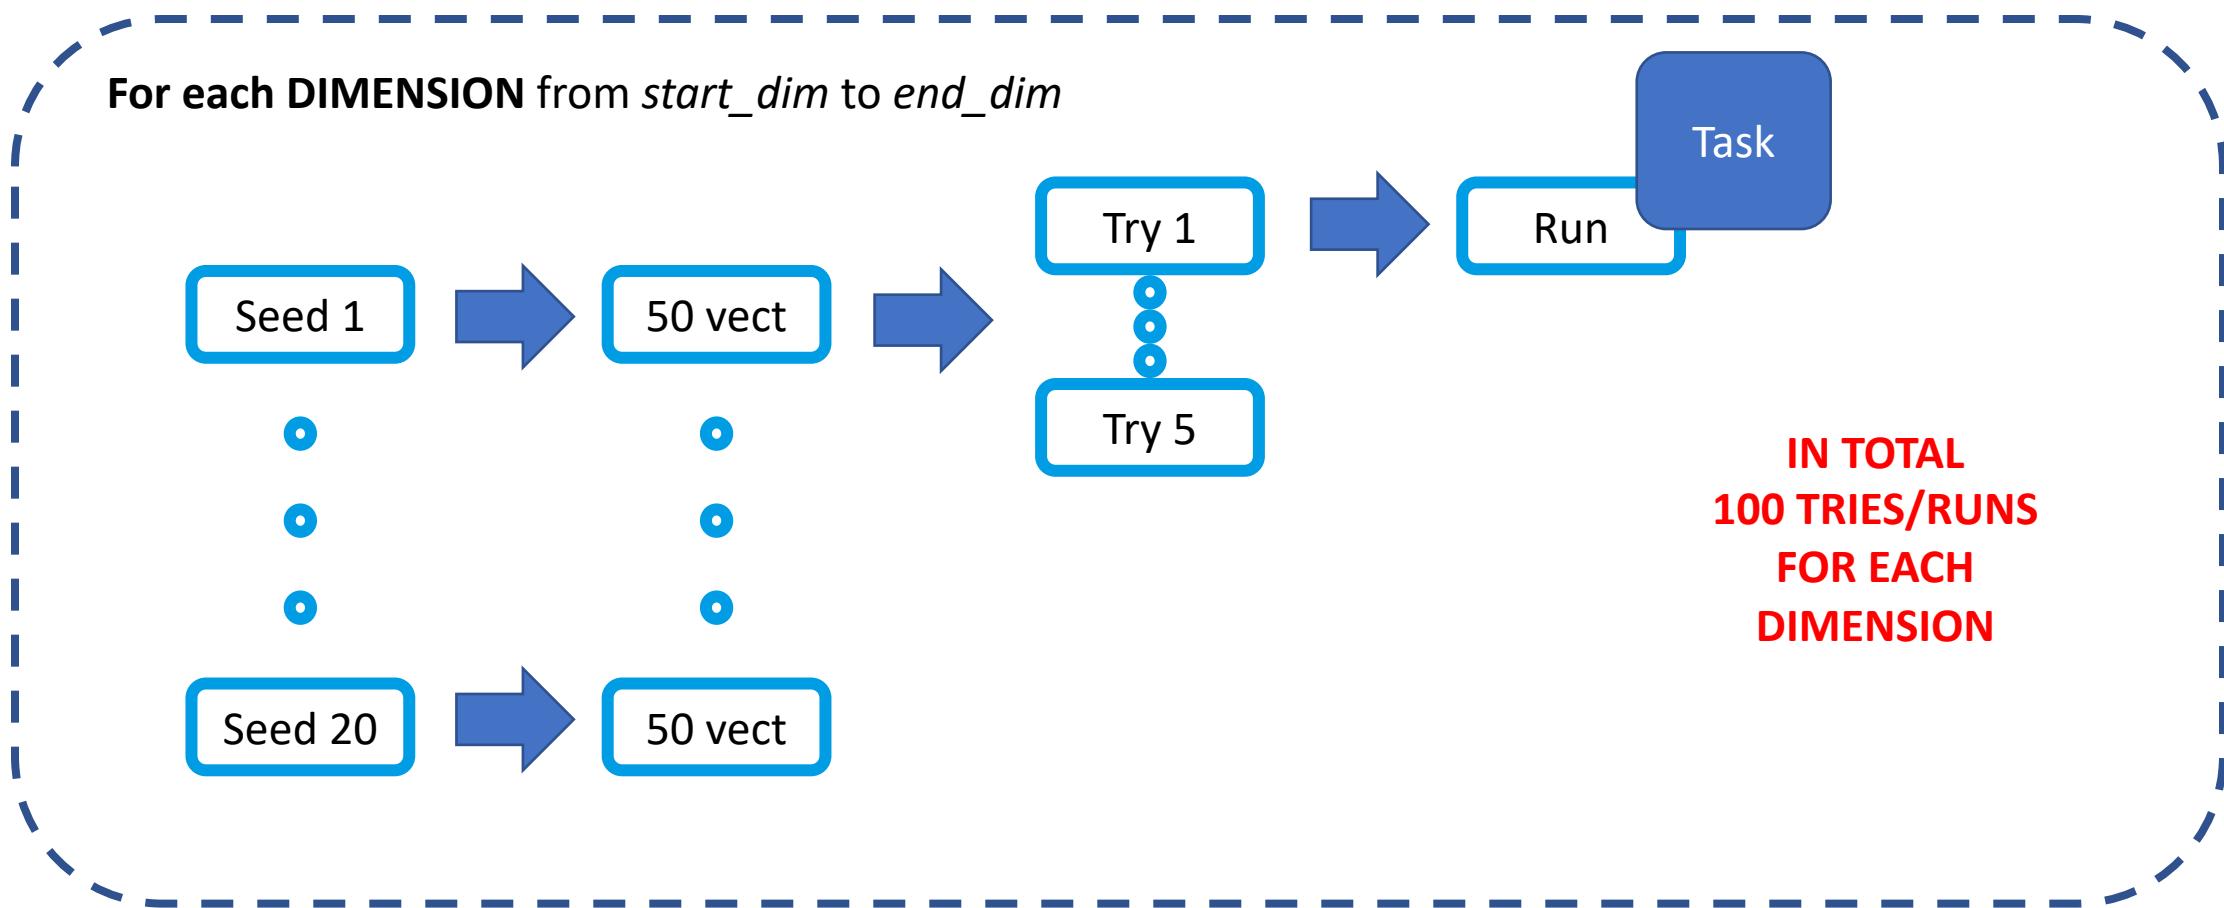
NumPy

Task 4.2:
From List

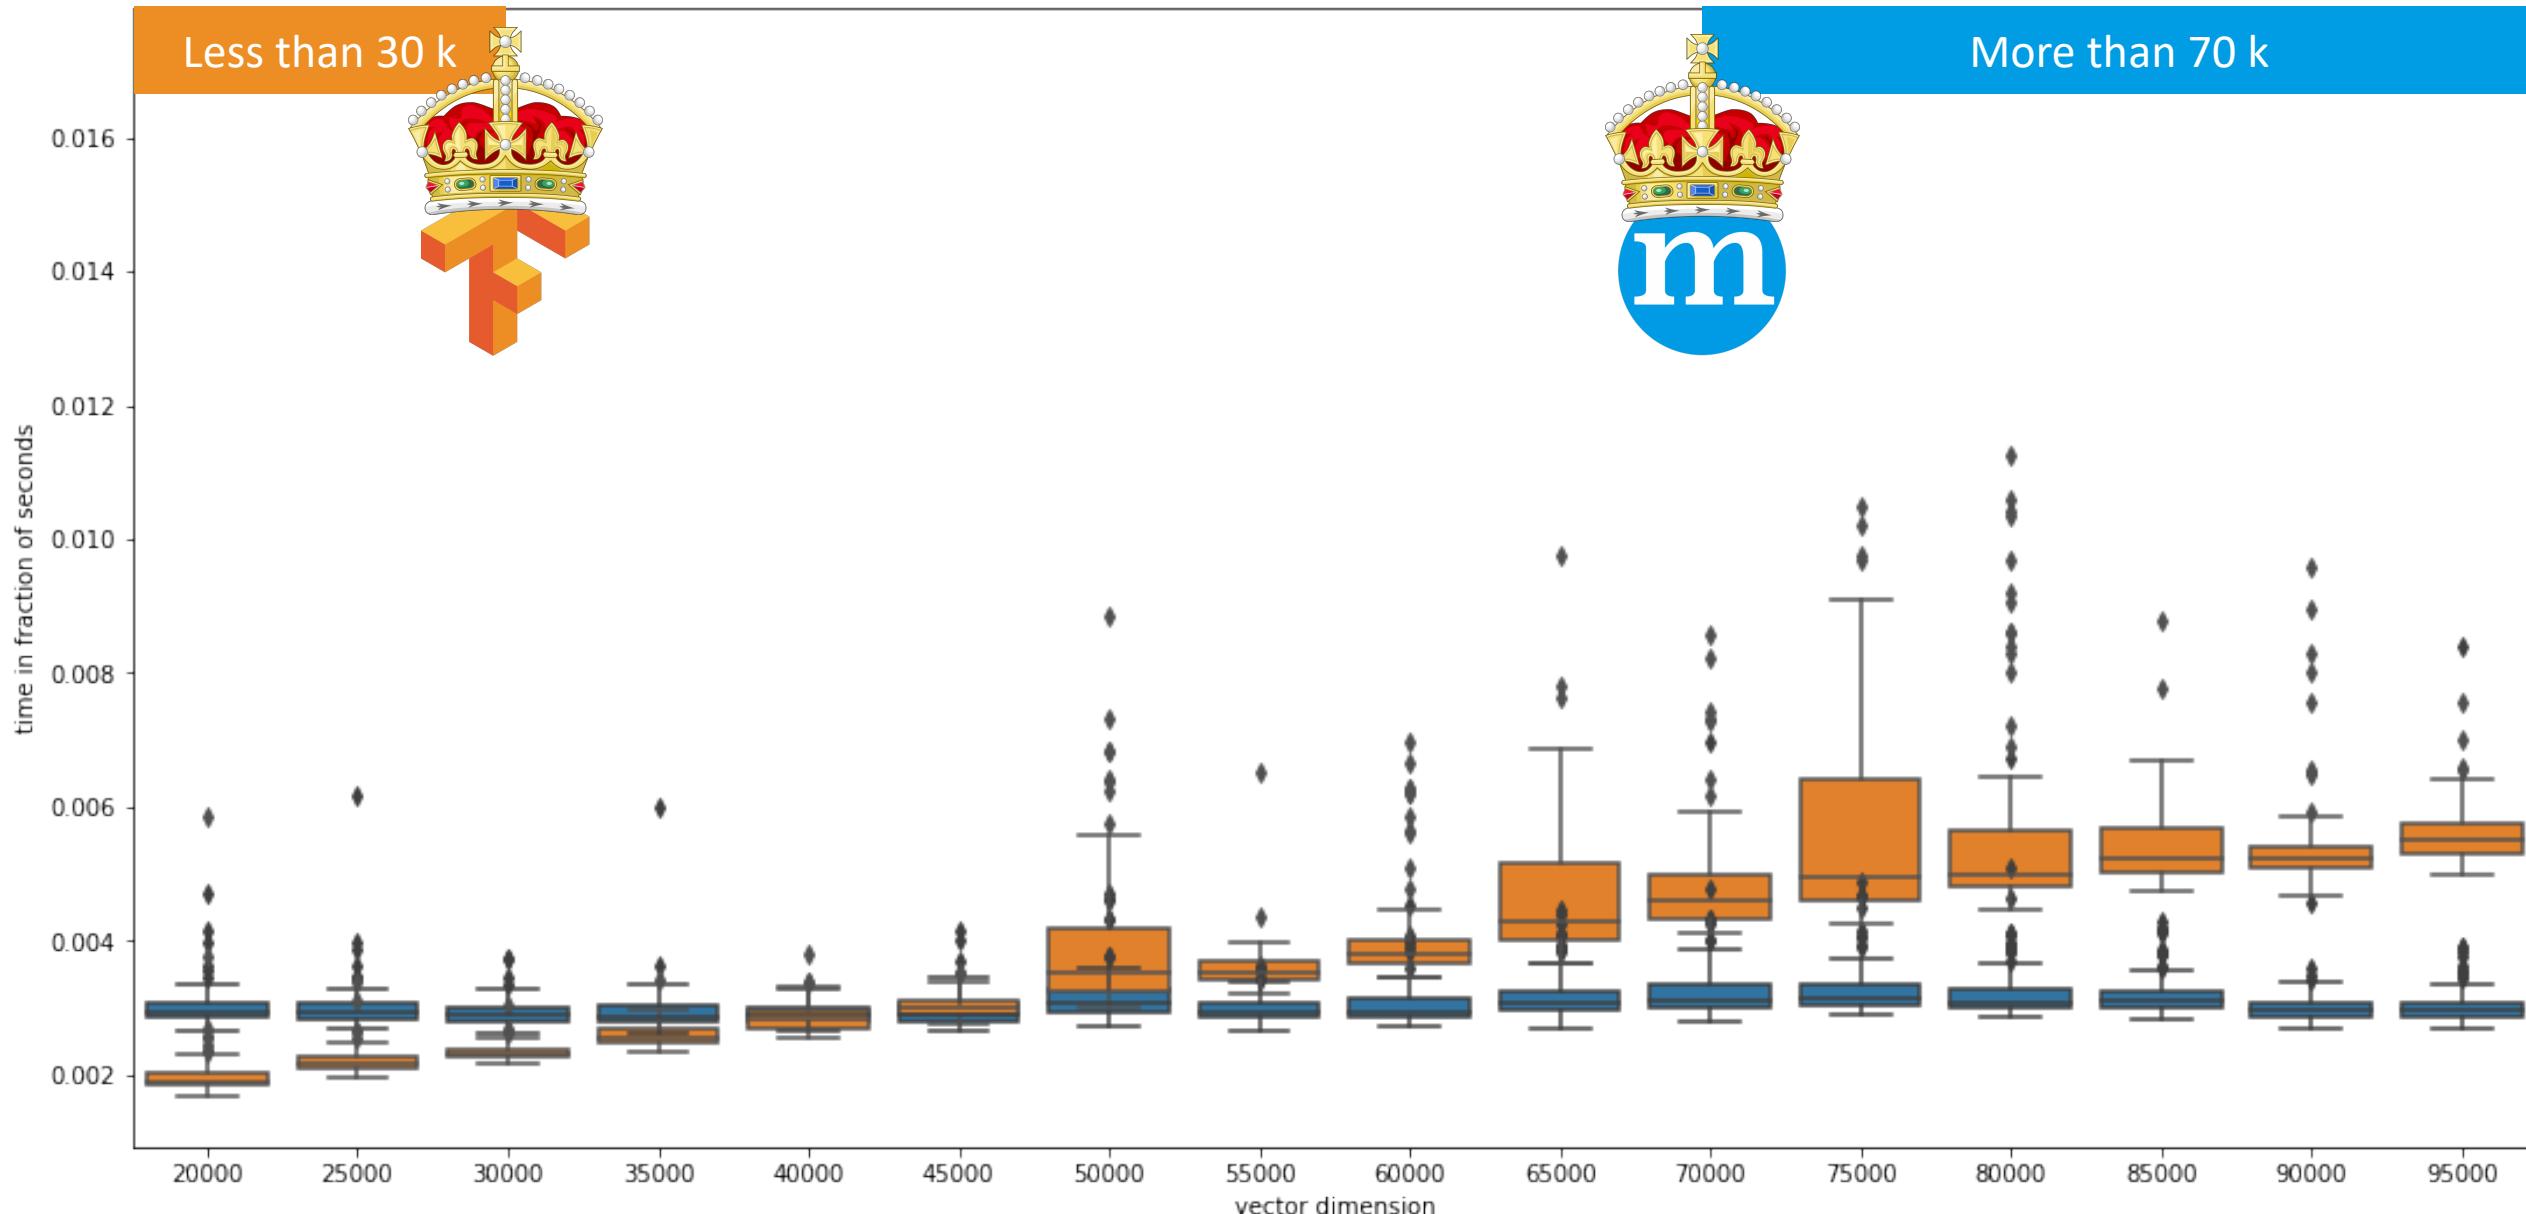


pythonTM list

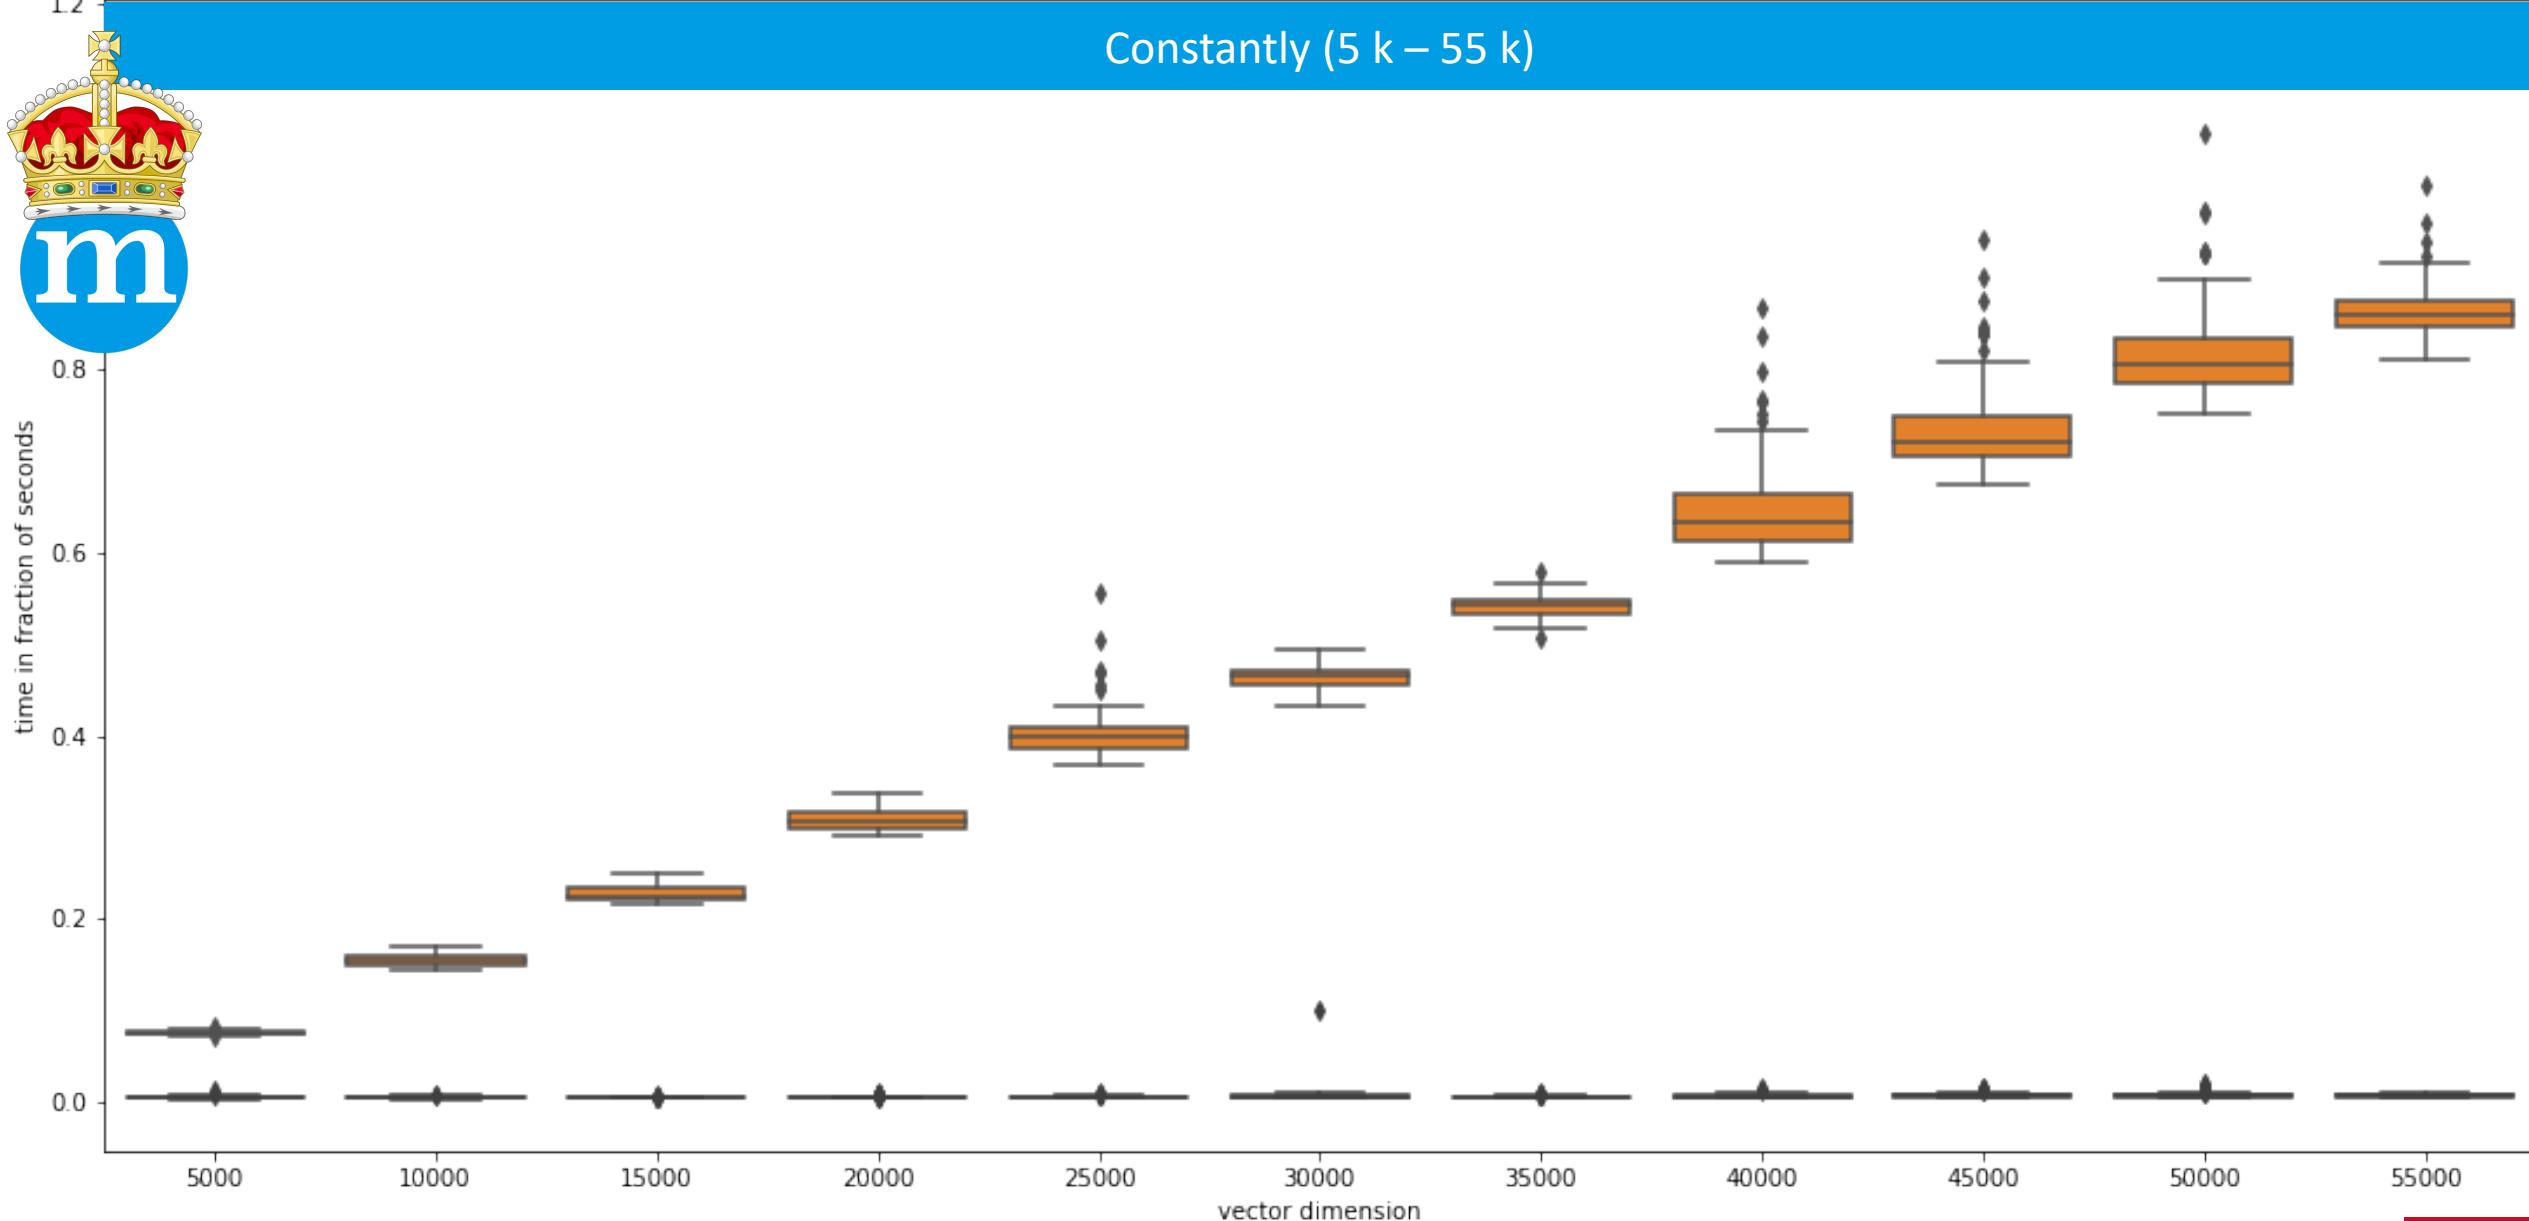
Experiment Description



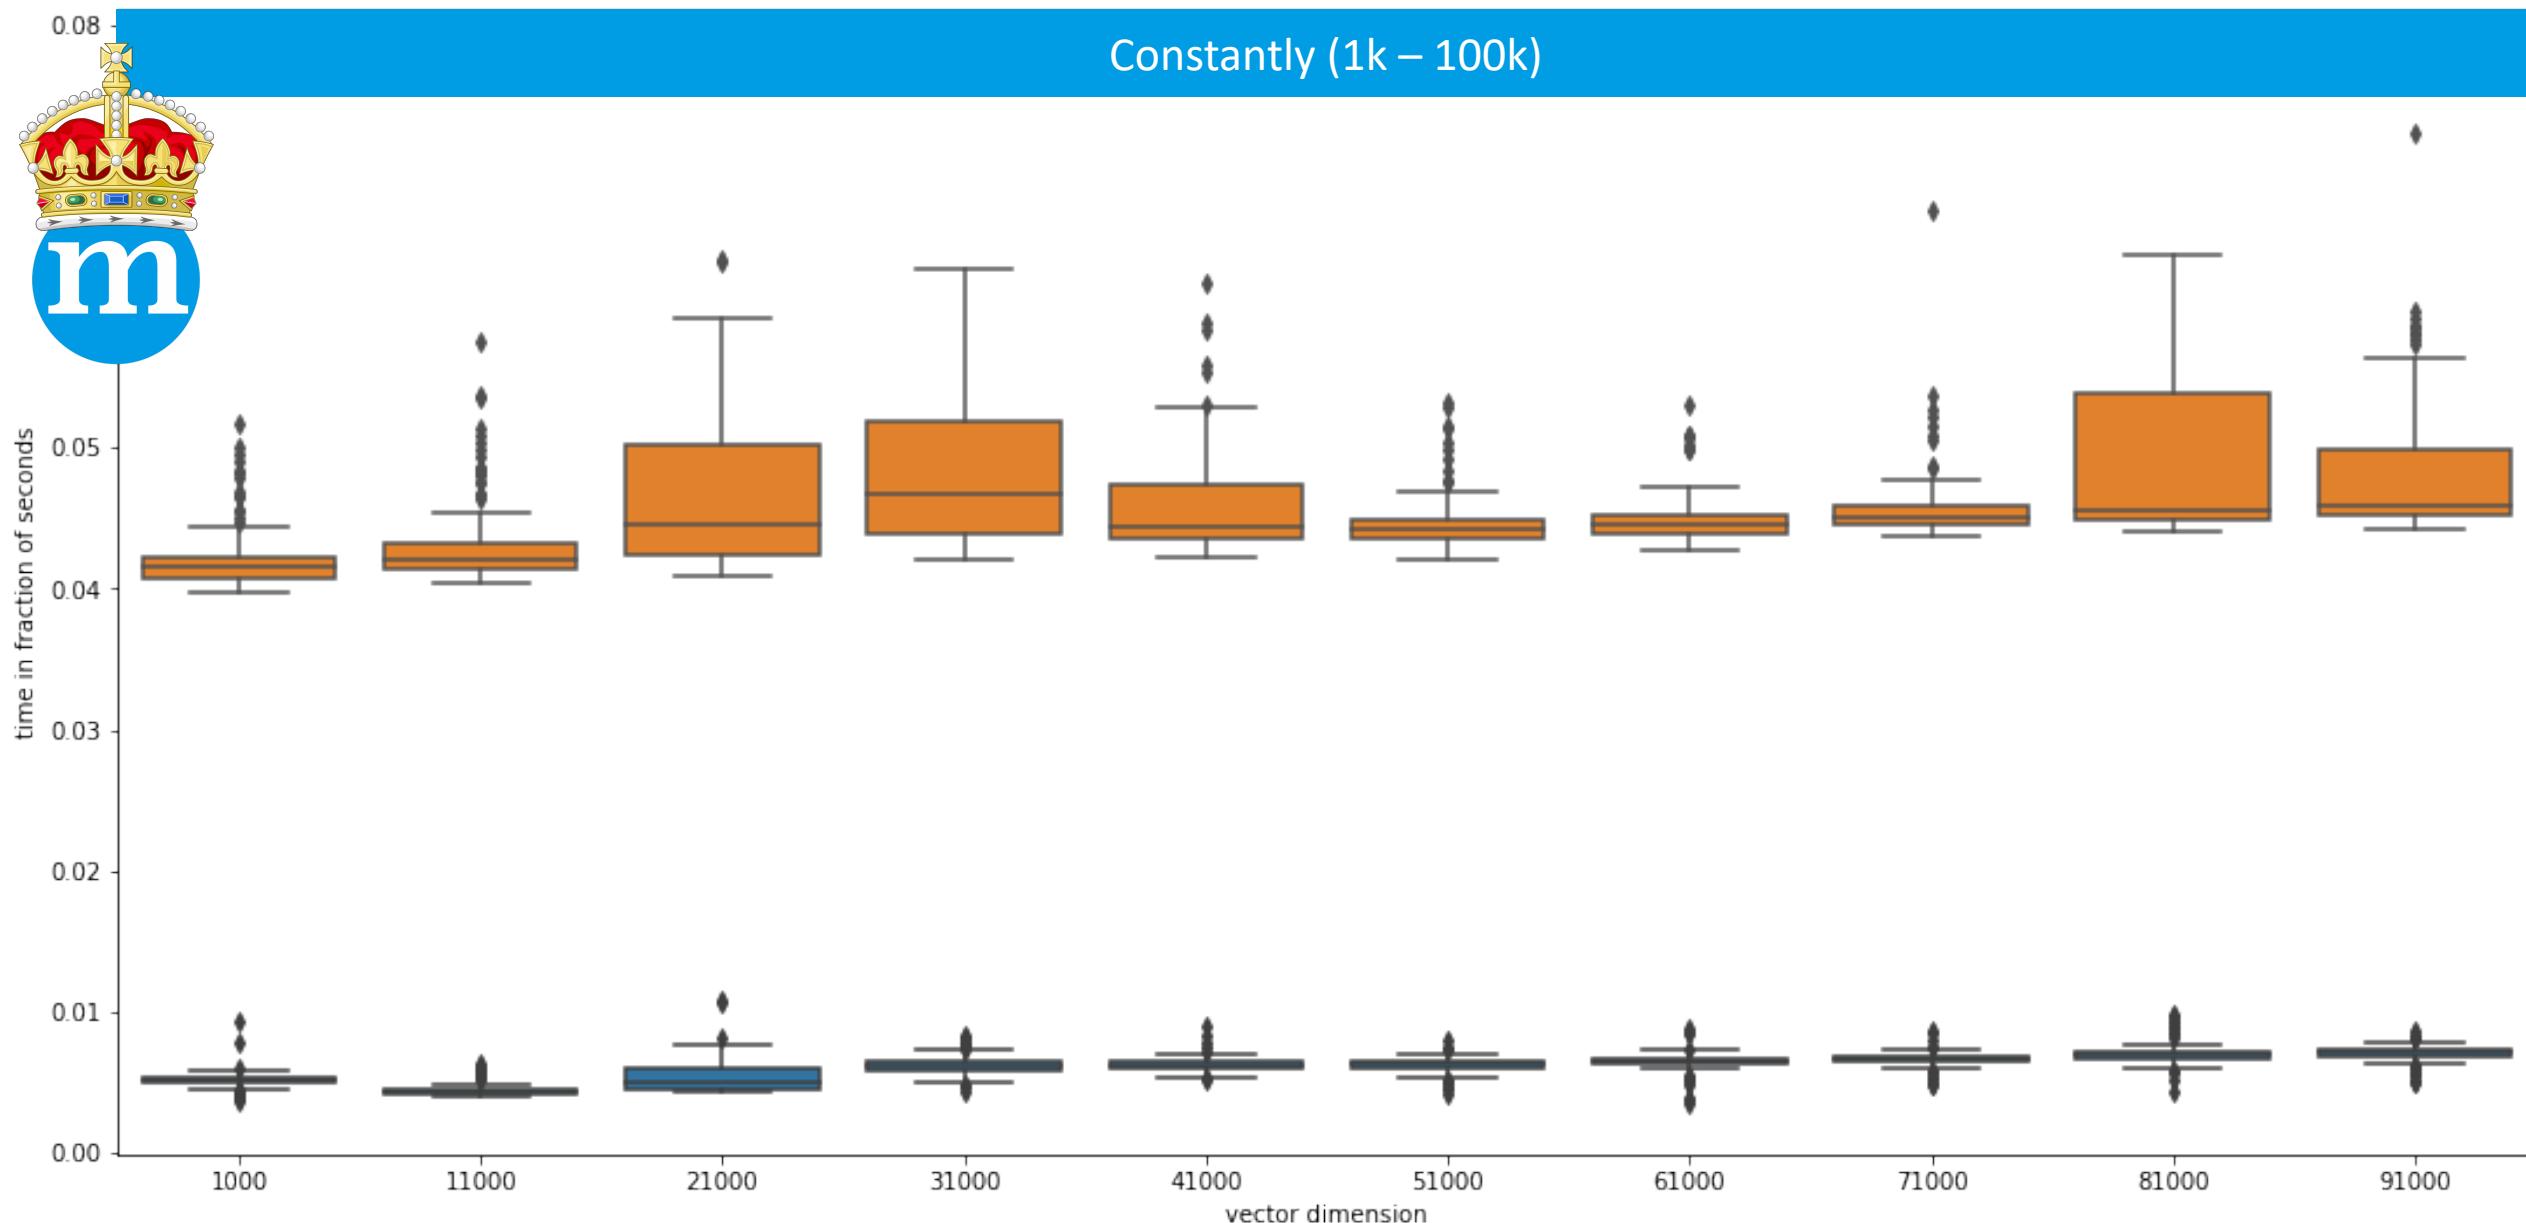
Task 1.1 - Initialization of 50 vectors - from Numpy input



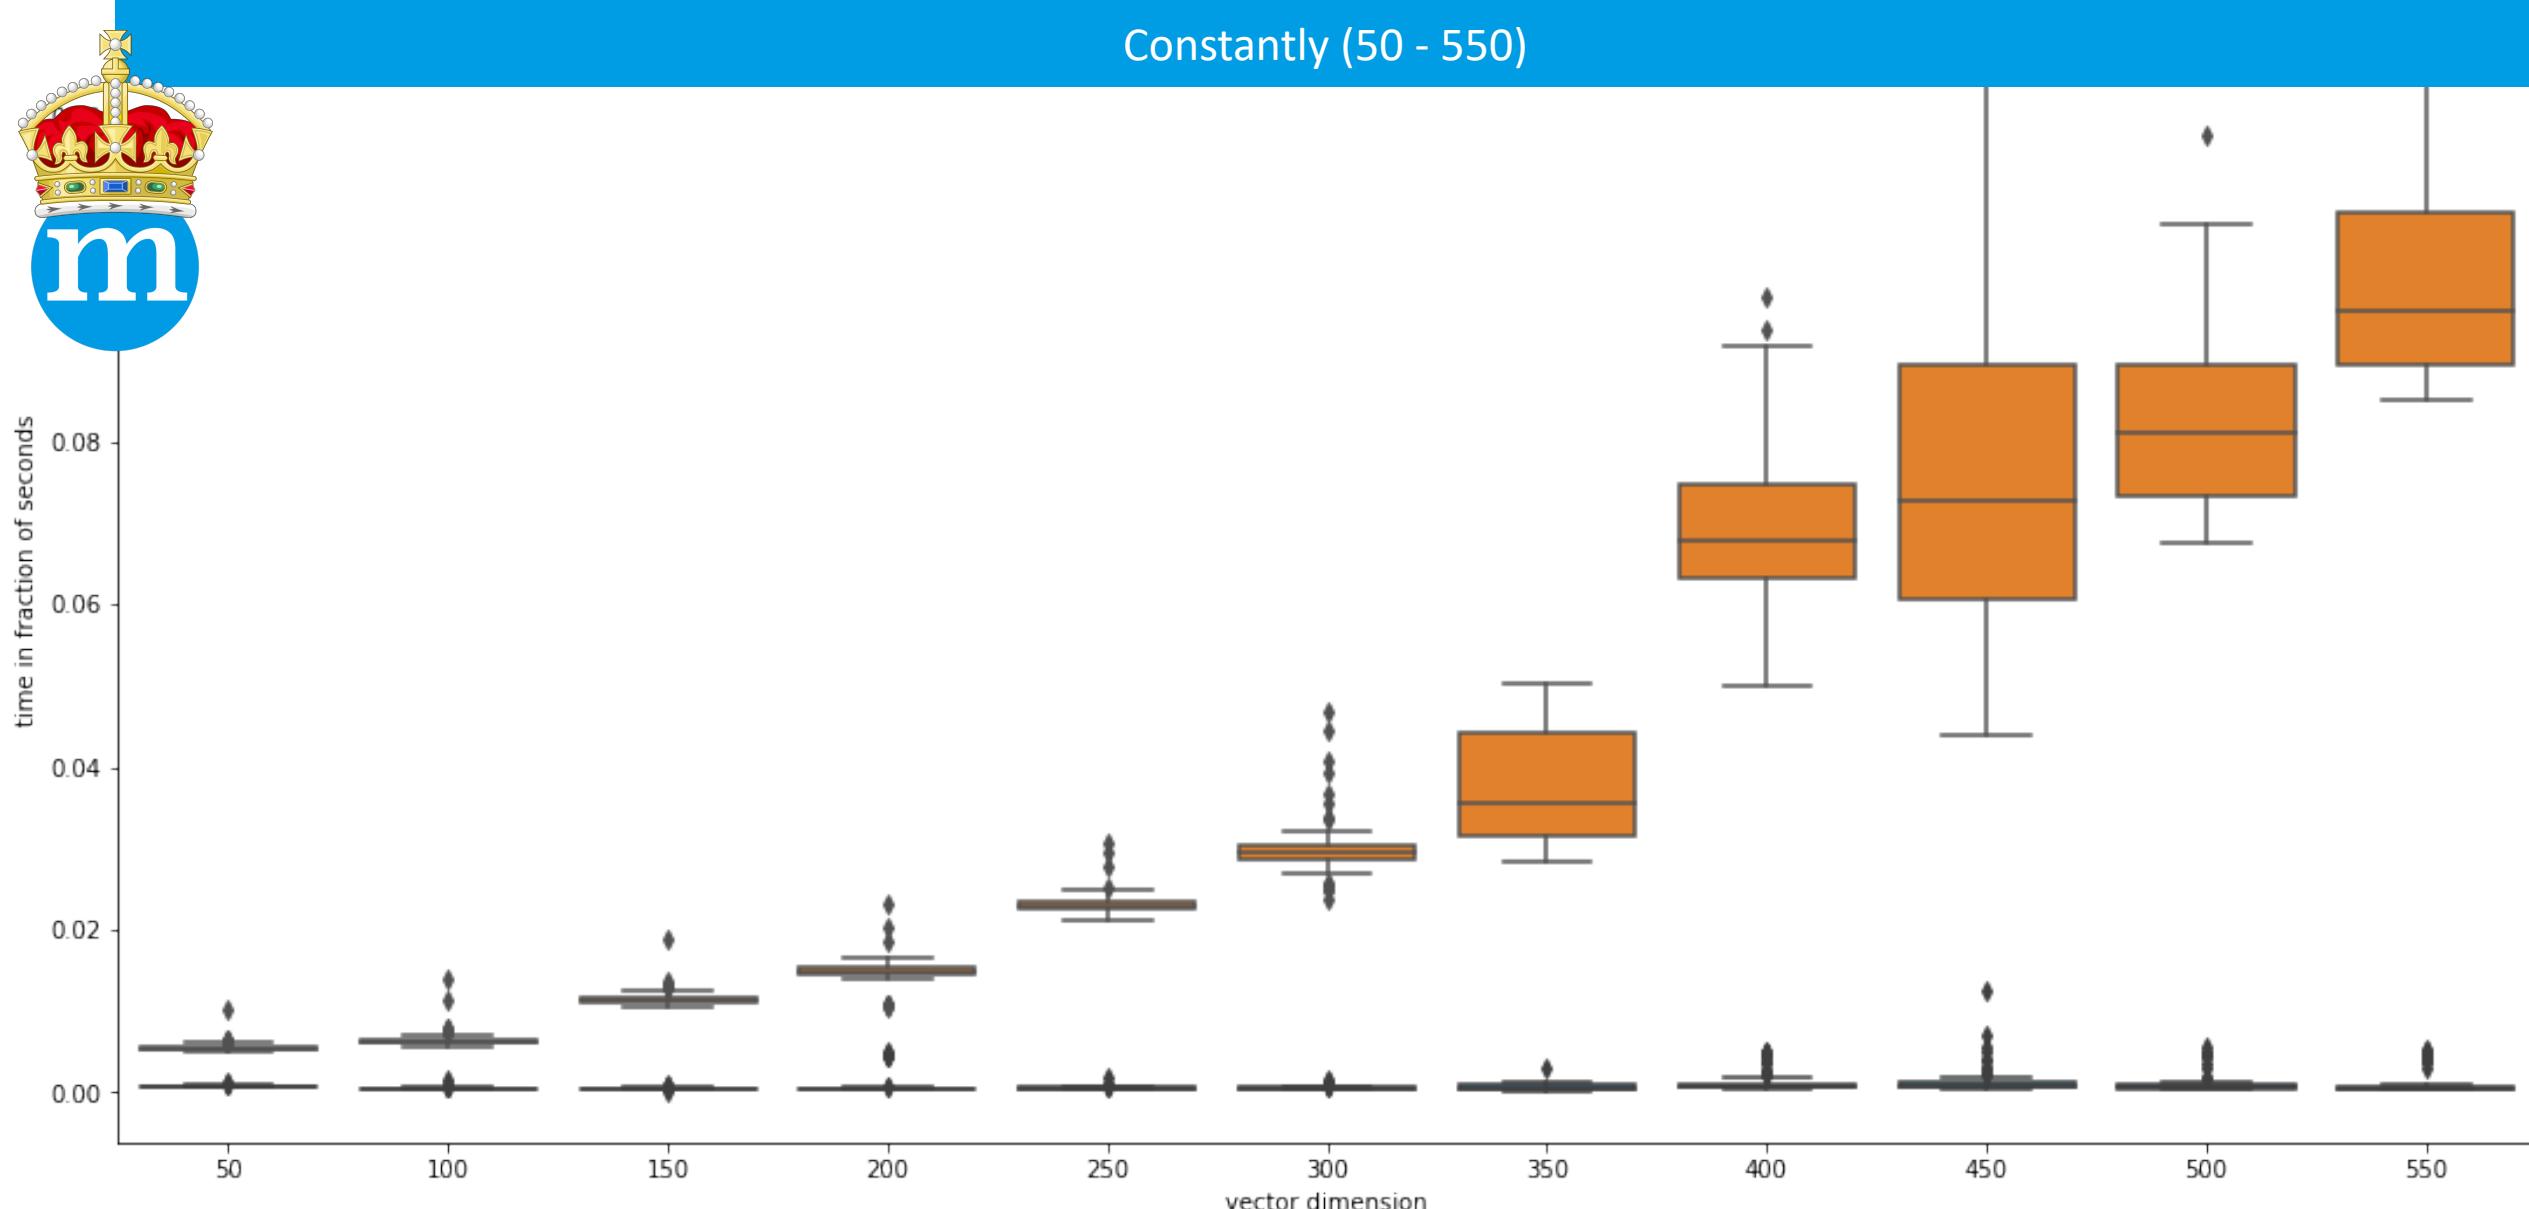
Task 1.2 - Initialization of 50 vectors - from list input



Task 2 - Dot Product of 50 vectors



Task 3 – 50 Matrix multiplication –between square matrices



Tasks 4 – Activation Functions

Task 4:
Execute the
activation
function on 50
vectors

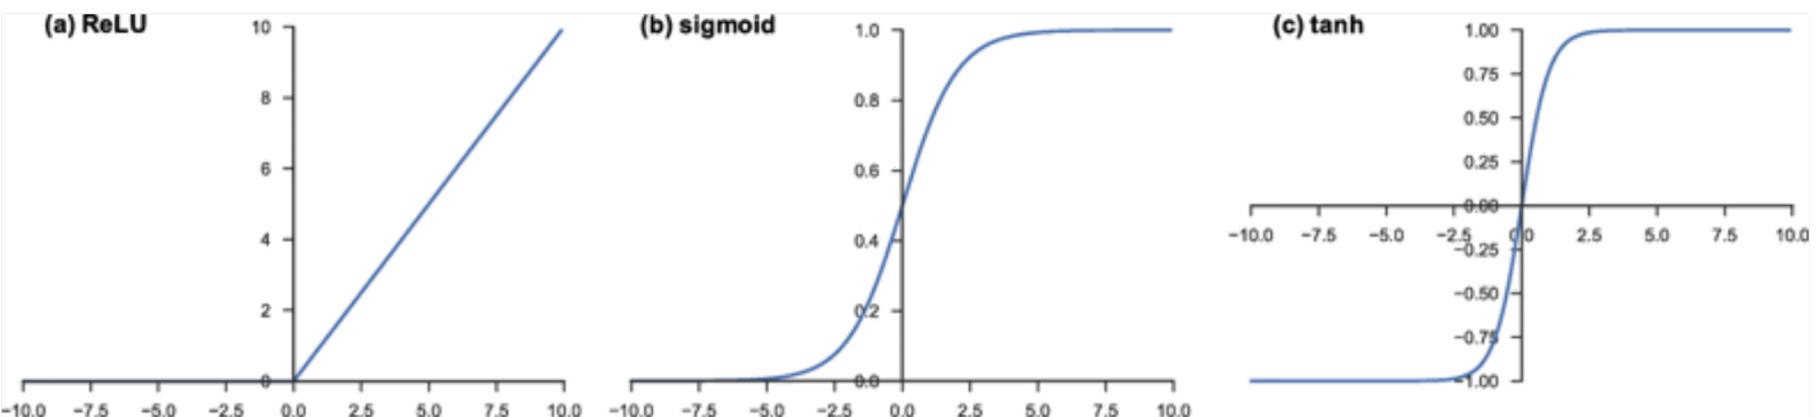
Task 4.1:
Sigmoid

Task 4.2:
ReLU

Task 4.3:
tanh

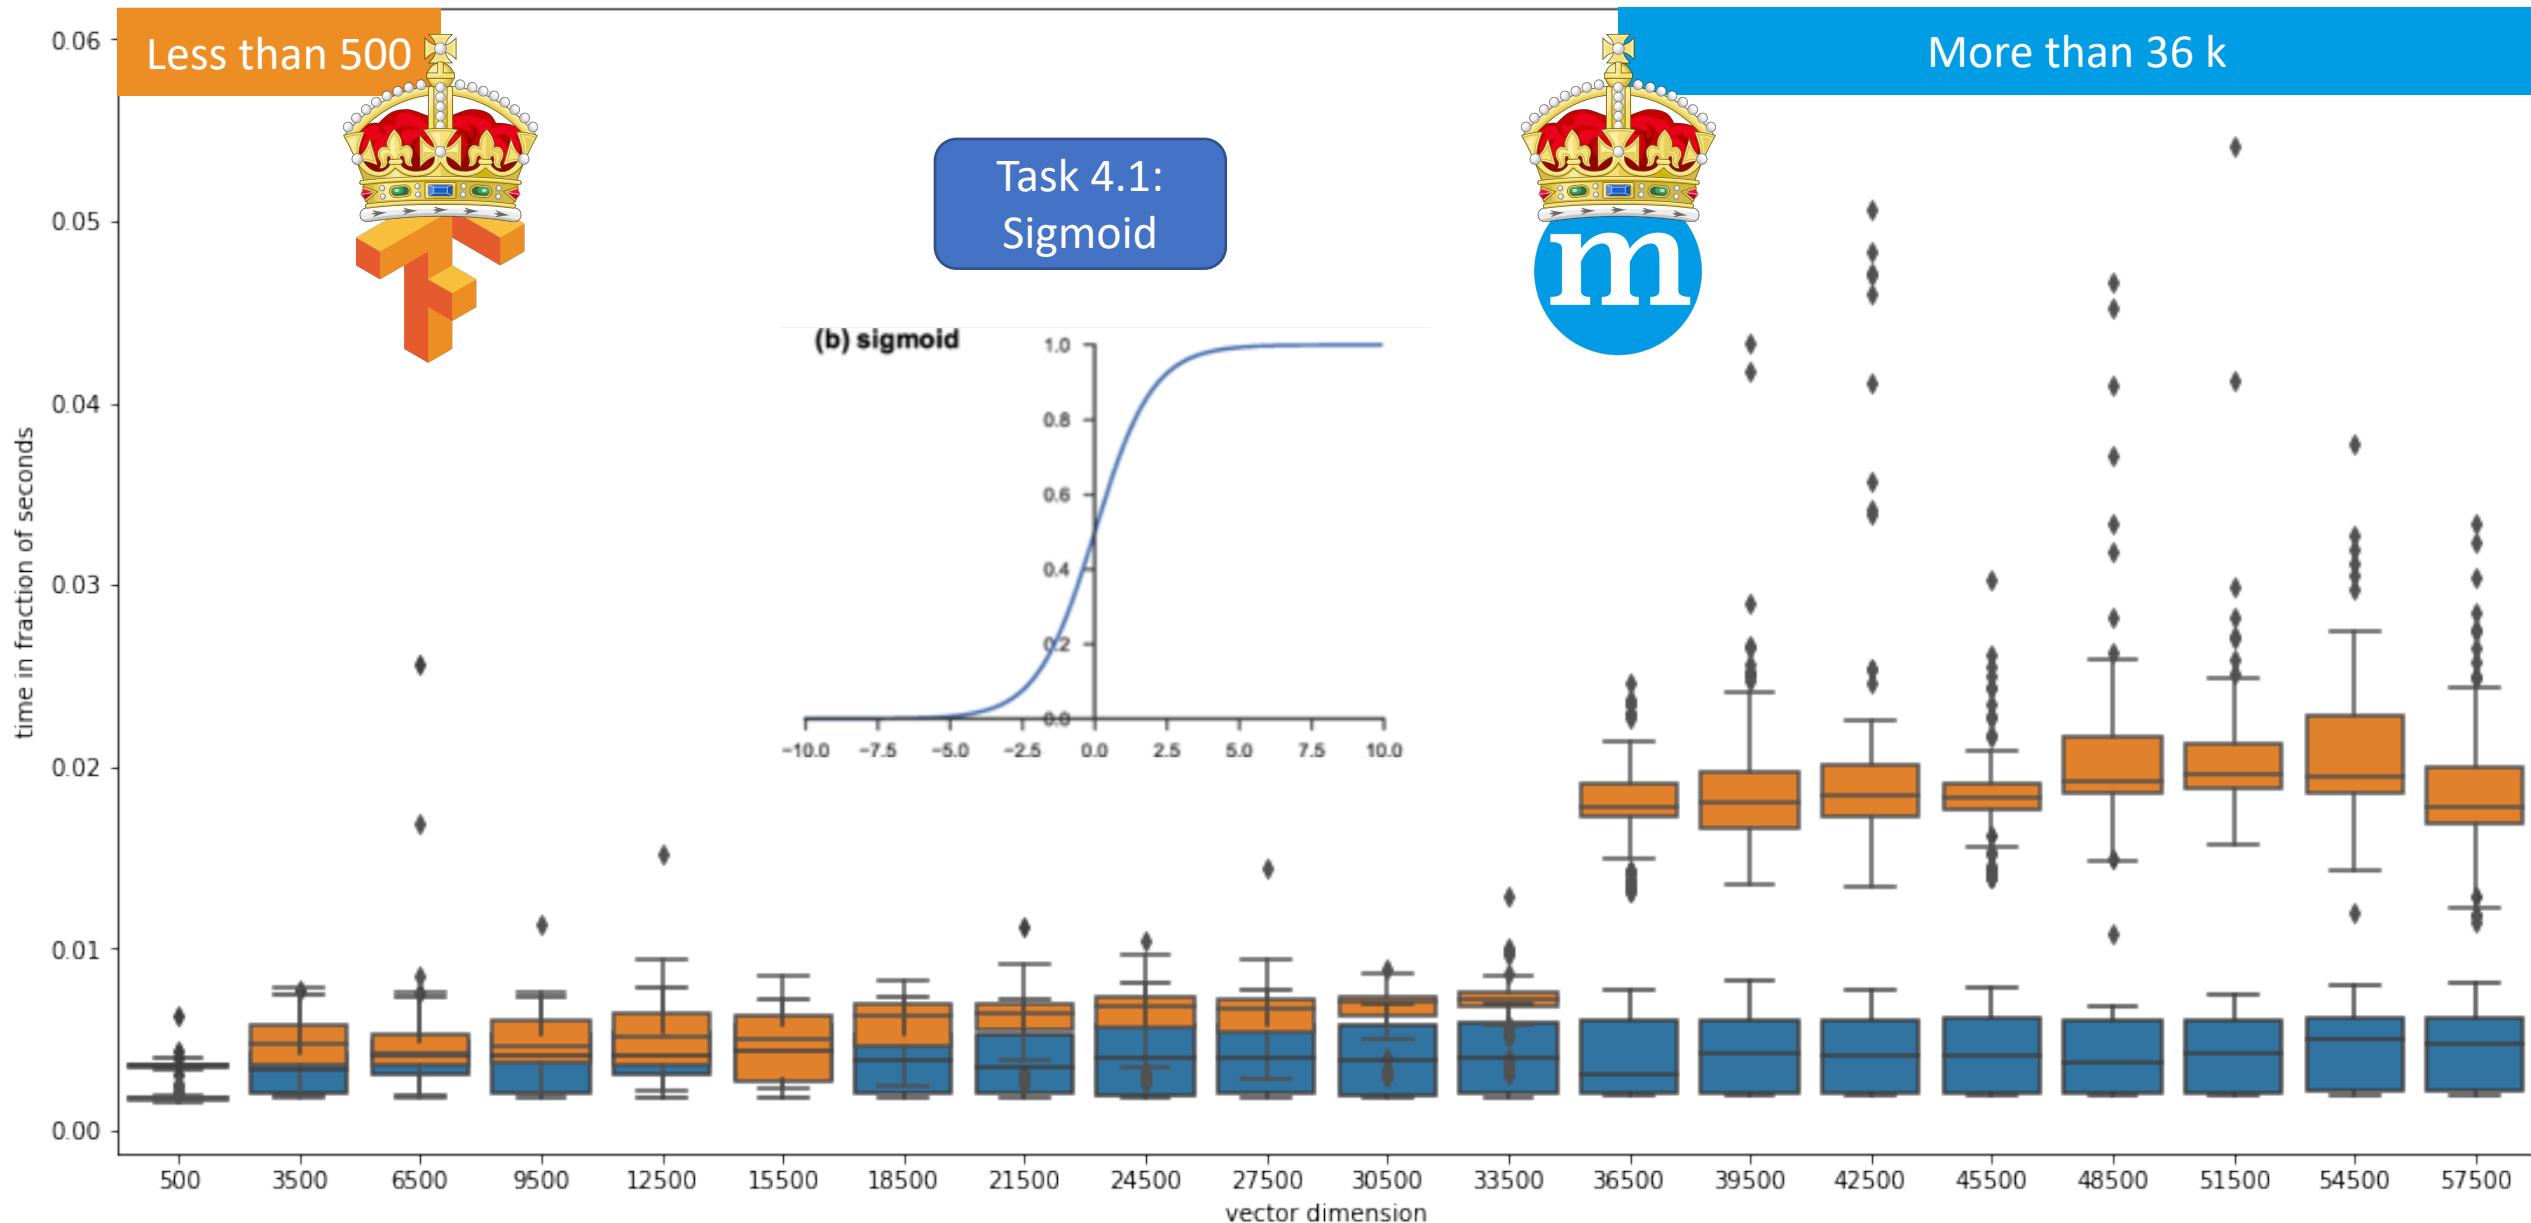
Task 4.4:
softReLU

Task 4.5:
softsign

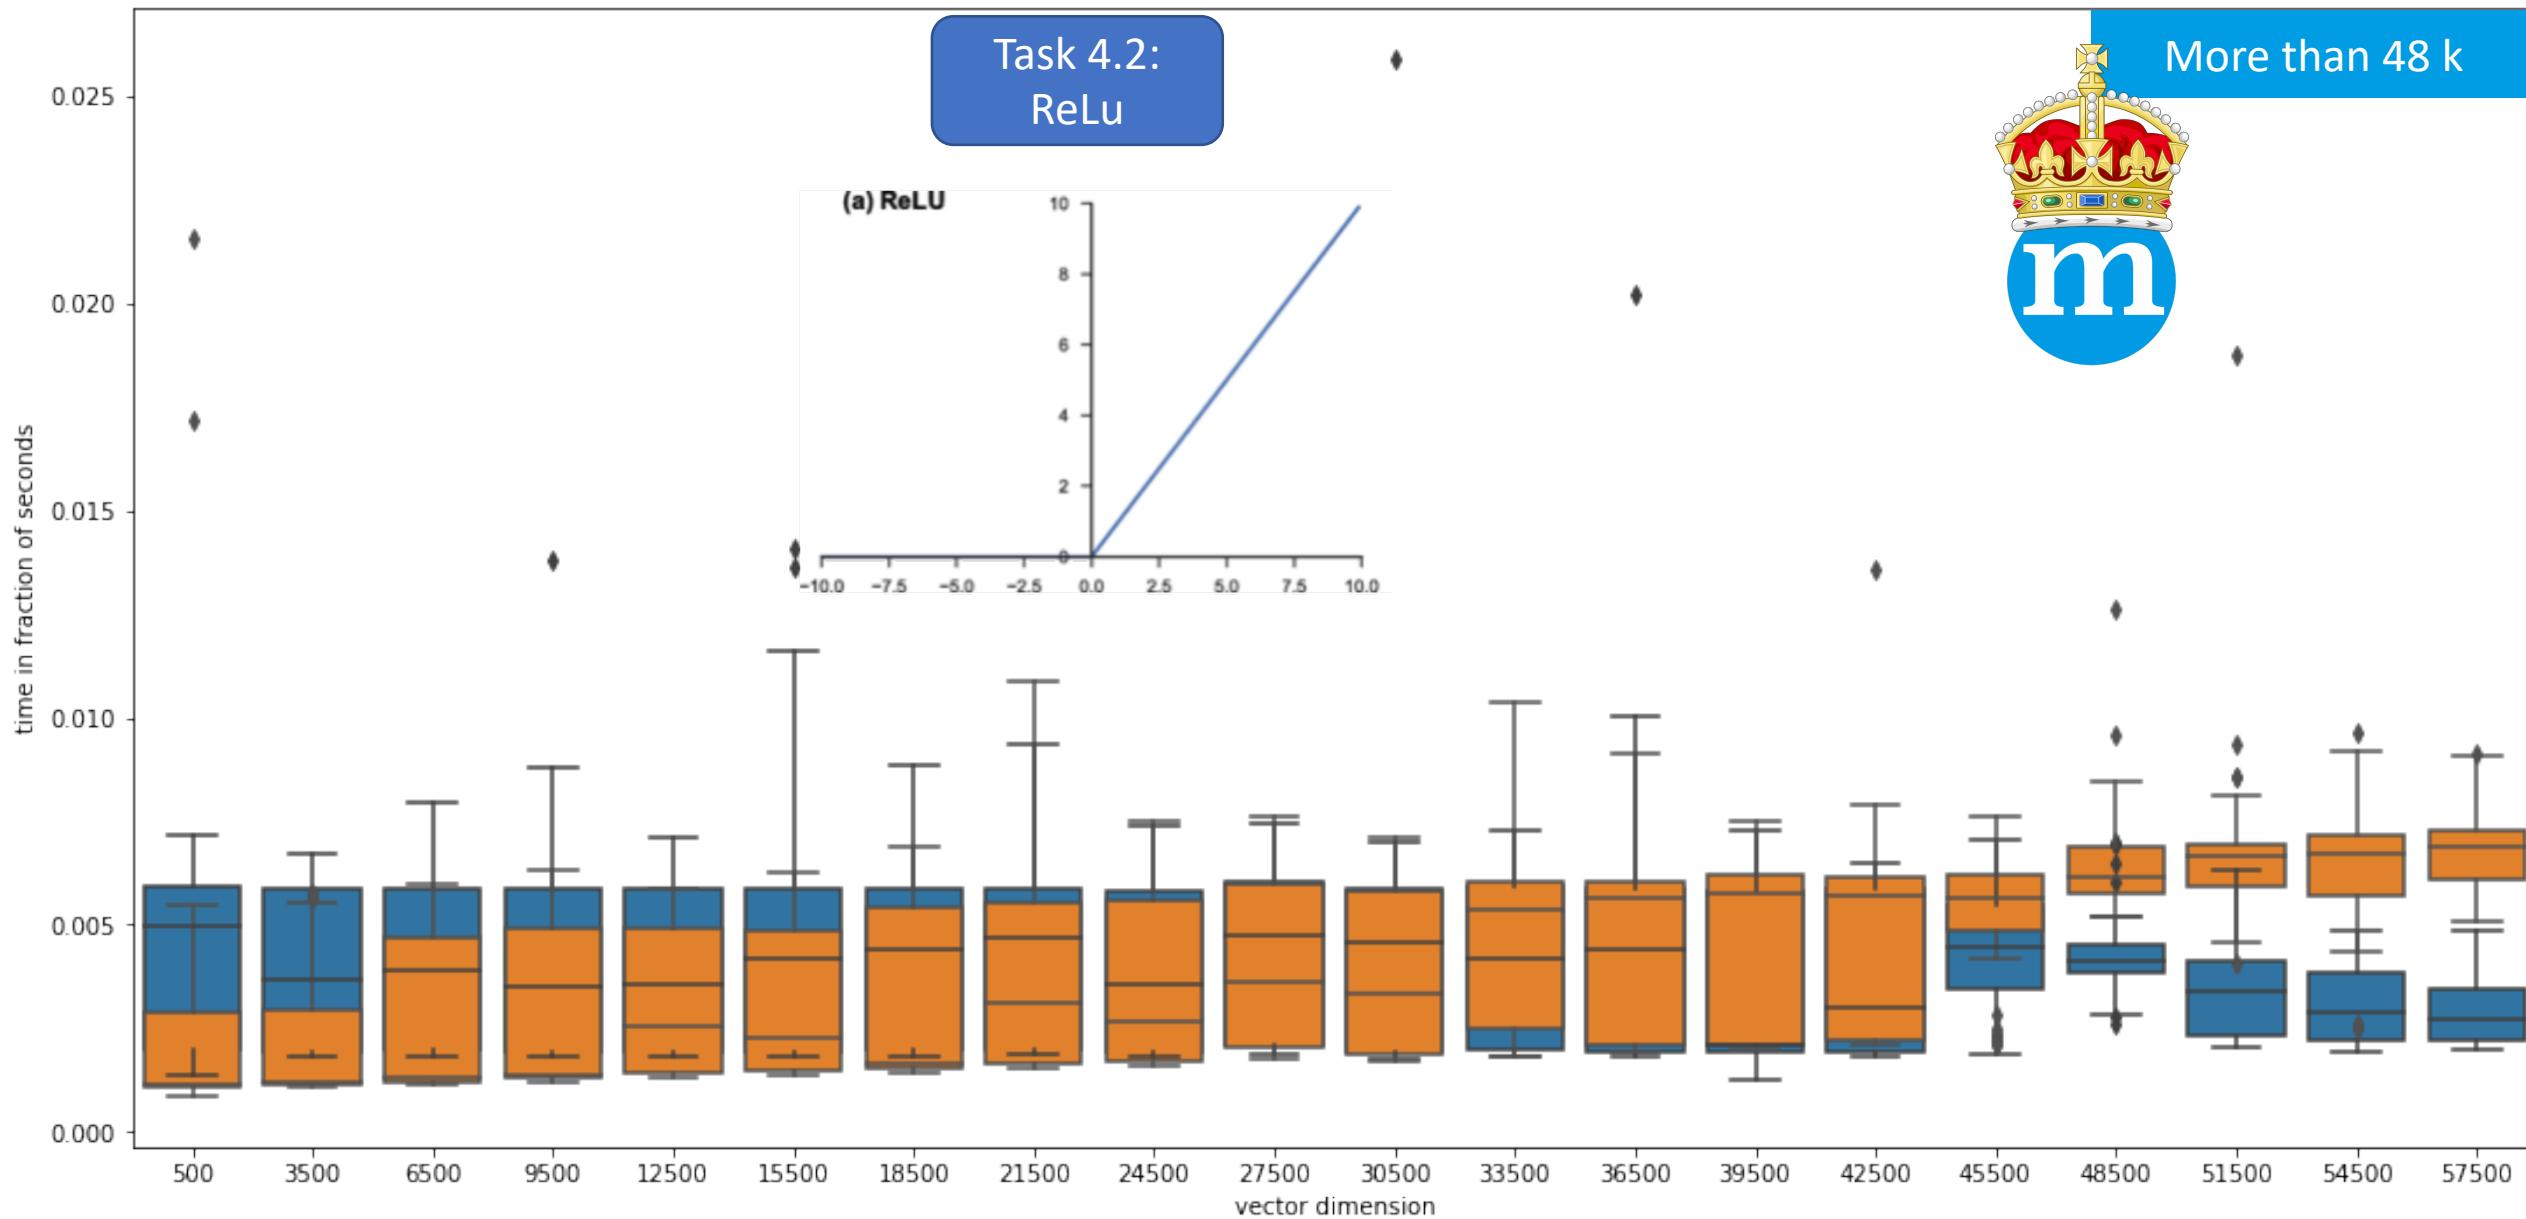


Activation functions commonly applied to neural networks: **a** rectified linear unit (ReLU), **b** sigmoid, and **c** hyperbolic tangent (tanh)

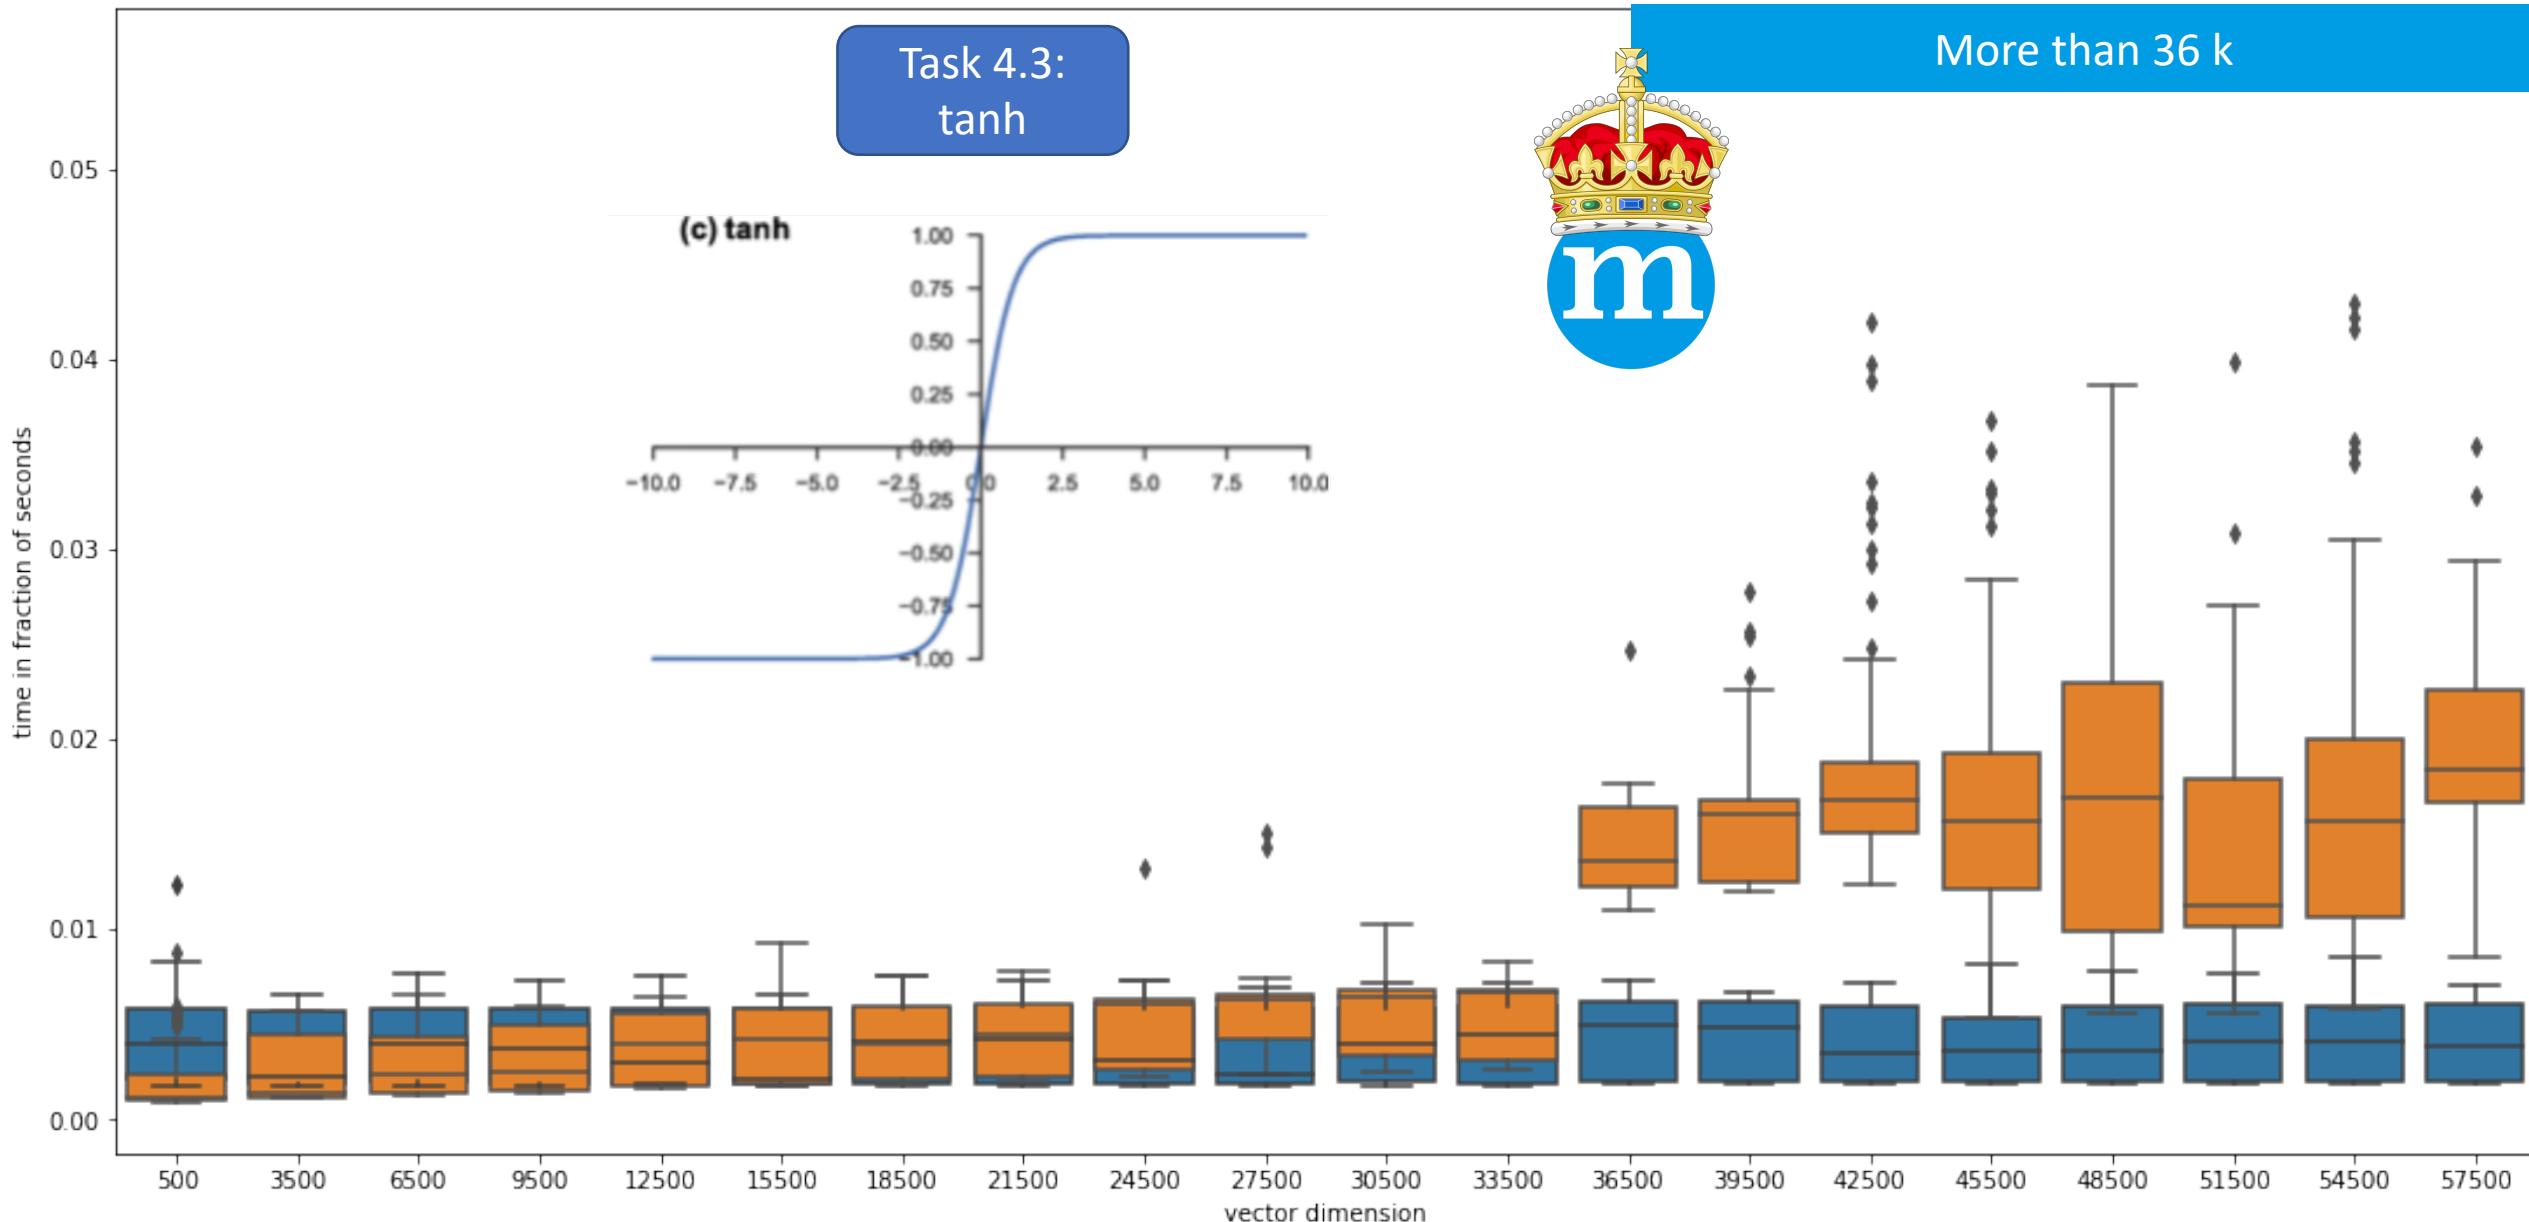
Task 4.1 - Activation on 50 vectors - Sigmoid



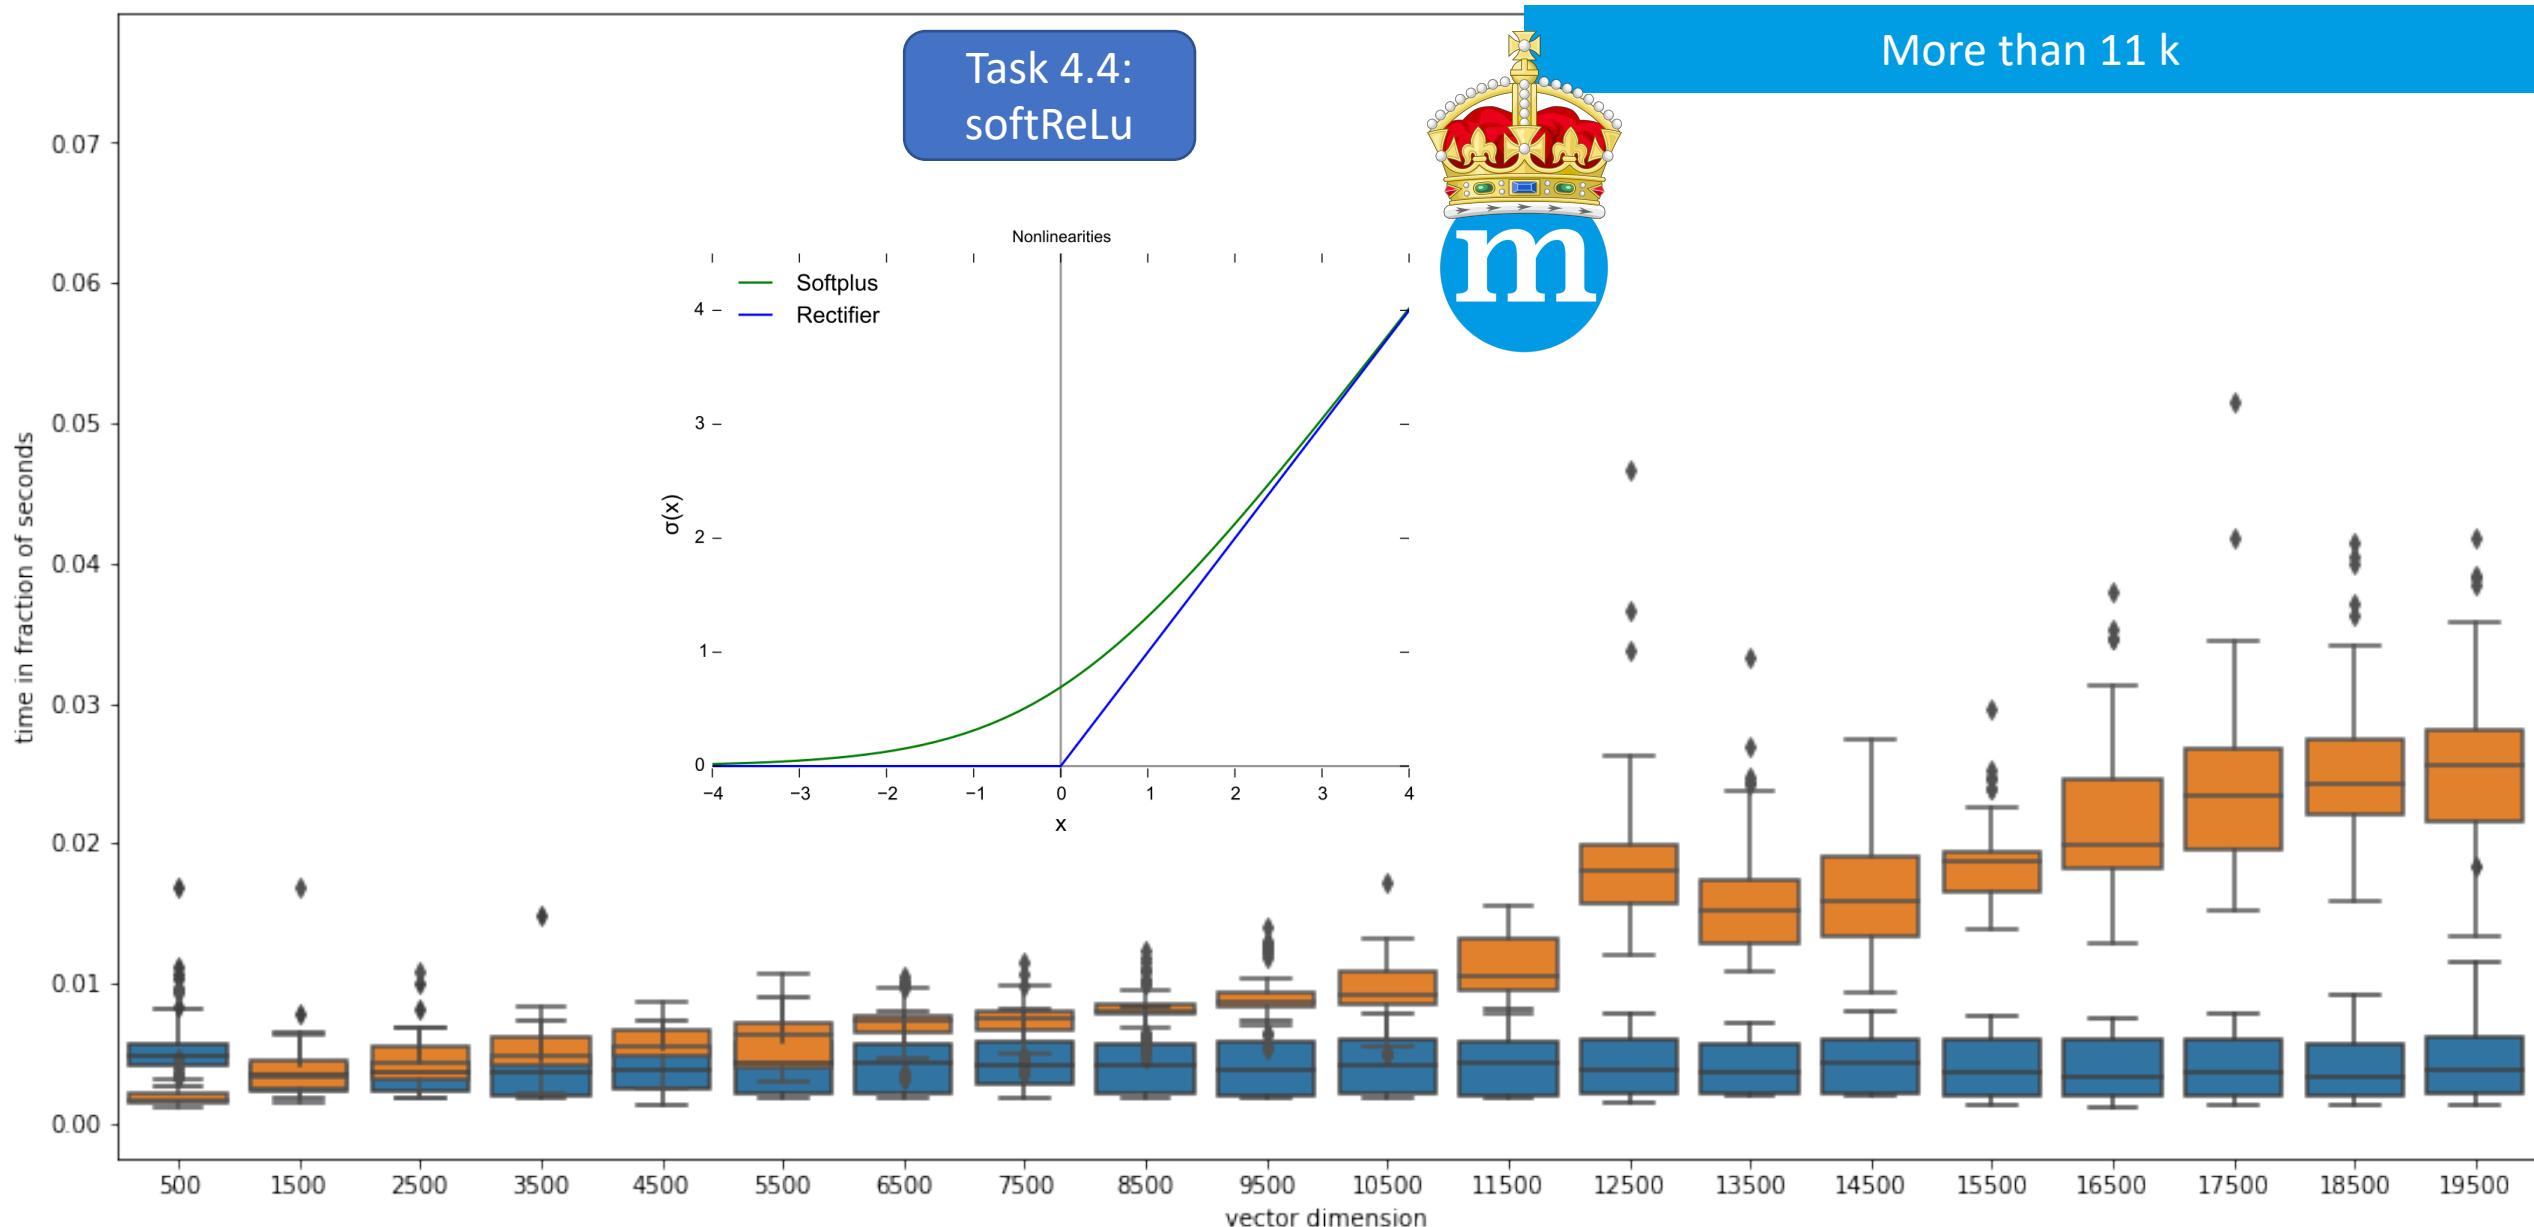
Task 4.2 - Activation on 50 vectors - Relu



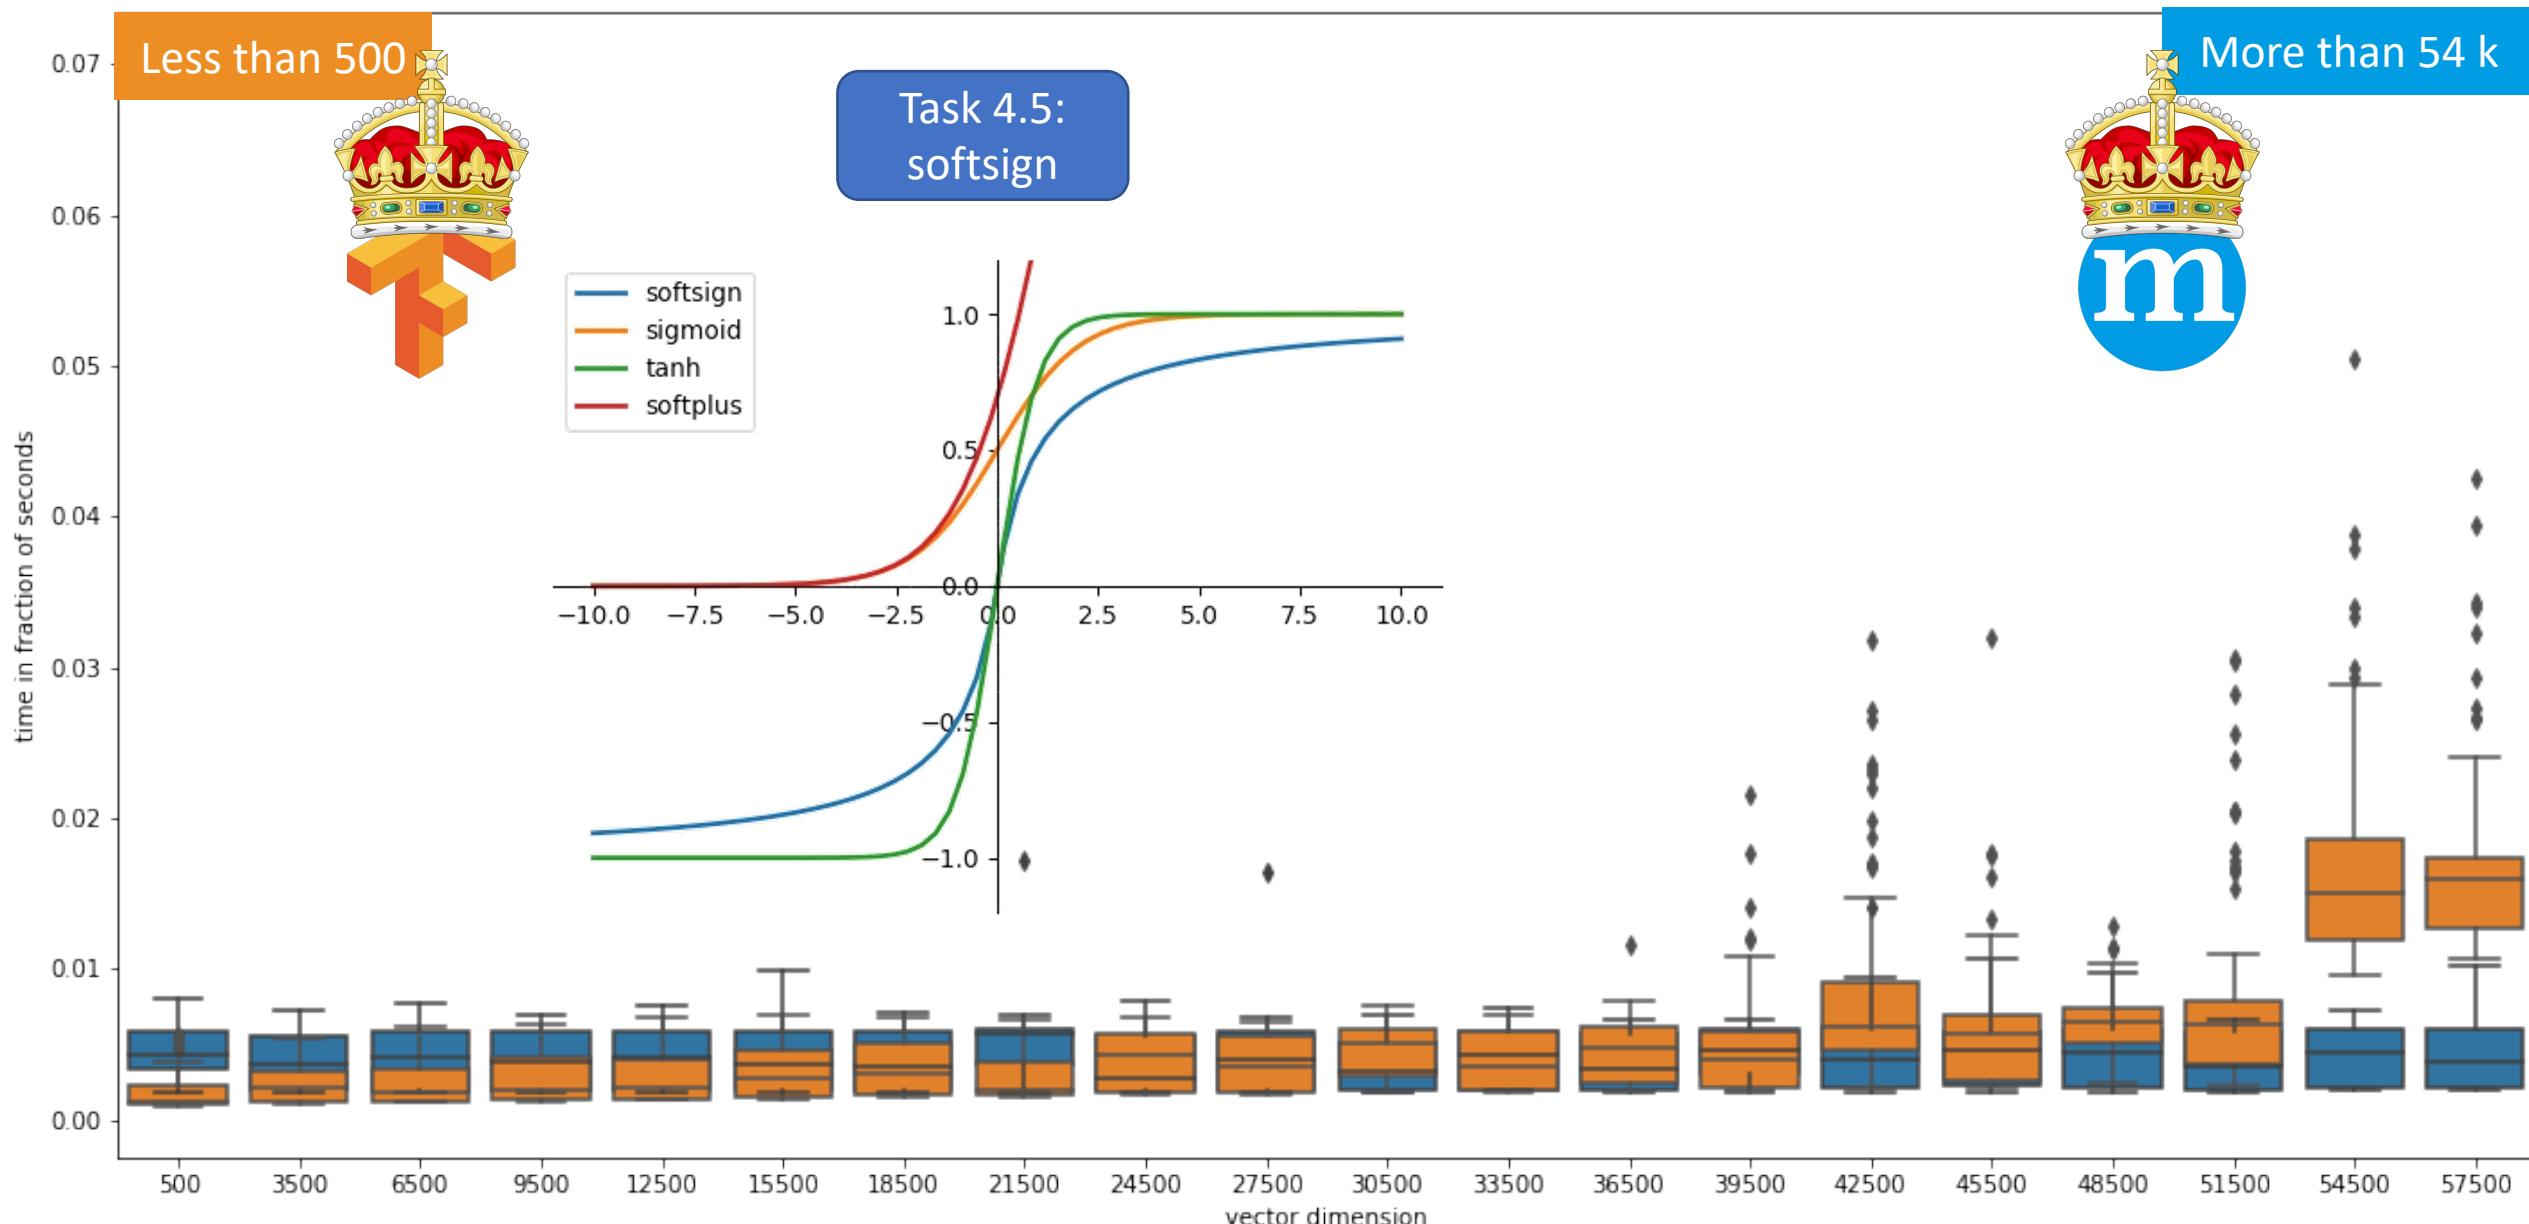
Task 4.3 - Activation on 50 vectors - TanH



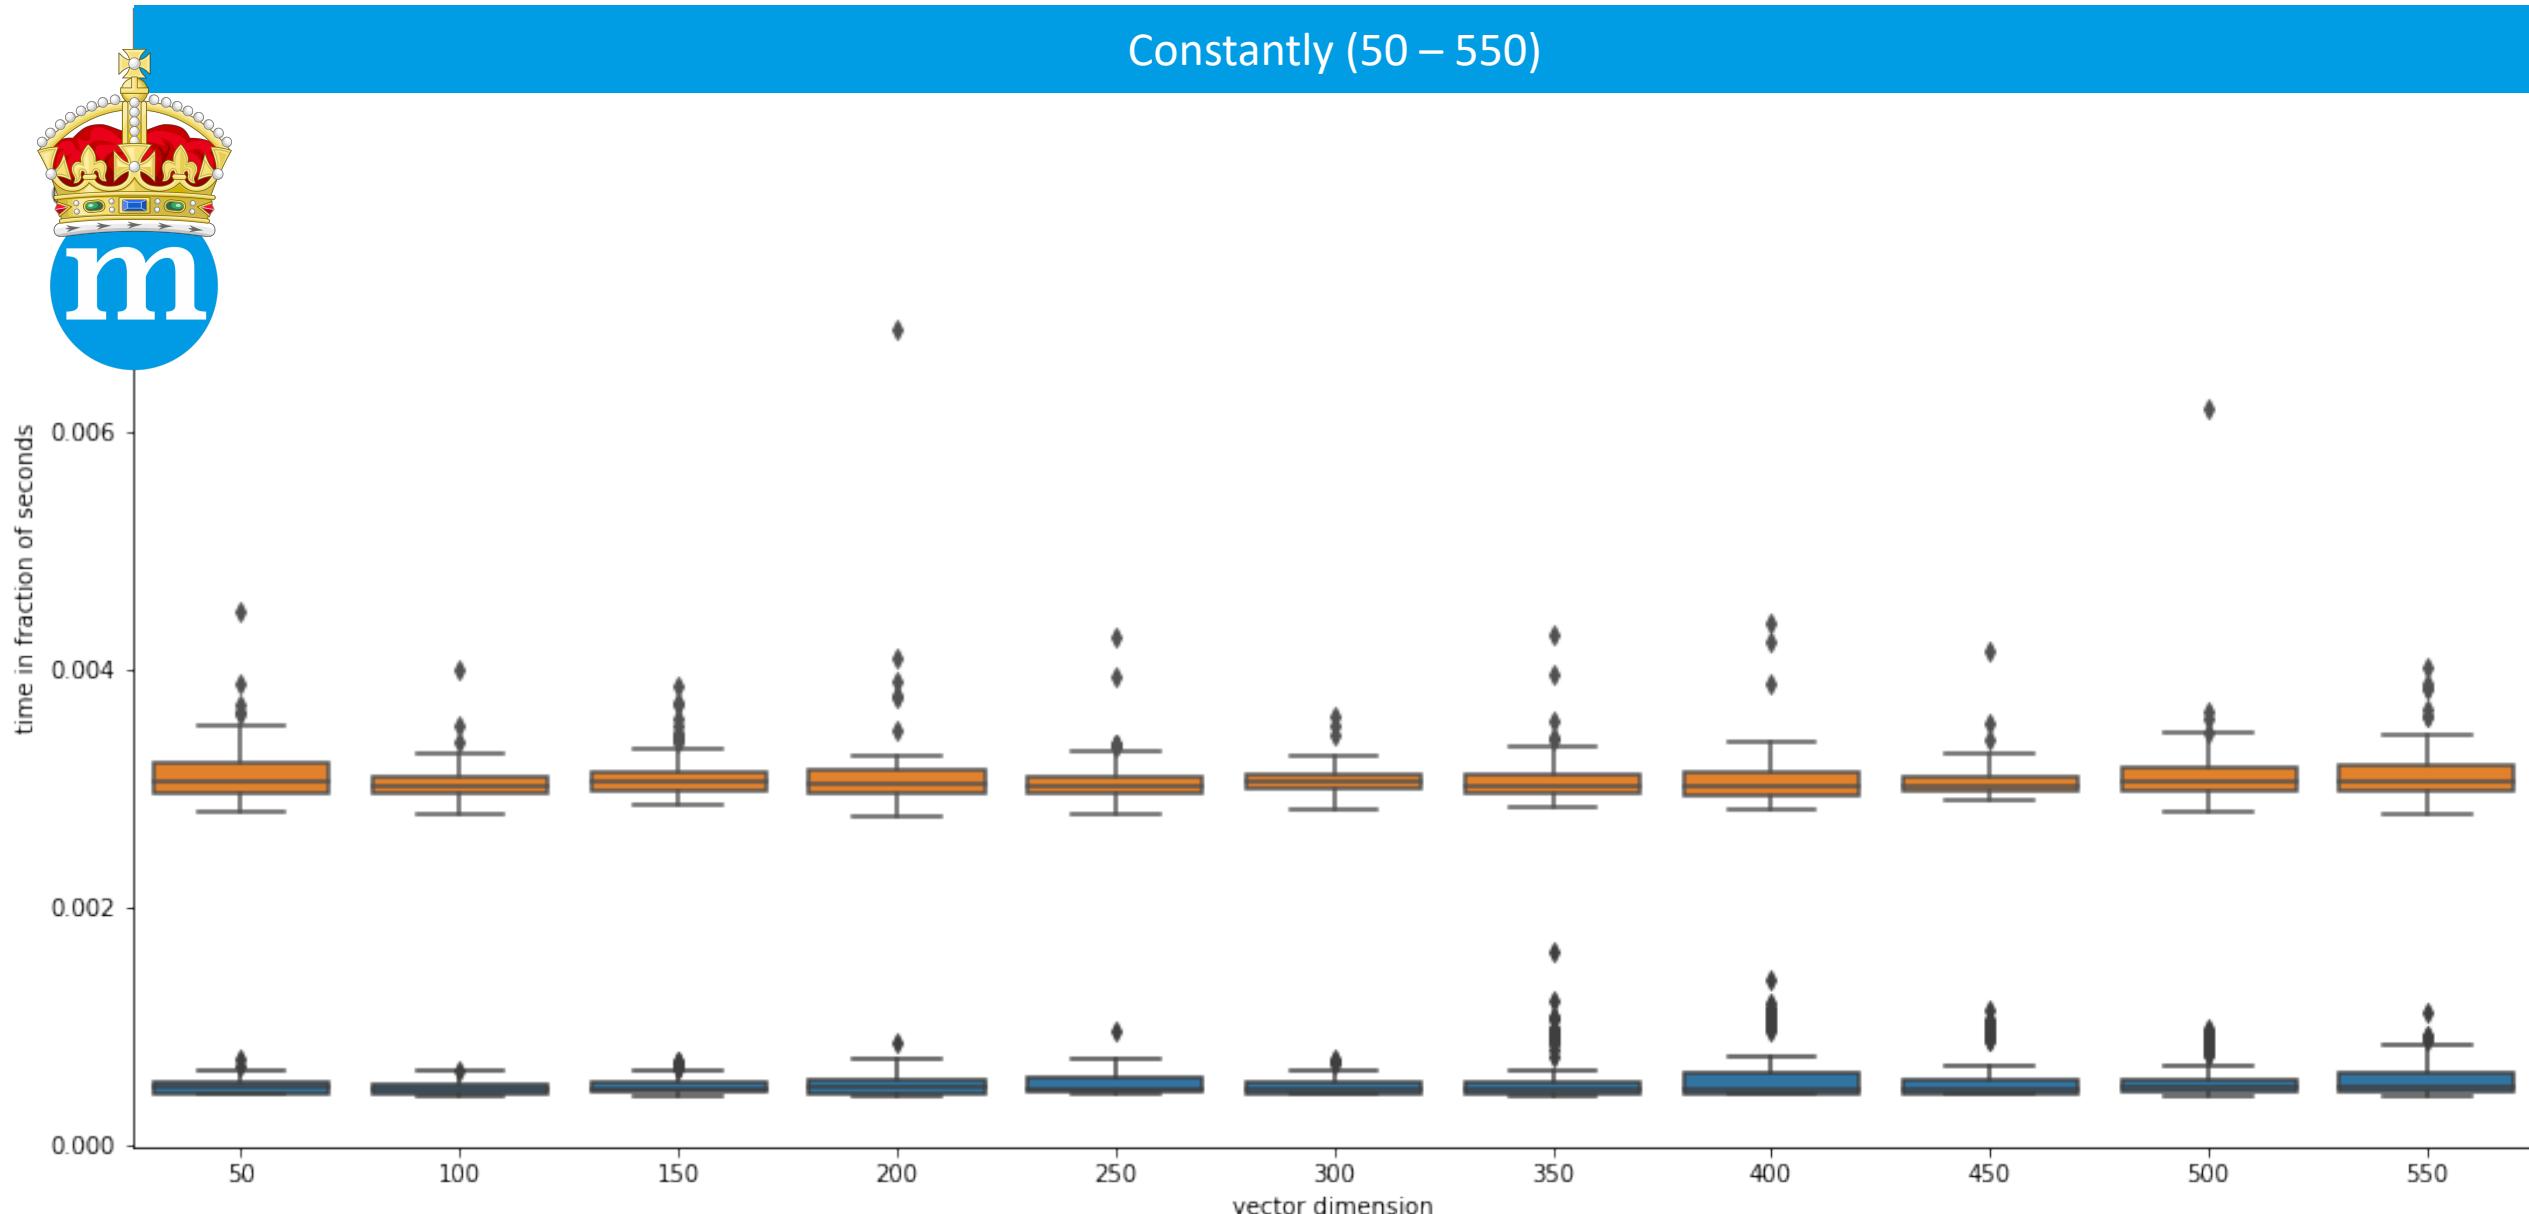
Task 4.4 - Activation on 50 vectors - SoftRelu



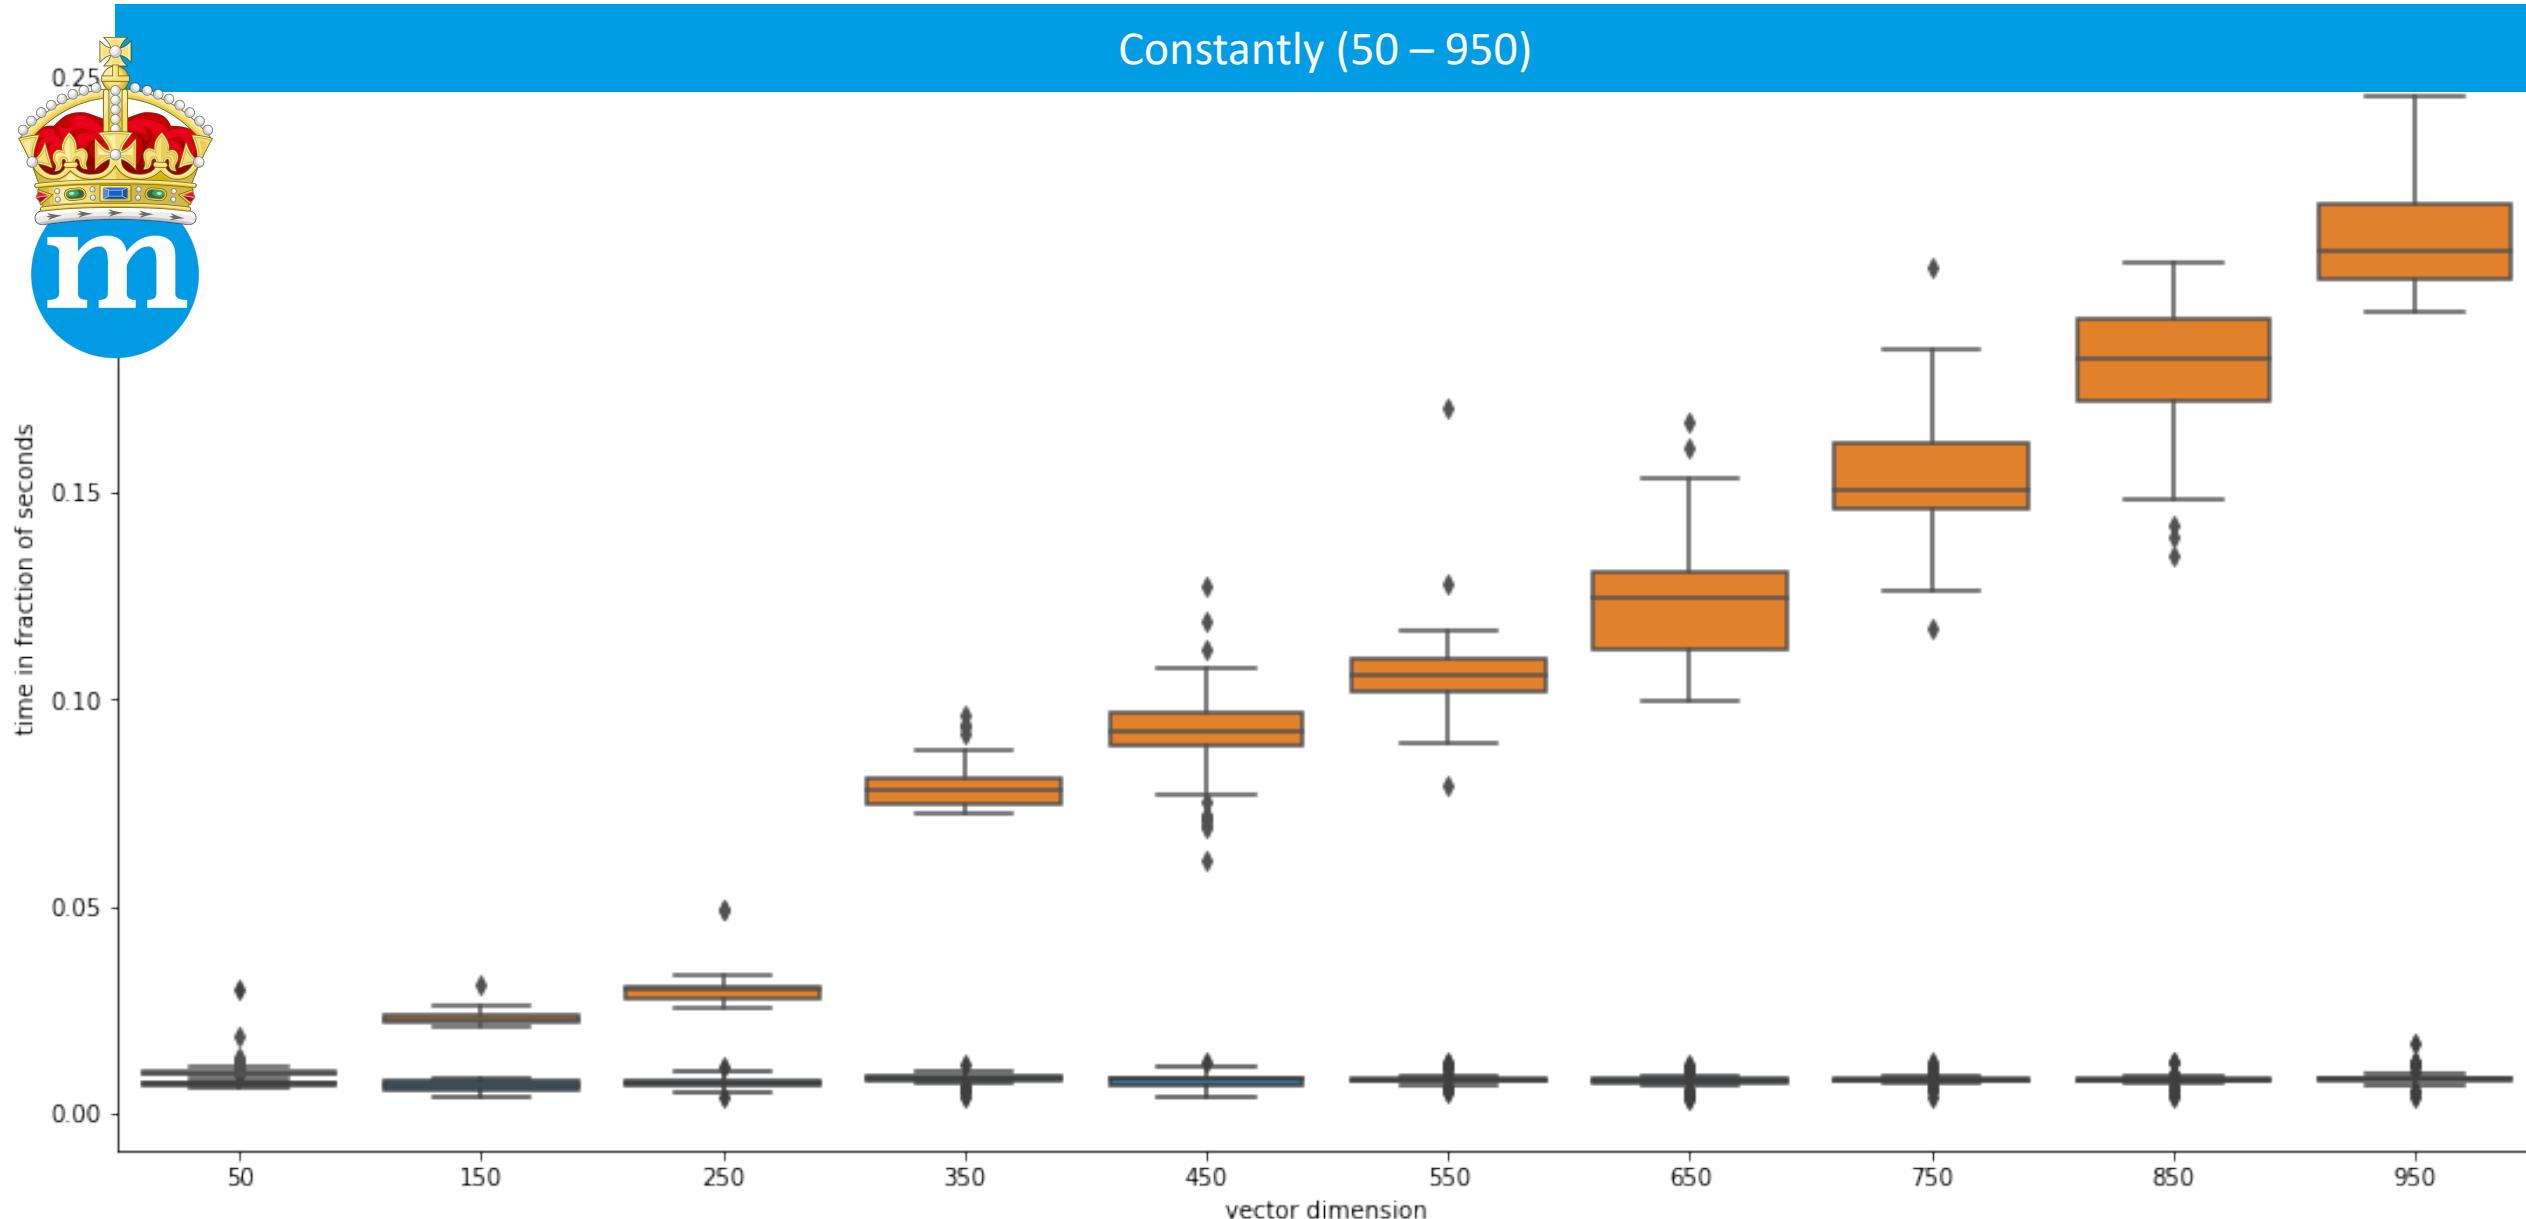
Task 4.5 - Activation on 50 vectors - SoftSign



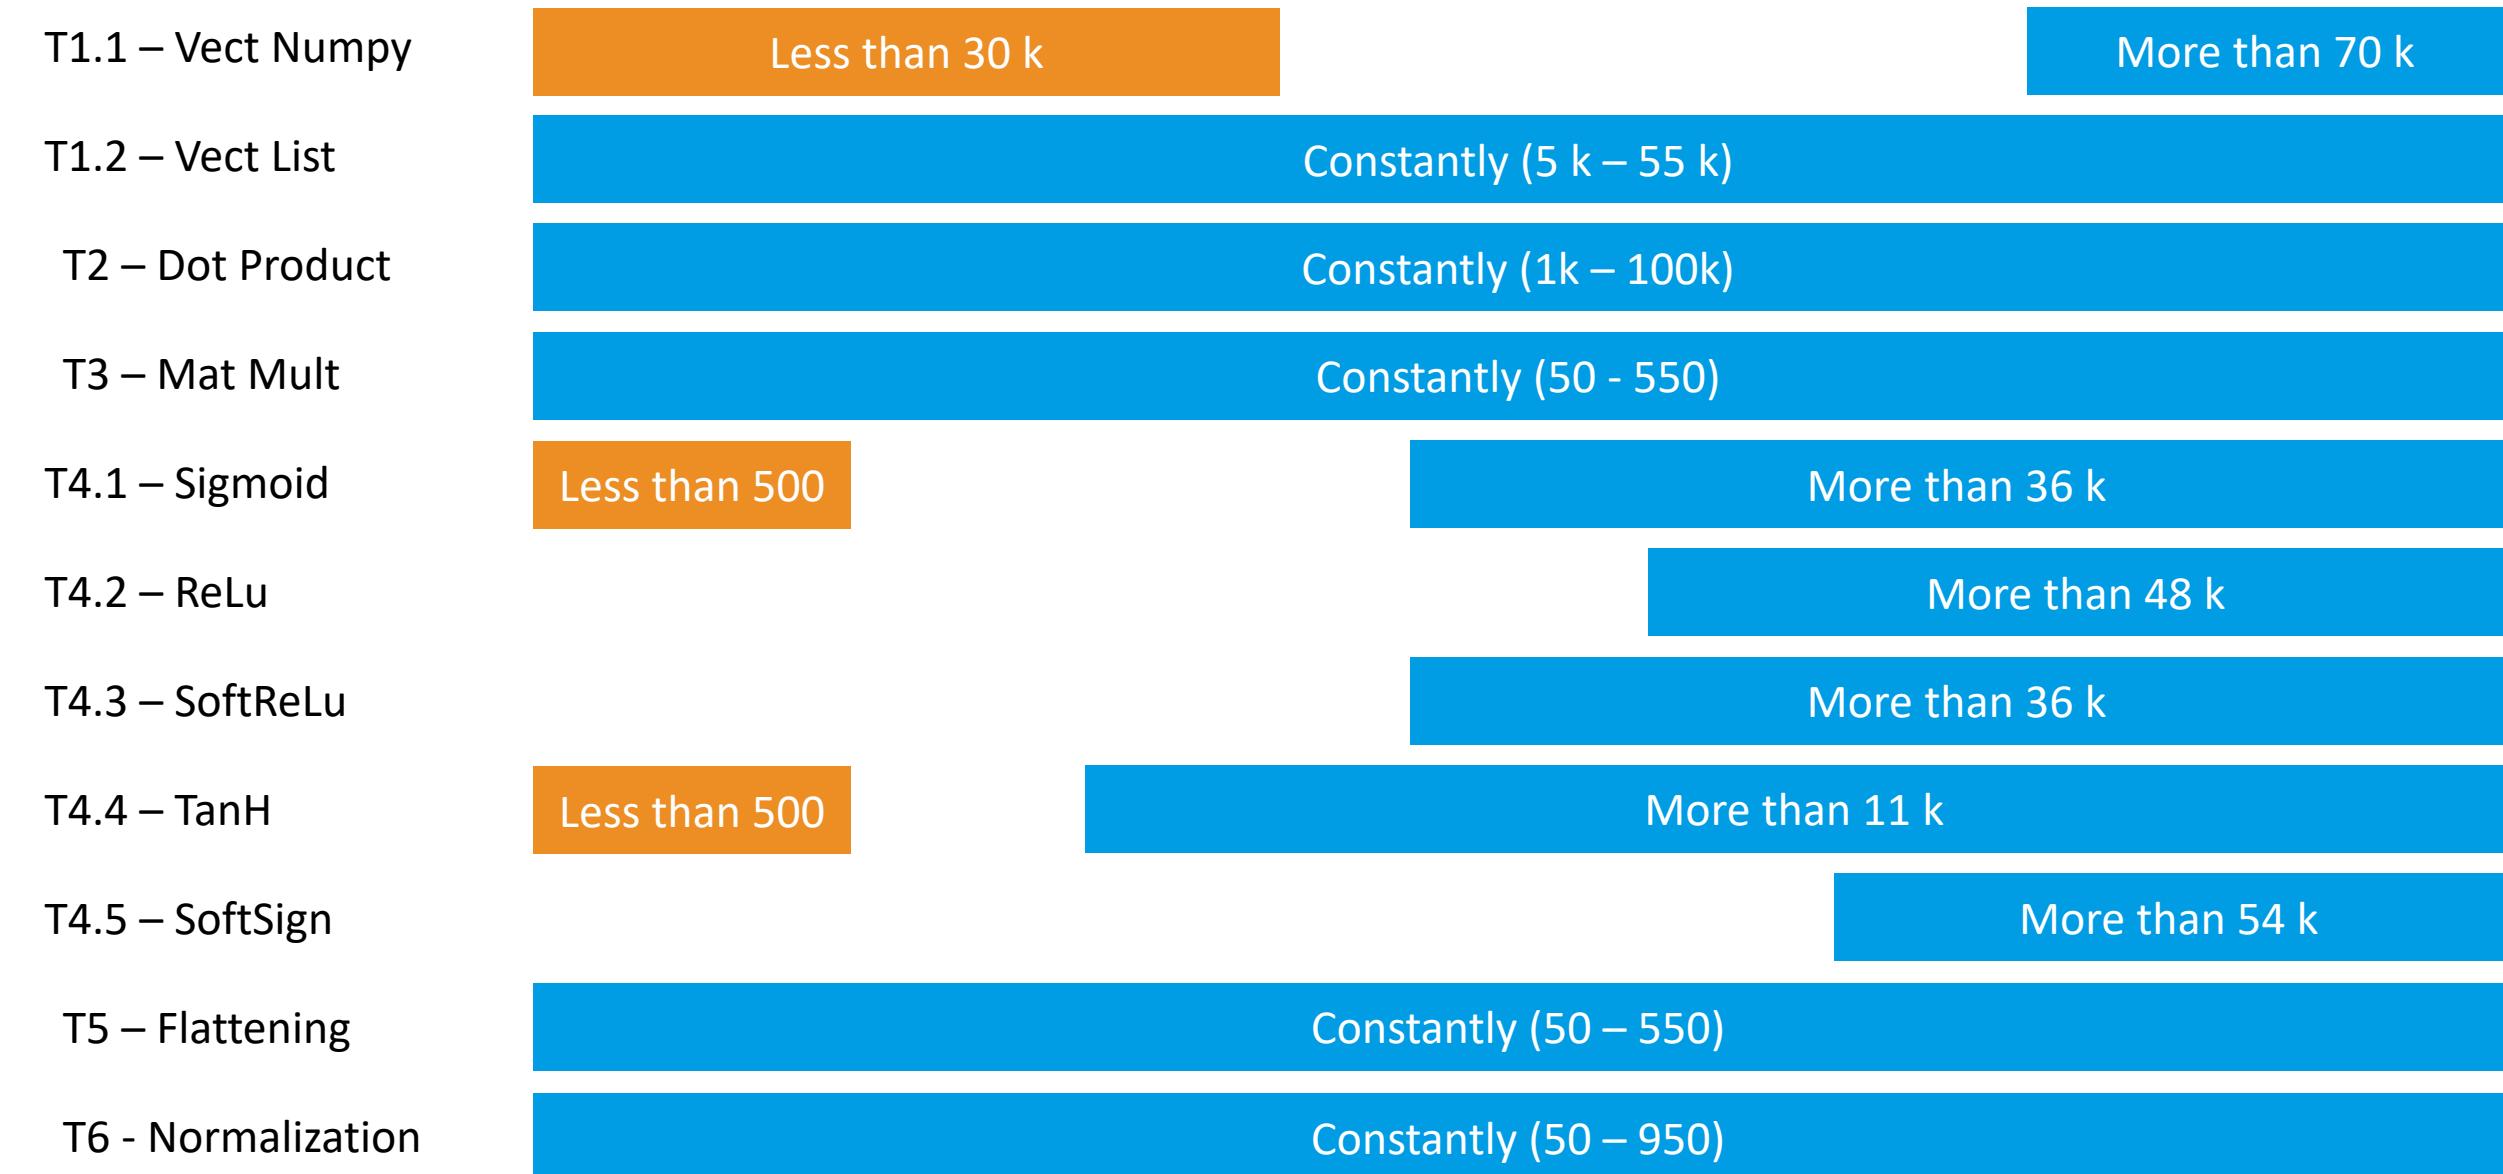
Task 5 – Flattening of 50 square matrices



Task 6 – Normalize 50 matrices



Benchmark Summary – Fundamental Operations





5. LeNet



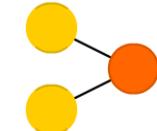
A mostly complete chart of

Neural Networks

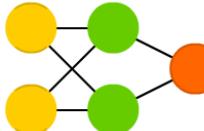
©2016 Fjodor van Veen - asimovinstitute.org

- (○) Backfed Input Cell
- (○) Input Cell
- (△) Noisy Input Cell
- (●) Hidden Cell
- (○) Probabilistic Hidden Cell
- (△) Spiking Hidden Cell
- (●) Output Cell
- (○) Match Input Output Cell
- (●) Recurrent Cell
- (○) Memory Cell
- (△) Different Memory Cell
- (●) Kernel
- (○) Convolution or Pool

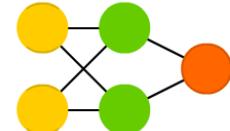
Perceptron (P)



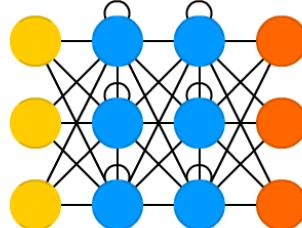
Feed Forward (FF)



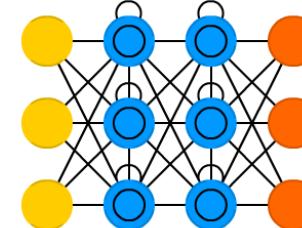
Radial Basis Network (RBF)



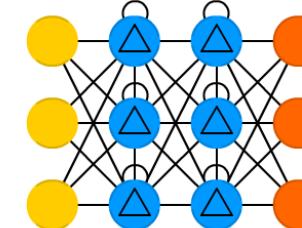
Recurrent Neural Network (RNN)



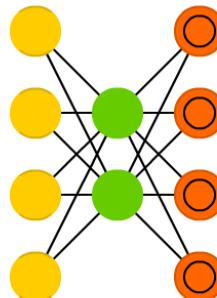
Long / Short Term Memory (LSTM)



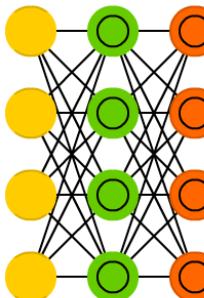
Gated Recurrent Unit (GRU)



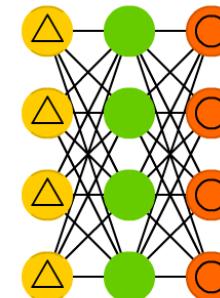
Auto Encoder (AE)



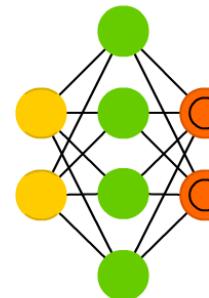
Variational AE (VAE)



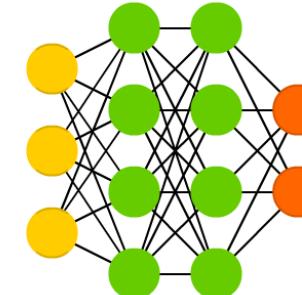
Denoising AE (DAE)



Sparse AE (SAE)

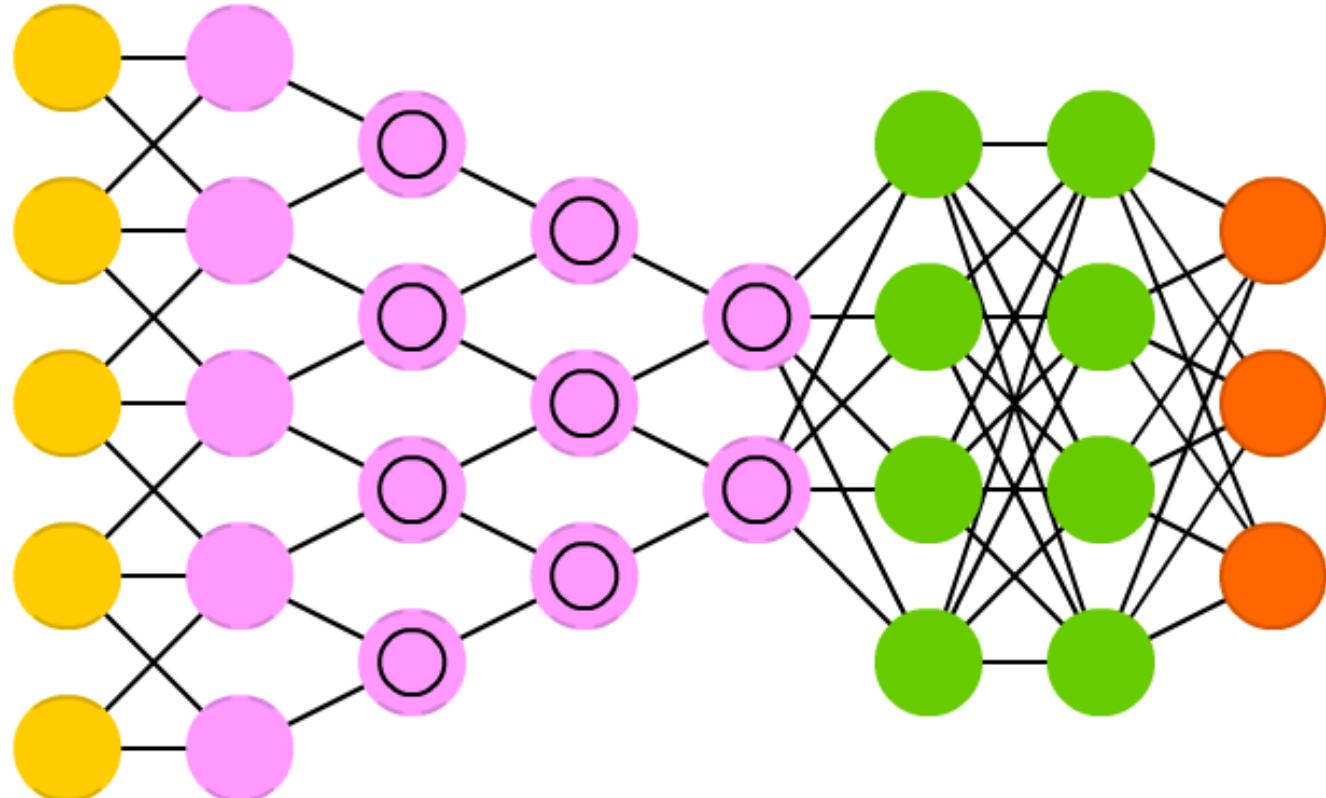


Deep Feed Forward (DFF)



Deep Convolutional Network (DCN)

-  Input Cell
-  Kernel
-  Convolution or Pool
-  Hidden Cell
-  Output Cell



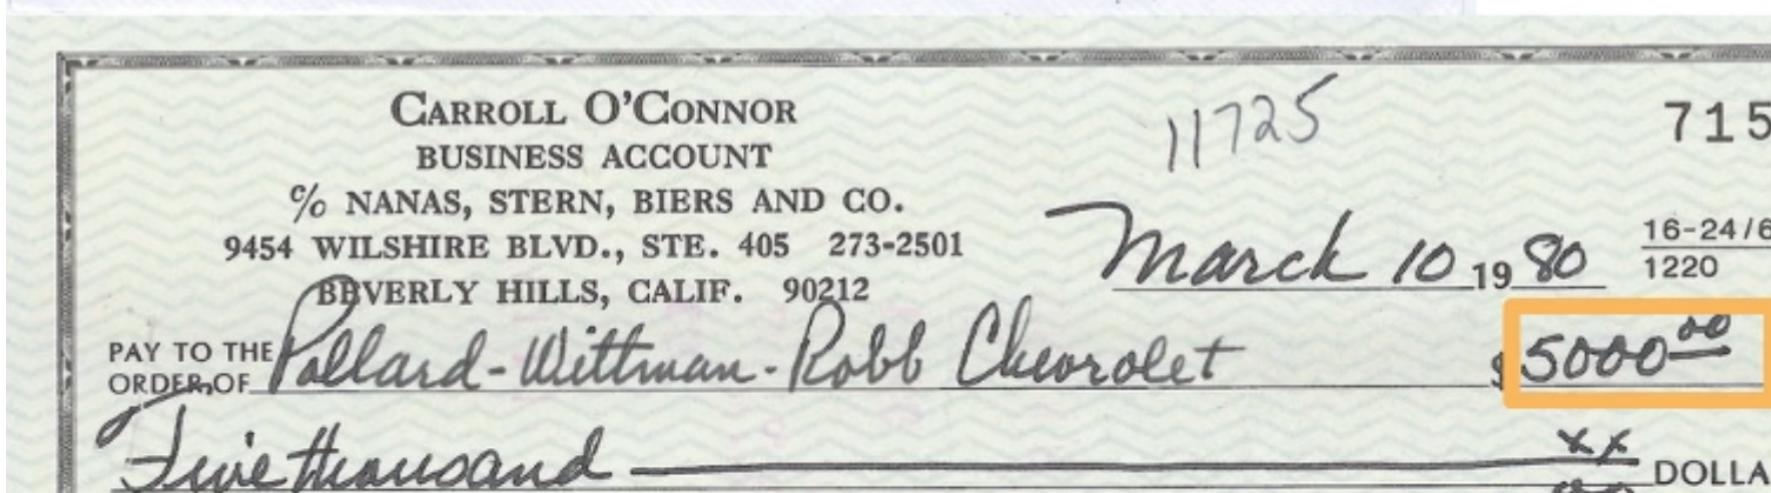
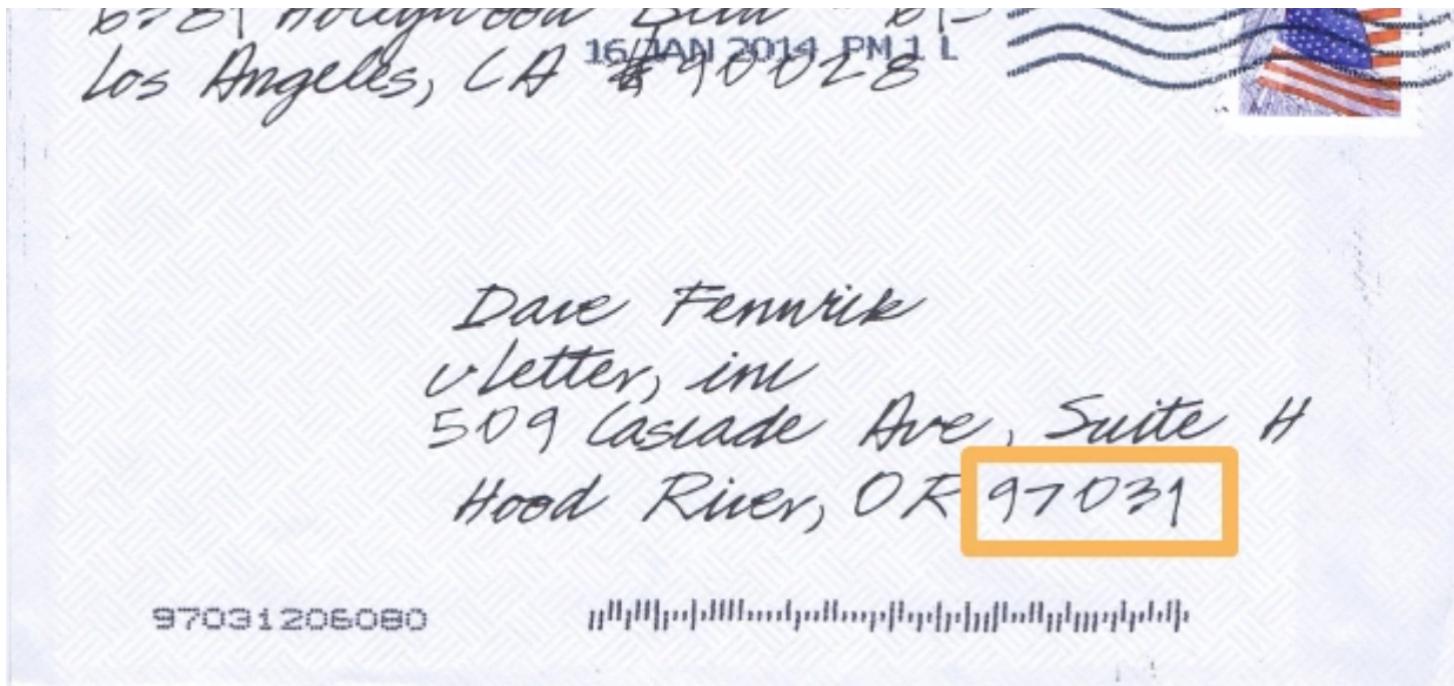


Image source: <https://www.youtube.com/watch?v=m3BrTjo2zUA>

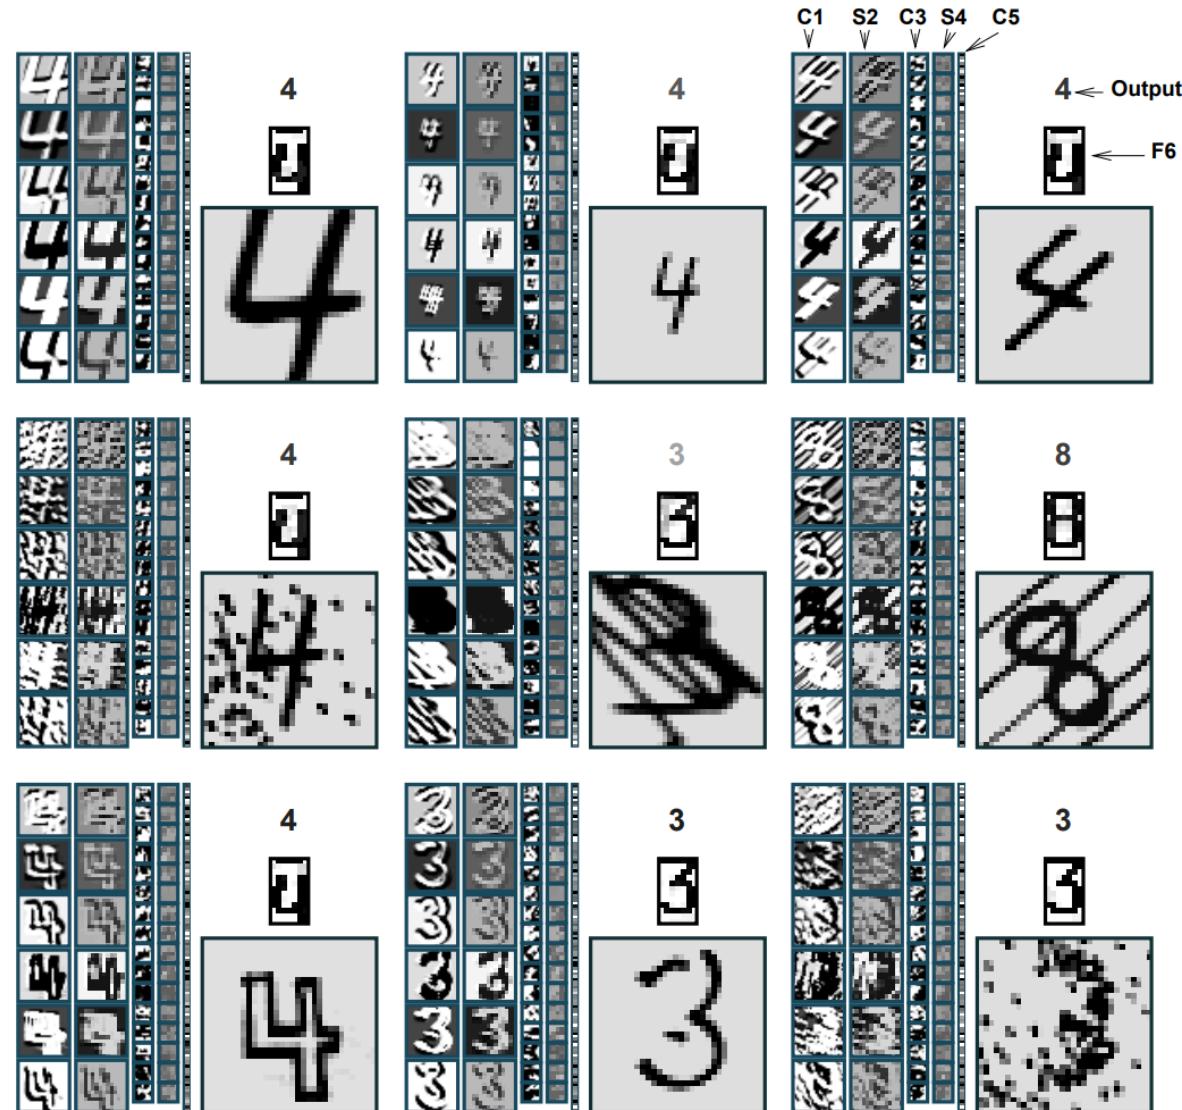


Fig. 13. Examples of unusual, distorted, and noisy characters correctly recognized by LeNet-5. The grey-level of the output label represents the penalty (lighter for higher penalties).

Image source: <http://yann.lecun.com/exdb/publis/pdf/lecun-01a.pdf>

LeNet Architecture

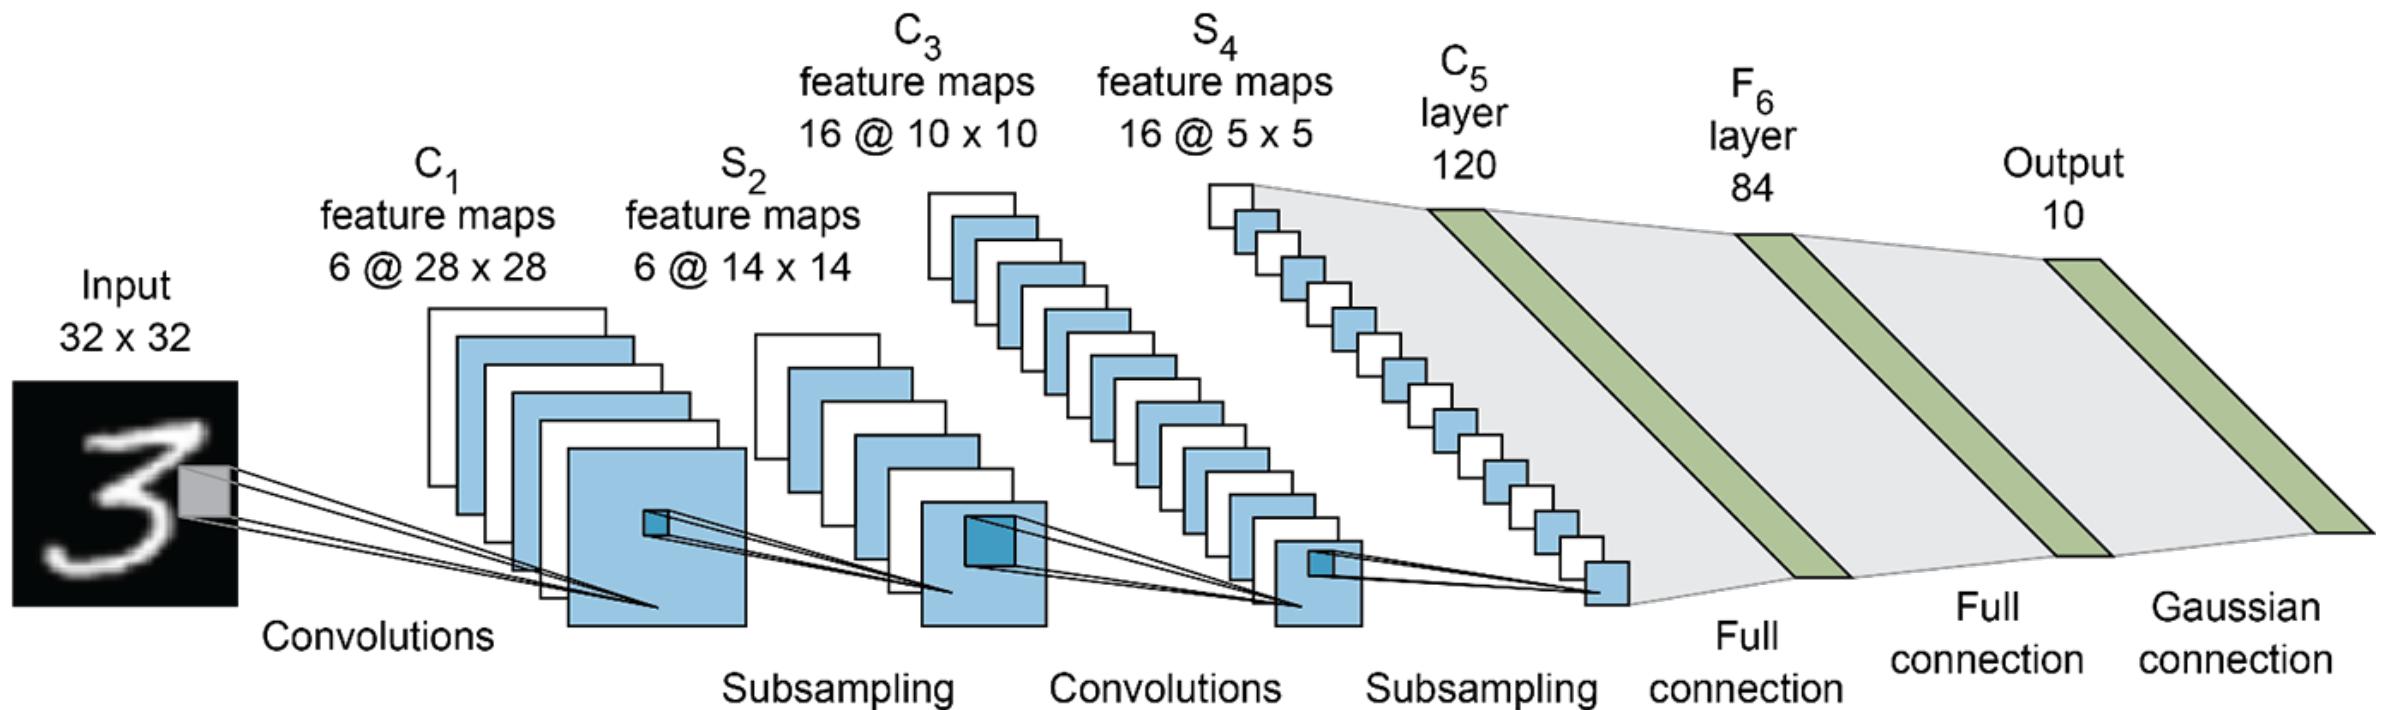


Image source: <https://link.springer.com/article/10.1007/s13244-018-0639-9>

A list of parameters and hyperparameters in a convolutional neural network (CNN)

	Parameters	Hyperparameters
Convolution layer	Kernels	Kernel size, number of kernels, stride, padding, activation function
Pooling layer	None	Pooling method, filter size, stride, padding
Fully connected layer	Weights	Number of weights, activation function
Others		Model architecture, optimizer, learning rate, loss function, mini-batch size, epochs, regularization, weight initialization, dataset splitting

Figure source: <https://link.springer.com/article/10.1007/s13244-018-0639-9>



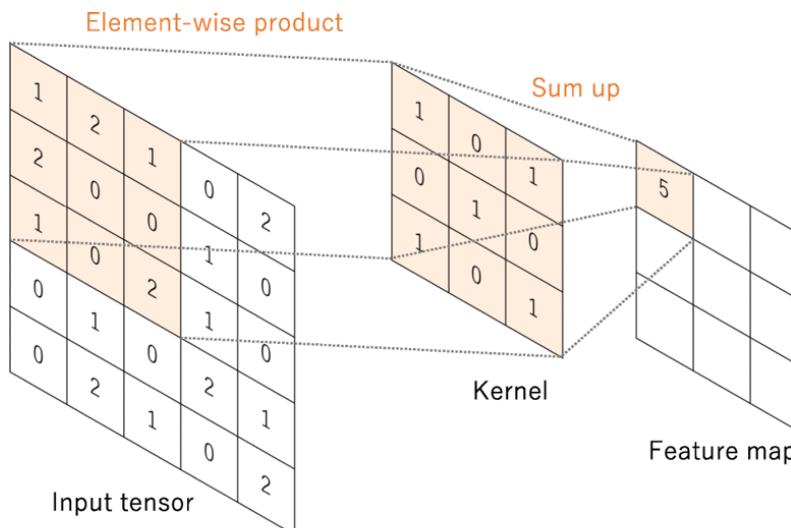
0	2	15	0	0	11	10	0	0	0	0	9	9	0	0	0	0
0	0	0	4	60	157	236	255	255	177	95	61	32	0	0	29	
0	10	16	119	238	255	244	245	243	250	249	255	222	103	10	0	
0	14	170	255	255	244	254	255	253	245	255	249	253	251	124	1	
2	98	255	228	255	251	254	211	141	116	122	215	251	238	255	49	
13	217	243	255	155	33	226	52	2	0	10	13	232	255	255	36	
16	229	252	254	49	12	0	0	7	7	0	70	237	252	235	62	
6	141	245	255	212	25	11	9	3	0	115	236	243	255	137	0	
0	87	252	250	248	215	60	0	1	121	252	255	248	144	6	0	
0	13	113	255	255	245	255	182	181	248	252	242	208	36	0	19	
1	0	5	117	251	255	241	255	247	255	241	162	17	0	7	0	
0	0	0	4	58	251	255	246	254	253	255	120	11	0	1	0	
0	0	4	97	255	255	255	248	252	255	244	255	182	10	0	4	
0	22	206	252	246	251	241	100	24	113	255	245	255	194	9	0	
0	111	255	242	255	158	24	0	0	6	39	255	232	230	56	0	
0	218	251	250	137	7	11	0	0	0	2	62	255	250	125	3	
0	173	255	255	101	9	20	0	13	3	13	182	251	245	61	0	
0	107	251	241	255	230	98	55	19	118	217	248	253	255	52	4	
0	18	146	250	255	247	255	255	249	255	240	255	129	0	5	0	
0	0	23	113	215	255	250	248	255	248	248	118	14	12	0	0	
0	0	6	1	0	52	153	233	255	252	147	37	0	0	4	1	
0	0	5	5	0	0	0	0	14	1	0	6	6	0	0	0	

0	2	15	0	0	11	10	0	0	0	0	9	9	0	0	0	0
0	0	0	4	60	157	236	255	255	177	95	61	32	0	0	29	
0	10	16	119	238	255	244	245	243	250	249	255	222	103	10	0	
0	14	170	255	255	244	254	255	253	245	255	249	253	251	124	1	
2	98	255	228	255	251	254	211	141	116	122	215	251	238	255	49	
13	217	243	255	155	33	226	52	2	0	10	13	232	255	255	36	
16	229	252	254	49	12	0	0	7	7	0	70	237	252	235	62	
6	141	245	255	212	25	11	9	3	0	115	236	243	255	137	0	
0	87	252	250	248	215	60	0	1	121	252	255	248	144	6	0	
0	13	113	255	255	245	255	182	181	248	252	242	208	36	0	19	
1	0	5	117	251	255	241	255	247	255	241	162	17	0	7	0	
0	0	0	4	58	251	255	246	254	253	255	120	11	0	1	0	
0	0	4	97	255	255	255	248	252	255	244	255	182	10	0	4	
0	22	206	252	246	251	241	100	24	113	255	245	255	194	9	0	
0	111	255	242	255	158	24	0	0	6	39	255	232	230	56	0	
0	218	251	250	137	7	11	0	0	0	2	62	255	250	125	3	
0	173	255	255	101	9	20	0	13	3	13	182	251	245	61	0	
0	107	251	241	255	230	98	55	19	118	217	248	253	255	52	4	
0	18	146	250	255	247	255	255	249	255	240	255	129	0	5	0	
0	0	23	113	215	255	250	248	255	248	248	118	14	12	0	0	
0	0	6	1	0	52	153	233	255	252	147	37	0	0	4	1	
0	0	5	5	0	0	0	0	14	1	0	6	6	0	0	0	

Image source: <https://link.springer.com/article/10.1007/s13244-018-0639-9>

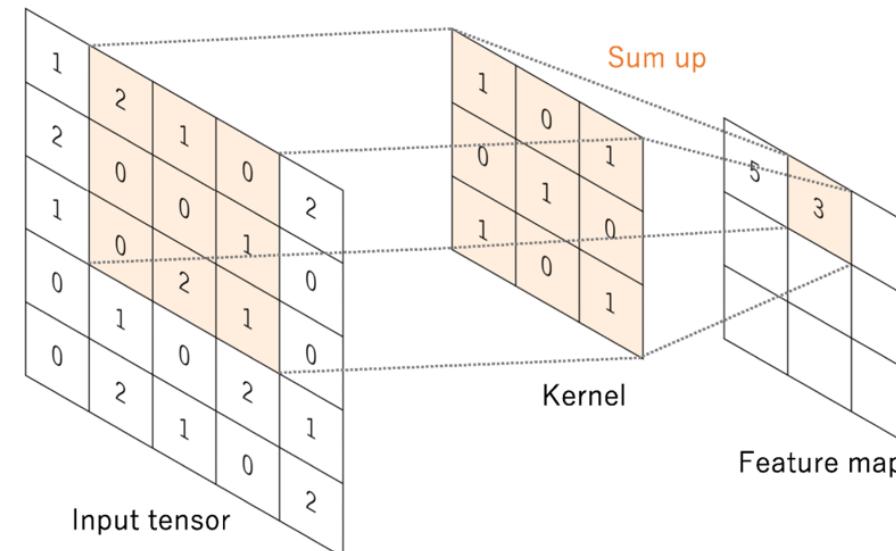
Convolution

a



b

Element-wise product



c

Element-wise product

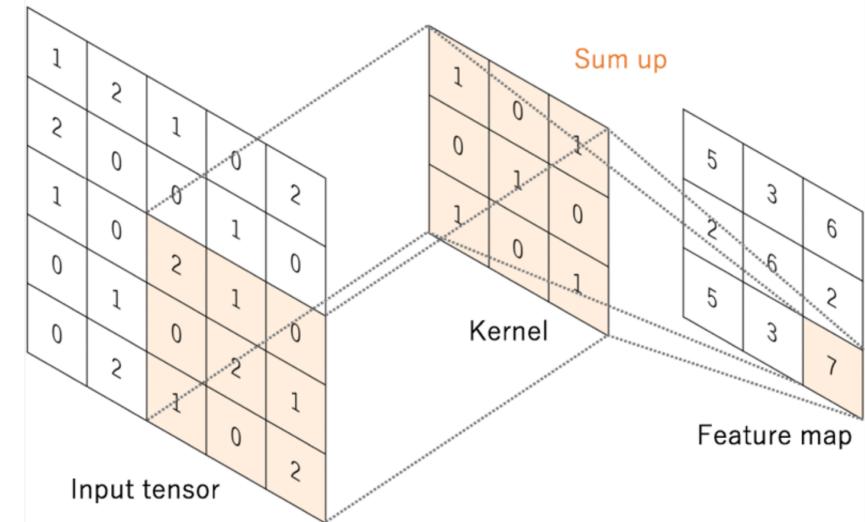
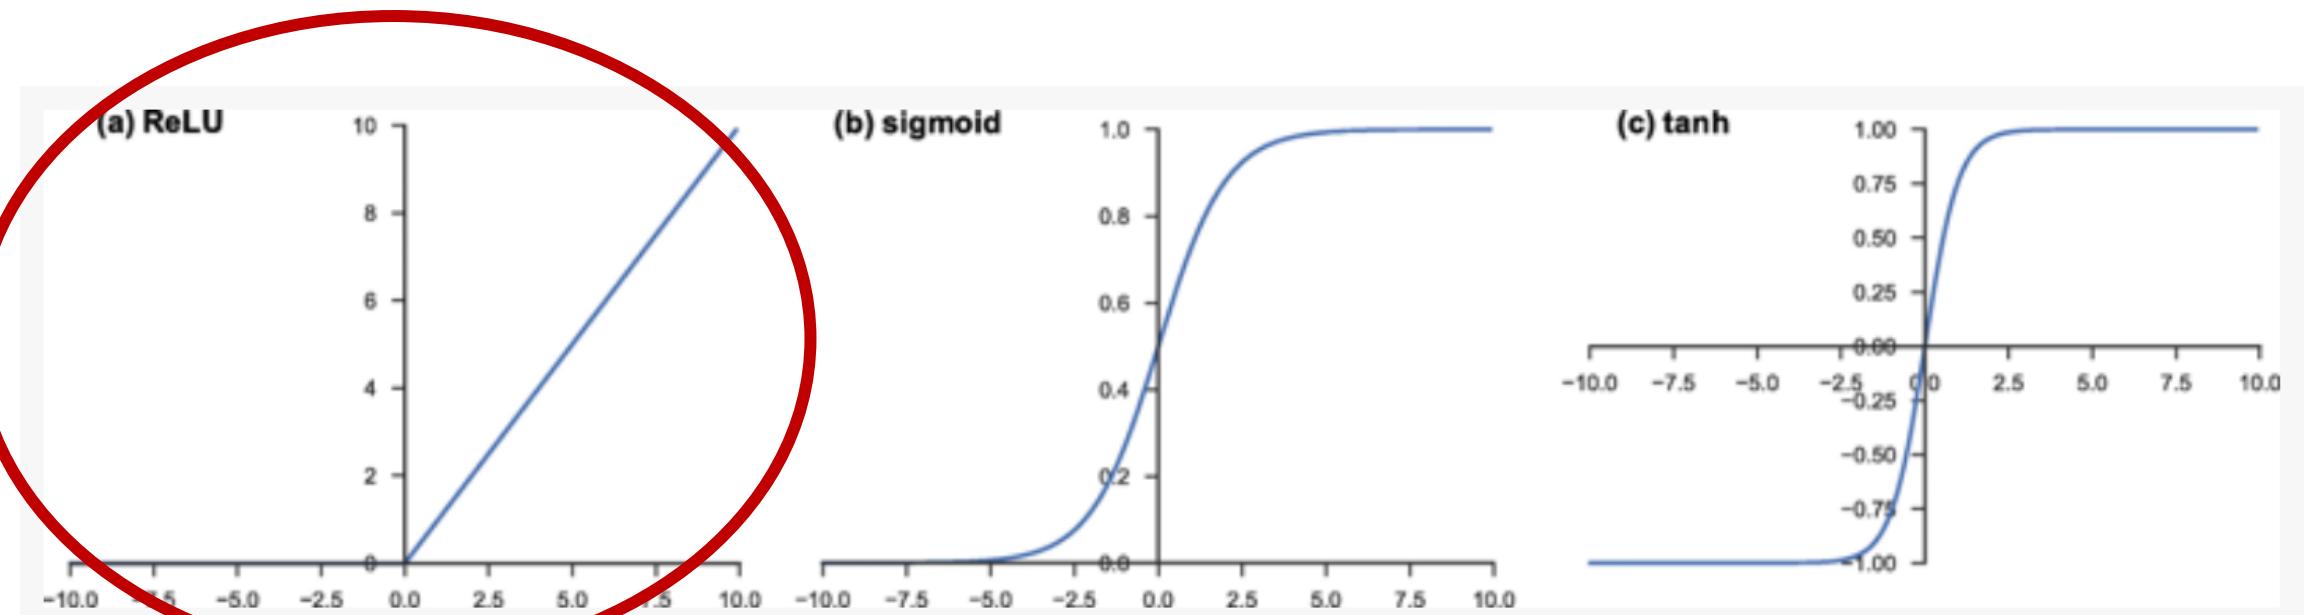
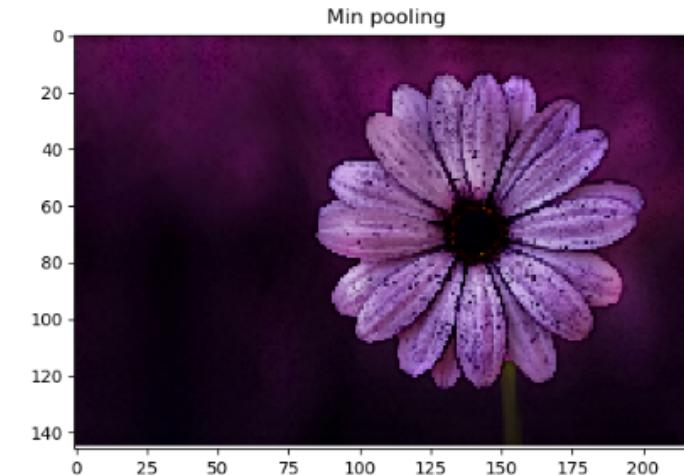
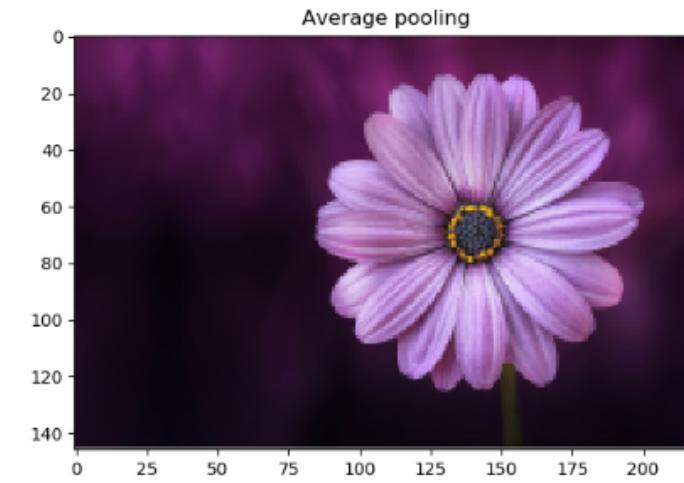
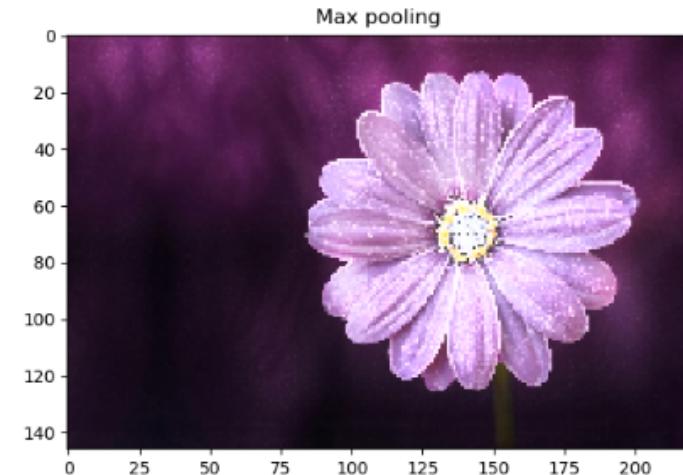


Image source: <https://link.springer.com/article/10.1007/s13244-018-0639-9>



Activation functions commonly applied to neural networks: **a** rectified linear unit (ReLU), **b** sigmoid, and **c** hyperbolic tangent (tanh)

Pooling



Average, Max and Min pooling of size 9x9 applied on an image

Image source: <https://medium.com/@bdhuma/which-pooling-method-is-better-maxpooling-vs-minpooling-vs-average-pooling-95fb03f45a9>

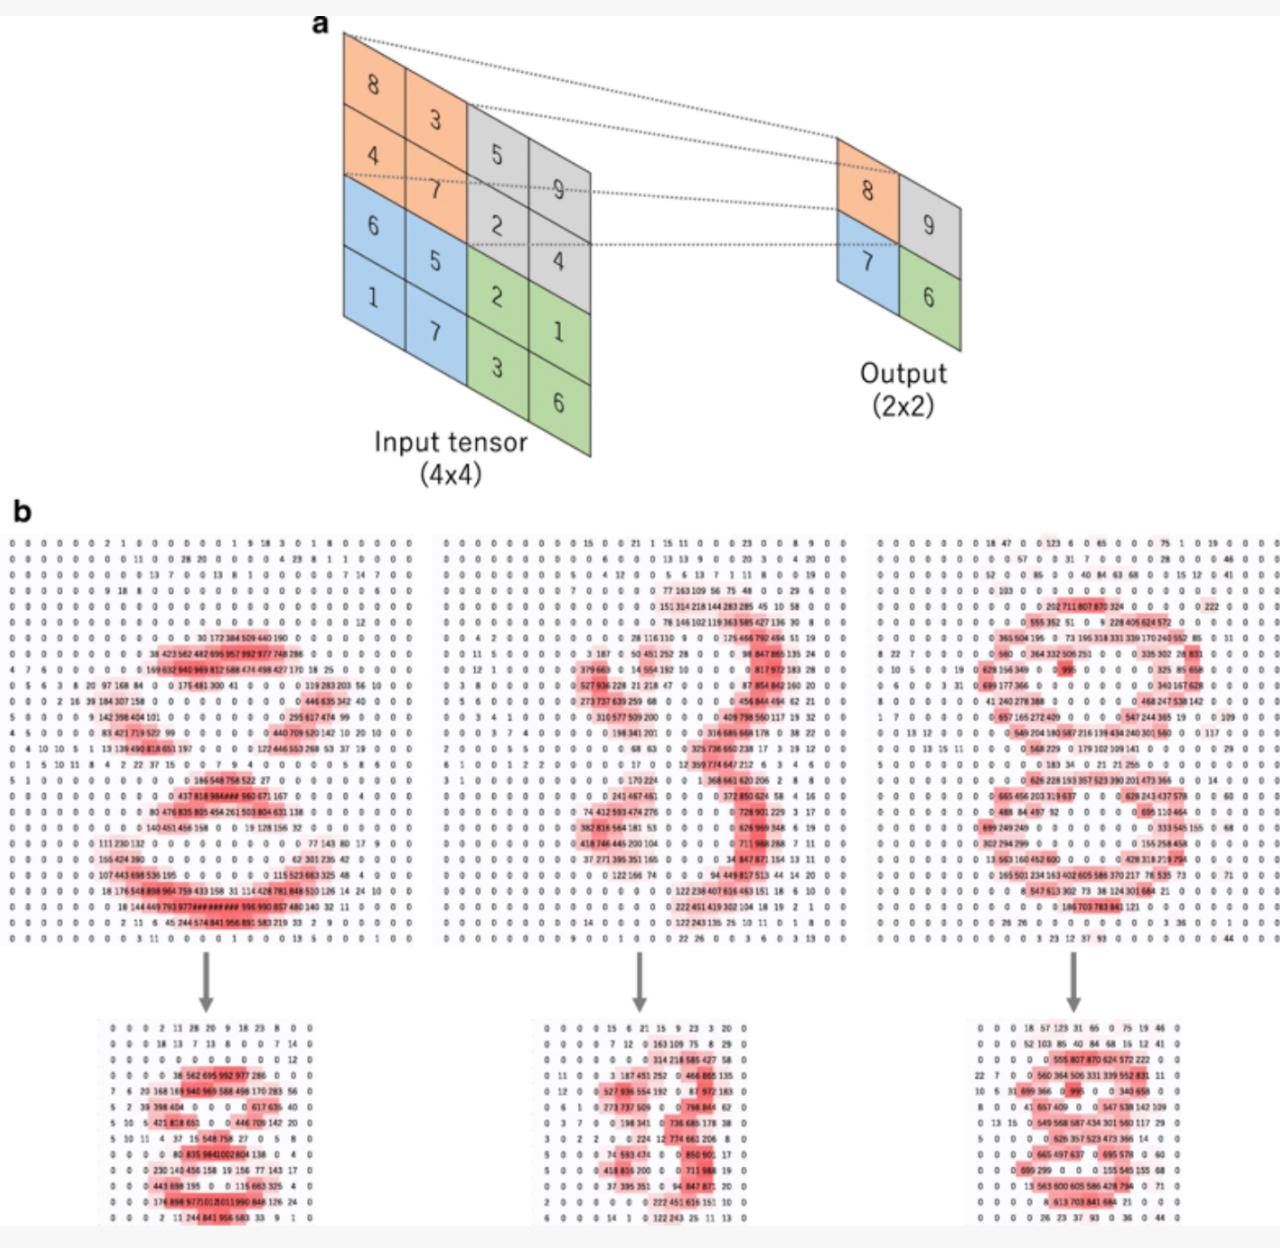


Image source: <https://link.springer.com/article/10.1007/s13244-018-0639-9>

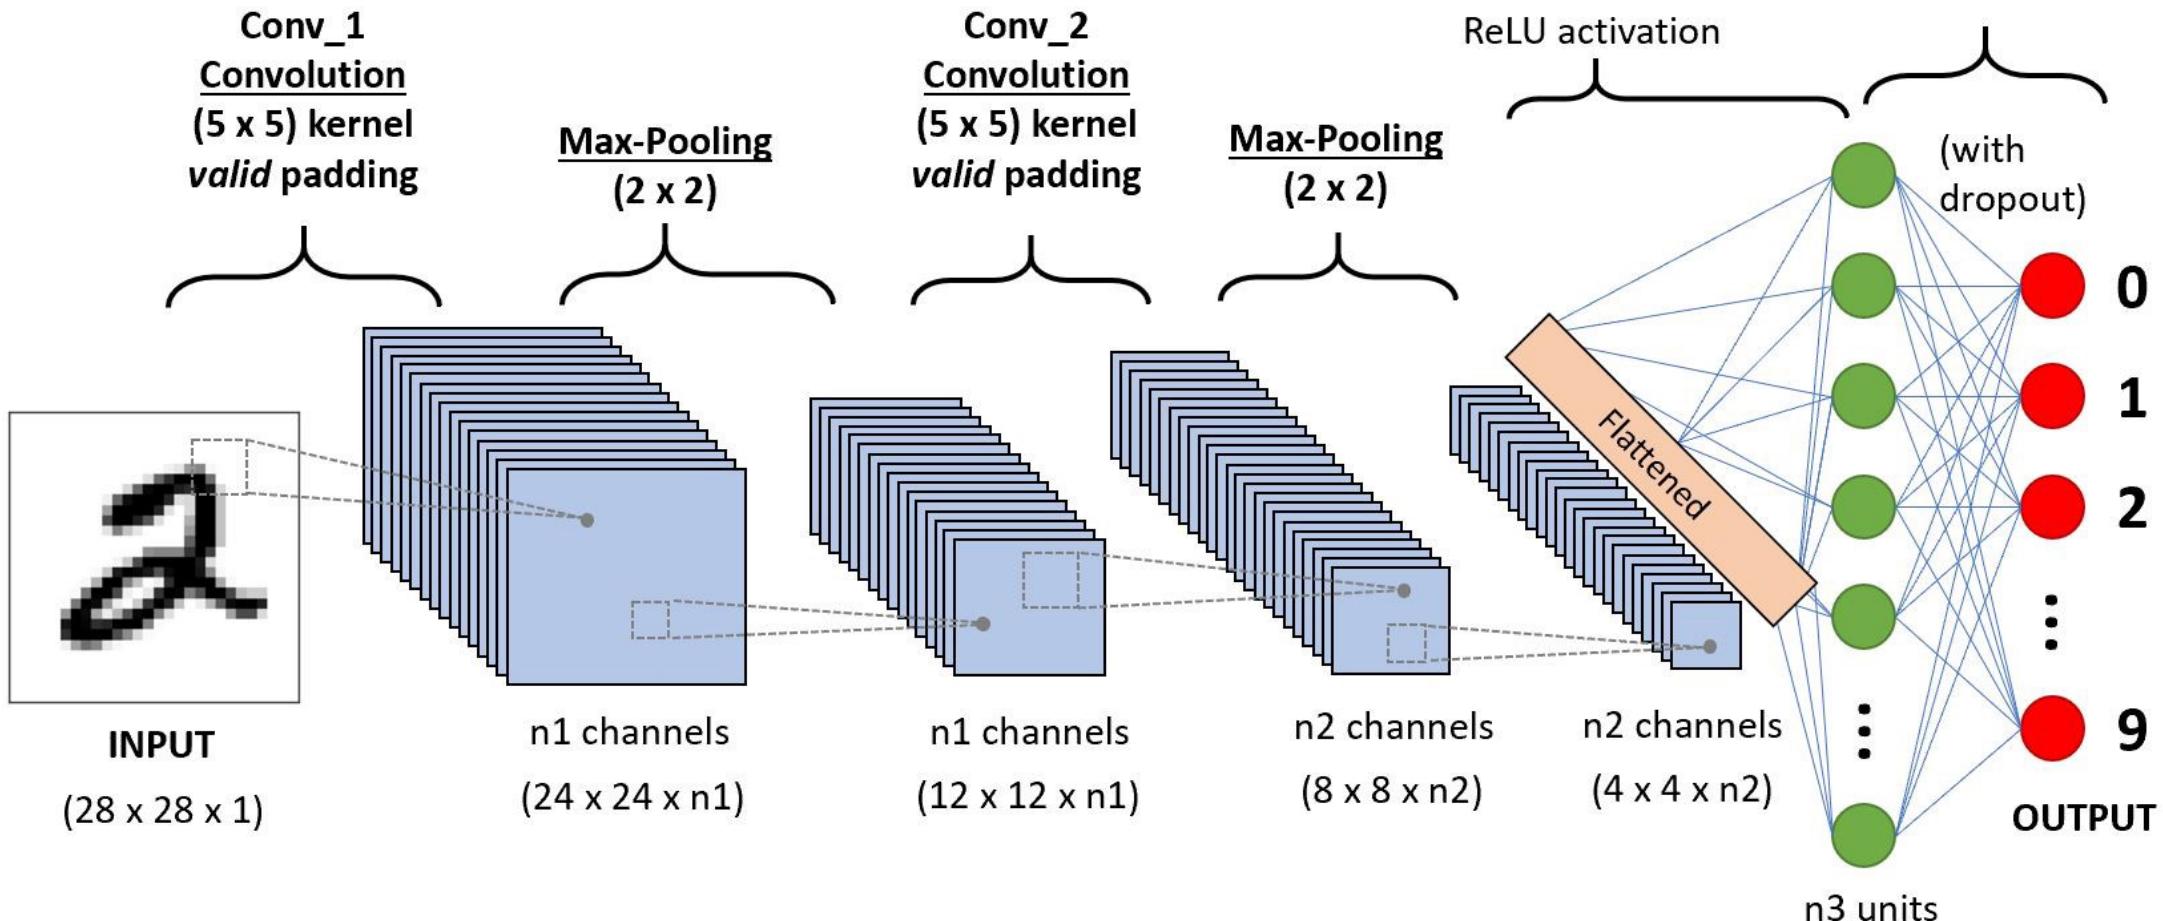


Image source: https://www.researchgate.net/figure/Toy-example-illustrating-the-drawbacks-of-max-pooling-and-average-pooling_fig2_300020038

LeNet Architecture

Layer		Feature Map	Input Size	Kernel Size	Stride	Activation
Input	Image	1	32*32	-	-	-
1	Convolution	6	30*30	5*5	1	Relu
2	Average Pooling	6	15*15	2	2	-
3	Convolution	16	13*13	3*3	1	Relu
4	Average Pooling	16	6*6	2	2	-
Flatten						
5	FC	-	120	-	-	Relu
6	FC	-	84	-	-	Relu
Output	FC	-	10	-	-	Softmax

LeNet Architecture

Hyperparameter	Value
Optimizer	Adam
Learning rate	0.2
Loss function	SOFTMAX-CROSSENTROPY
mini-batch size	200
epochs	10
# images	70000
# training images	60000
# test images	10000
# training images	48000
# validation images	12000

LeNet Construction



```
handwritten_net = nn.Sequential()  
handwritten_net.add(nn.Conv2D(channels=6, kernel_size=5, activation='relu'),  
    nn.MaxPool2D(pool_size=2, strides=2),  
    nn.Conv2D(channels=16, kernel_size=3, activation='relu'),  
    nn.MaxPool2D(pool_size=2, strides=2),  
    nn.Flatten(),  
    nn.Dense(120, activation="relu"),  
    nn.Dense(84, activation="relu"),  
    nn.Dense(10))
```



```
model = keras.Sequential()  
model.add(layers.Conv2D(filters=6, kernel_size=(5, 5), activation='relu', input_shape=(32,32,1)))  
model.add(layers.AveragePooling2D())  
model.add(layers.Conv2D(filters=16, kernel_size=(3, 3), activation='relu'))  
model.add(layers.AveragePooling2D())  
model.add(layers.Flatten())  
model.add(layers.Dense(units=120, activation='relu'))  
model.add(layers.Dense(units=84, activation='relu'))  
model.add(layers.Dense(units=10, activation = 'softmax'))
```

Training Loop



```
In [6]: # TRAIN THE NETWORK WITH ACCURACY
epoch = 15
# CHECK IF GPUs ARE PRESENT
gpus = mxnet.test_utils.list_gpus()
ctx = [mxnet.gpu()] if gpus else [mxnet.cpu(0), mxnet.cpu(1)]
handwritten_net.initialize(mxnet.init.Xavier(), ctx=ctx, force_reinit=True)
trainer = gluon.Trainer(handwritten_net.collect_params(), 'sgd', {'learning_rate': 0.02})
# Use Accuracy as the evaluation metric.
metric = mxnet.metric.Accuracy()
softmax_cross_entropy_loss = gluon.loss.SoftmaxCrossEntropyLoss()

...
# Ask the profiler to start recording
profiler.set_state('run')
...

for i in range(epoch):
    # Reset the train data iterator.
    train_data.reset()
    # Loop over the train data iterator.
    for batch in train_data:
        # Splits train data into multiple slices along batch_axis
        # and copy each slice into a context.
        data = gluon.utils.split_and_load(batch.data[0], ctx_list=ctx, batch_axis=0)
        # Splits train labels into multiple slices along batch_axis
        # and copy each slice into a context.
        label = gluon.utils.split_and_load(batch.label[0], ctx_list=ctx, batch_axis=0)
        outputs = []
        # Inside training scope
        with autograd.record():
            for x, y in zip(data, label):
                z = handwritten_net(x)
                # Computes softmax cross entropy loss.
                loss = softmax_cross_entropy_loss(z, y)
                # Backpropagate the error for one iteration.
                loss.backward()
                outputs.append(z)
        # Updates internal evaluation
        metric.update(label, outputs)
        # Make one step of parameter update. Trainer needs to know the
        # batch size of data to normalize the gradient by 1/batch_size.
        trainer.step(batch.data[0].shape[0])
    # Gets the evaluation result.
    name, acc = metric.get()
    # Reset evaluation result to initial state.
    metric.reset()
    print('training acc at epoch %d: %s=%f'%(i, name, acc))
...
```



```
In [18]: EPOCHS = 15
BATCH_SIZE = 200

In [19]: X_train, y_train = train['features'], to_categorical(train['labels'])
X_validation, y_validation = validation['features'], to_categorical(validation['labels'])

train_generator = ImageDataGenerator().flow(X_train, y_train, batch_size=BATCH_SIZE)
validation_generator = ImageDataGenerator().flow(X_validation, y_validation, batch_size=BATCH_SIZE)

In [20]: print('# of training images:', train['features'].shape[0])
print('# of validation images:', validation['features'].shape[0])

steps_per_epoch = X_train.shape[0]//BATCH_SIZE
validation_steps = X_validation.shape[0]//BATCH_SIZE

#log_dir = "logs/profile/{}".format(time())
#tensorboard_callback = tf.keras.callbacks.TensorBoard(log_dir = log_dir, histogram_freq=1, profile_batch = 3)

model.fit_generator(train_generator, steps_per_epoch=steps_per_epoch, epochs=EPOCHS,
                    validation_data=validation_generator, validation_steps=validation_steps)
# shuffle=True, callbacks=[tensorboard_callback])
```

Accuracy Results



```
mxnet.nd.waitall()
# Ask the profiler to stop recording
profiler.set_state('stop')
# Dump all results to log file before download
profiler.dump()
...

training acc at epoch 0: accuracy=0.553133
training acc at epoch 1: accuracy=0.883817
training acc at epoch 2: accuracy=0.912917
training acc at epoch 3: accuracy=0.927500
training acc at epoch 4: accuracy=0.937967
training acc at epoch 5: accuracy=0.944550
training acc at epoch 6: accuracy=0.949750
training acc at epoch 7: accuracy=0.954333
training acc at epoch 8: accuracy=0.956900
training acc at epoch 9: accuracy=0.960633
training acc at epoch 10: accuracy=0.963350
training acc at epoch 11: accuracy=0.965267
training acc at epoch 12: accuracy=0.967050
training acc at epoch 13: accuracy=0.968633
training acc at epoch 14: accuracy=0.970317

Out[6]: "\n# Make sure all operations have completed\nmxnet.nd.waitall()\n# Ask the profiler to stop recording\nprofiler.set_state('stop')\n# Dump all results to log file before download\nprofiler.dump() \n"

In [7]: # TEST THE NETWORK
metric = mxnet.metric.Accuracy()
# Reset the validation data iterator.
val_data.reset()
# Loop over the validation data iterator.
for batch in val_data:
    # Splits validation data into multiple slices along batch_axis
    # and copy each slice into a context.
    data = gluon.utils.split_and_load(batch.data[0], ctx_list=ctx, batch_axis=0)
    # Splits validation label into multiple slices along batch_axis
    # and copy each slice into a context.
    label = gluon.utils.split_and_load(batch.label[0], ctx_list=ctx, batch_axis=0)
    outputs = []
    for x in data:
        outputs.append(handwritten_net(x))
    # Updates internal evaluation
    metric.update(label, outputs)
print('validation acc: %s=%f'%metric.get())
assert metric.get()[1] > 0.90
validation acc: accuracy=0.970400
```



```
In [18]: EPOCHS = 15
BATCH_SIZE = 200

In [19]: X_train, y_train = train['features'], to_categorical(train['labels'])
X_validation, y_validation = validation['features'], to_categorical(validation['labels'])

train_generator = ImageDataGenerator().flow(X_train, y_train, batch_size=BATCH_SIZE)
validation_generator = ImageDataGenerator().flow(X_validation, y_validation, batch_size=BATCH_SIZE)

In [20]: print('# of training images:', train['features'].shape[0])
print('# of validation images:', validation['features'].shape[0])

steps_per_epoch = X_train.shape[0]//BATCH_SIZE
validation_steps = X_validation.shape[0]//BATCH_SIZE

#log_dir = "logs/profile/{}".format(time())
#tensorboard_callback = tf.keras.callbacks.TensorBoard(log_dir = log_dir, histogram_freq=1, profile_batch = 3)

model.fit_generator(train_generator, steps_per_epoch=steps_per_epoch, epochs=EPOCHS,
                    validation_data=validation_generator, validation_steps=validation_steps,
                    # shuffle=True, callbacks=[tensorboard_callback])

# of training images: 48000
# of validation images: 12000
Epoch 1/15
240/240 [=====] - 98s 408ms/step - loss: 0.5154 - accuracy: 0.9023 - val_loss: 0.0899
- val_accuracy: 0.9718
Epoch 2/15
240/240 [=====] - 92s 385ms/step - loss: 0.0821 - accuracy: 0.9756 - val_loss: 0.0758
- val_accuracy: 0.9776
Epoch 3/15
240/240 [=====] - 68s 282ms/step - loss: 0.0563 - accuracy: 0.9826 - val_loss: 0.0572
- val_accuracy: 0.9825
Epoch 4/15
240/240 [=====] - 70s 292ms/step - loss: 0.0409 - accuracy: 0.9875 - val_loss: 0.0682
- val_accuracy: 0.9788
Epoch 5/15
240/240 [=====] - 80s 333ms/step - loss: 0.0331 - accuracy: 0.9890 - val_loss: 0.0571
- val_accuracy: 0.9829
Epoch 6/15
240/240 [=====] - 95s 395ms/step - loss: 0.0275 - accuracy: 0.9911 - val_loss: 0.0570
- val_accuracy: 0.9835
Epoch 7/15
240/240 [=====] - 97s 406ms/step - loss: 0.0242 - accuracy: 0.9919 - val_loss: 0.0524
- val_accuracy: 0.9854
Epoch 8/15
240/240 [=====] - 101s 422ms/step - loss: 0.0175 - accuracy: 0.9945 - val_loss: 0.0551
```



TF

	5 epochs			10 epochs			15 epochs		
#	Time	Accuracy	#	Time	Accuracy	#	Time	Accuracy	
1	19.4	0.312883	1	24.68	0.560933	1	23.91	0.553133	
2	17.16	0.816533	2	27.1	0.882233	2	20.29	0.883817	
3	34.13	0.897083	3	21.48	0.9128	3	19.82	0.912917	
4	16.96	0.91745	4	20.09	0.9284	4	19.96	0.9275	
5	1.04	0.932433	5	20.34	0.93835	5	24.16	0.937967	
Total validation	128.71		6	20.27	0.9458	6	23.3	0.94455	
			7	20.67	0.950633	7	18.58	0.94975	
			8	20.16	0.954017	8	19.61	0.954333	
			9	30.3	0.957183	9	20.18	0.9569	
			10	10.31	0.960133	10	20.69	0.960633	
			Total validation	215.4		11	19.8	0.960633	
					0.9602	12	21.12	0.96335	
						13	19.49	0.96705	
						14	19.97	0.968633	
						15	14.99	0.970317	
MXnet	Total validation	305.87							
						Total validation	0.9704		
#	Time	Accuracy	#	Time	Accuracy	#	Time	Accuracy	
1	68	0.8415	1	90	0.8959	1	98	0.9023	
2	75	0.9756	2	101	0.9753	2	92	0.9756	
3	84	0.9827	3	98	0.9826	3	68	0.9826	
4	65	0.9864	4	83	0.9867	4	70	0.9875	
5	67	0.9899	5	93	0.9885	5	80	0.989	
Total validation	359		6	99	0.9905	6	95	0.9911	
			7	101	0.9926	7	97	0.9919	
			8	105	0.9935	8	101	0.9945	
			9	100	0.9936	9	78	0.9944	
			10	100	0.9953	10	76	0.9957	
			Total validation	970		11	93	0.9951	
					0.9857	12	100	0.9955	
						13	78	0.9952	
						14	75	0.997	
						15	77	0.9964	
TF	Total validation	1278							
						Total validation	0.9863		



TF

5 epochs			10 epochs			15 epochs			
#	Time	Accuracy	#	Time	Accuracy	#	Time	Accuracy	
1	19.4	0.312883	1	24.68	0.560933	1	23.91	0.553133	
2	17.16	0.816533	2	27.1	0.882233	2	20.29	0.883817	
3	34.13	0.897083	3	21.48	0.9128	3	19.82	0.912917	
4	16.96	0.91745	4	20.09	0.9284	4	19.96	0.9275	
5	1.04	0.932433	5	20.34	0.93835	5	24.16	0.937967	
Total validation	128.71		6	20.27	0.9458	6	23.3	0.94455	
	0.9336		7	20.67	0.950633	7	18.58	0.94975	
			8	20.16	0.954017	8	19.61	0.954333	
			9	30.3	0.957183	9	20.18	0.9569	
			10	10.31	0.960133	10	20.69	0.960633	
	215.4		Total validation	0.9602		11	19.8	0.960633	
			Total validation	305.87		12	21.12	0.96335	
				0.9704		13	19.49	0.96705	
				14		14	19.97	0.968633	
				15		15	14.99	0.970317	
				0.9704					
MXnet			#	Time	Accuracy	#	Time	Accuracy	
1	68	0.8415	1	90	0.8959	1	98	0.9023	
2	75	0.9756	2	101	0.9753	2	92	0.9756	
3	84	0.9827	3	98	0.9826	3	68	0.9826	
4	65	0.9864	4	83	0.9867	4	70	0.9875	
5	67	0.9899	5	93	0.9885	5	80	0.989	
Total validation	359		6	99	0.9905	6	95	0.9911	
	0.9847		7	101	0.9926	7	97	0.9919	
			8	105	0.9935	8	101	0.9945	
			9	100	0.9936	9	78	0.9944	
			10	100	0.9953	10	76	0.9957	
	970		Total validation	0.9857		11	93	0.9951	
				1278		12	100	0.9955	
				0.9863		13	78	0.9952	
				14		14	75	0.997	
				15		15	77	0.9964	

LeNet Comparison - Summary

TensorFlow - Keras

- **More accuracy**
- Takes way more time to train
- **Faster convergence**
- More declarative
- Easier to implement
- Extensive documentation
- Wider community
- Minor reproducibility issues

MXNet - Gluon

- Slightly less accuracy
- **More Speed in training**
- **'Slower' convergence**
- Less declarative
- 'harder' to implement
- Fewer documentation
- Smaller community
- **Reproducibility issues**

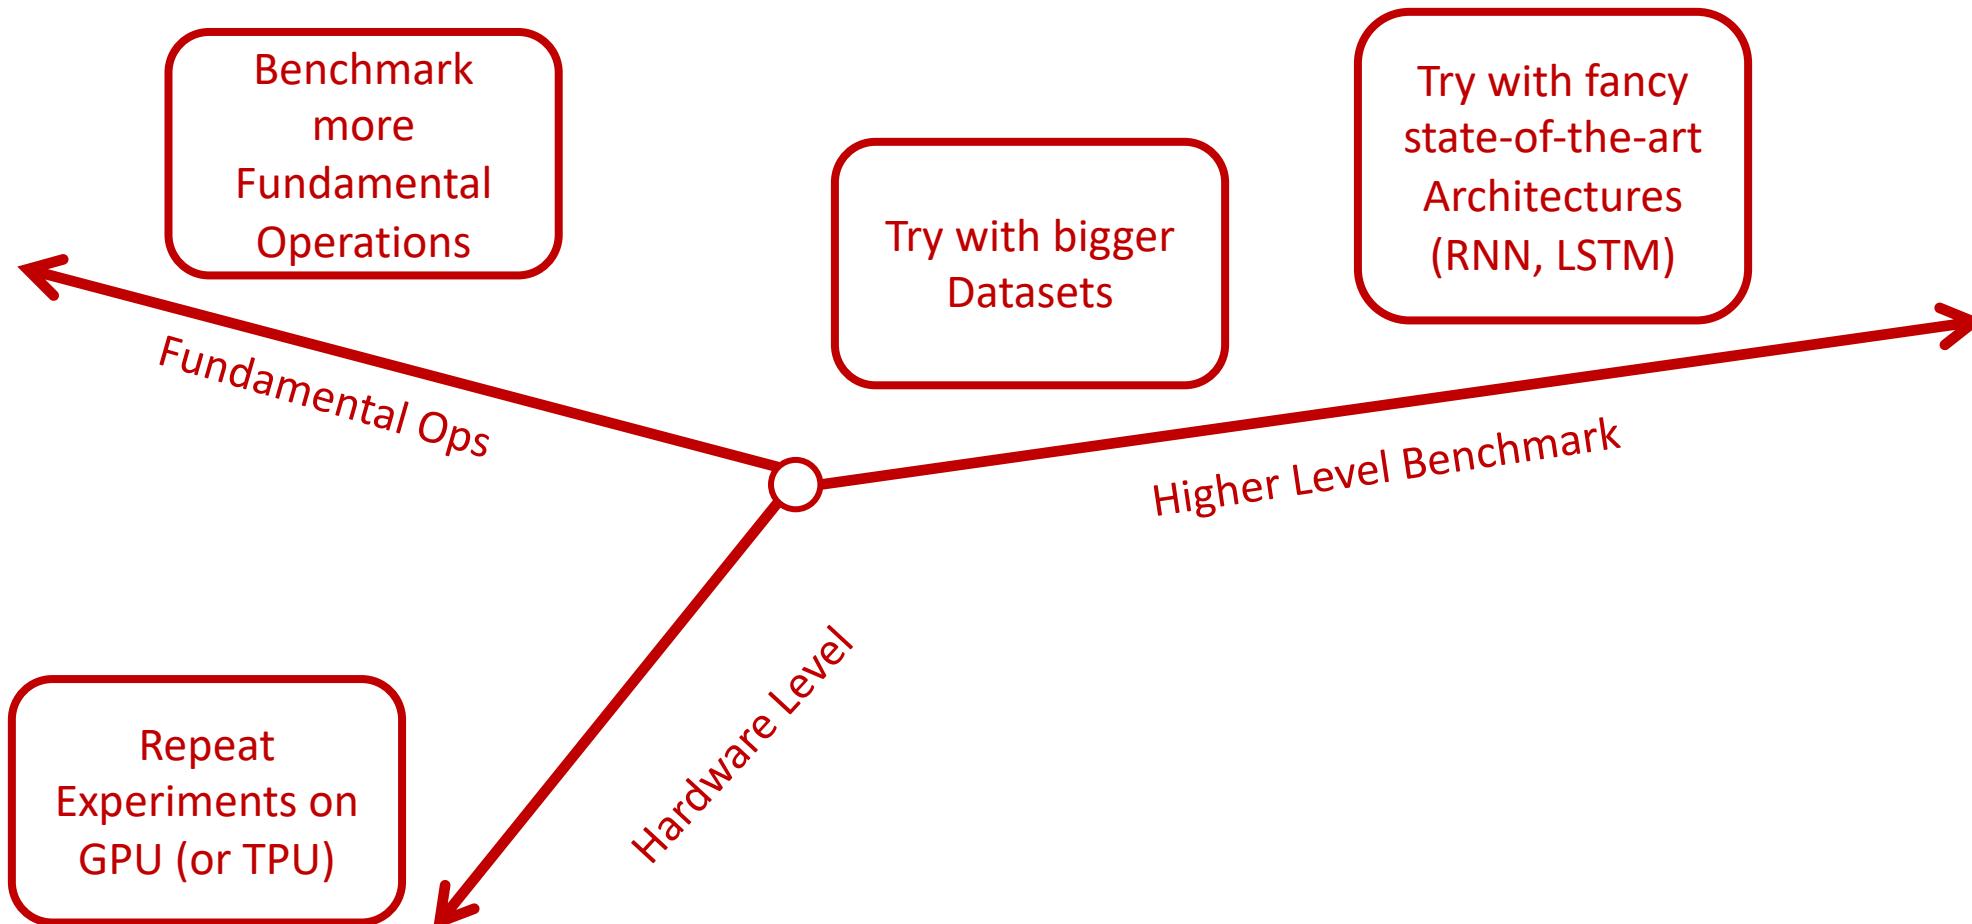


6. Summary – Future Work

Summary

- Two architectures are very similar in terms of high level features
- Same mixed paradigm approach: declarative (more efficient) and imperative (more intuitive)
- Very different in terms of performance speed (on CPU) both for fundamental operations and complete training
- Two huge documentations: TF bigger but not well organized, Mxnet smaller but in constant expansion with high quality resources.

Future Work – Possible Directions



Questions?

Suggestions?

Thank you!

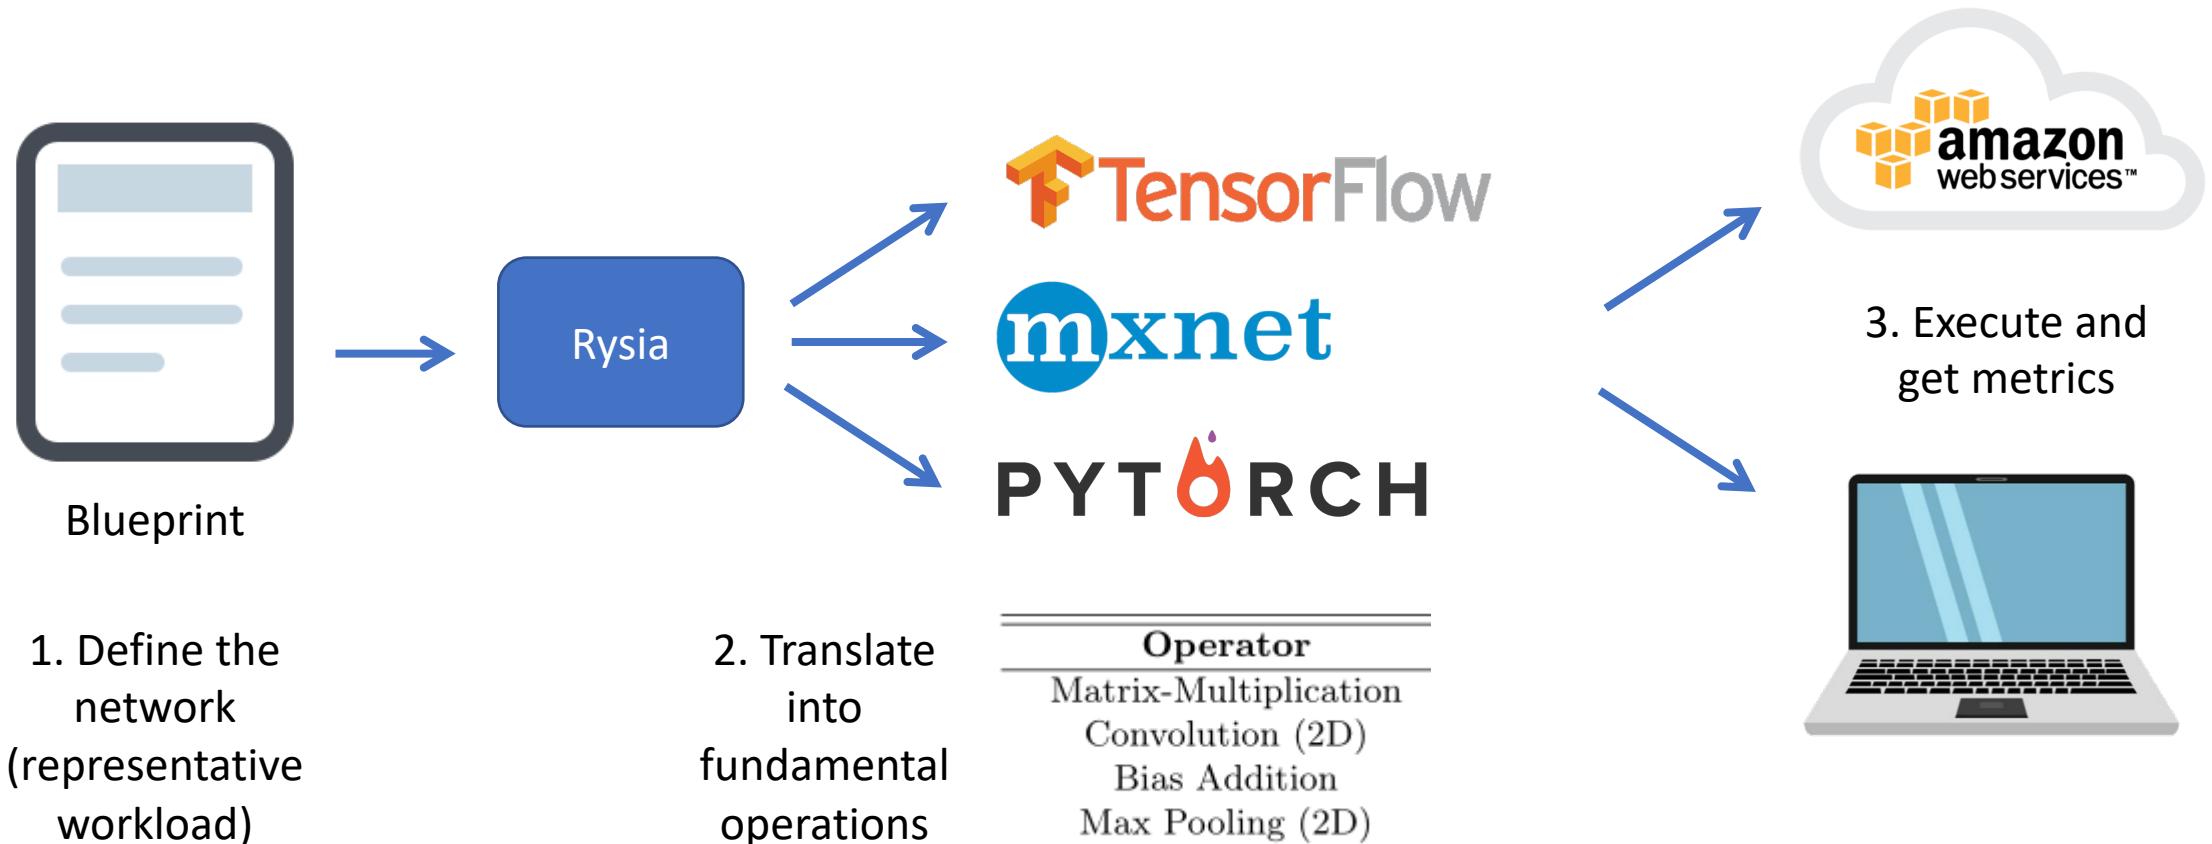




x. Porting Rysia Framework to TensorFlow 2.0

Paper: End-to-End Benchmarking of Deep Learning Platforms

Rysia: Declarative Benchmarking Framework



Porting

Major Release
1.12.2 -> 2.0.0



1.4.1 -> 1.5.1



1.12.2

```
W = tf.Variable(  
    tf.glorot_uniform_initializer()  
        (10, 10))  
b = tf.Variable(tf.zeros(10))  
c = tf.Variable(0)  
  
x = tf.placeholder(tf.float32)  
ctr = c.assign_add(1)  
with tf.control_dependencies([ctr]):  
    y = tf.matmul(x, W) + b  
init =  
    tf.global_variables_initializer()  
  
with tf.Session() as sess:  
    sess.run(init)  
    print(sess.run(y,  
        feed_dict={x: make_input_value()}))  
    assert int(sess.run(c)) == 1
```

2.0.0

```
W = tf.Variable(  
    tf.glorot_uniform_initializer()  
        (10, 10))  
b = tf.Variable(tf.zeros(10))  
c = tf.Variable(0)  
  
@tf.function  
def f(x):  
    c.assign_add(1)  
    return tf.matmul(x, W) + b  
  
print(f(make_input_value()))  
assert int(c) == 1
```